



**MySQL**



---

# Il linguaggio SQL

Giovani programmati e sviluppatori nel settore ICT



# Introduction to MySQL

# MySQL Websites

---

- <http://www.mysql.com> includes:
  - Product information
  - Services (Training, Certification, Consulting, and Support)
  - White papers, webinars, and other resources
  - MySQL Enterprise Edition downloads (trial versions)
- <http://dev.mysql.com> includes:
  - Developer Zone (forums, articles, Planet MySQL, and more)
  - Documentation
  - Downloads
- <https://github.com/mysql>
  - Source code for MySQL Server and other MySQL products



# Community Resources

---

- Mailing lists
- Forums
- Developer articles
- MySQL Newsletter  
(published monthly)
- Planet MySQL blogs
- Social media channels
  - Facebook, Twitter, and Google+
- Physical and virtual events, including:
  - Developer days
  - MySQL Tech Tours
  - Webinars
- Bug tracking
- Github repositories



# MySQL-Supported Operating Systems

---

- MySQL:
  - Provides control and flexibility for users
  - Supports multiple commodity platforms, including:
    - Windows (x86, x86\_64)
    - Linux (x86, x86\_64)
    - Oracle Solaris (SPARC\_64, x86\_64, x86)
    - Mac OS X (x86, x86\_64)
  - Can be compiled to run on other platforms

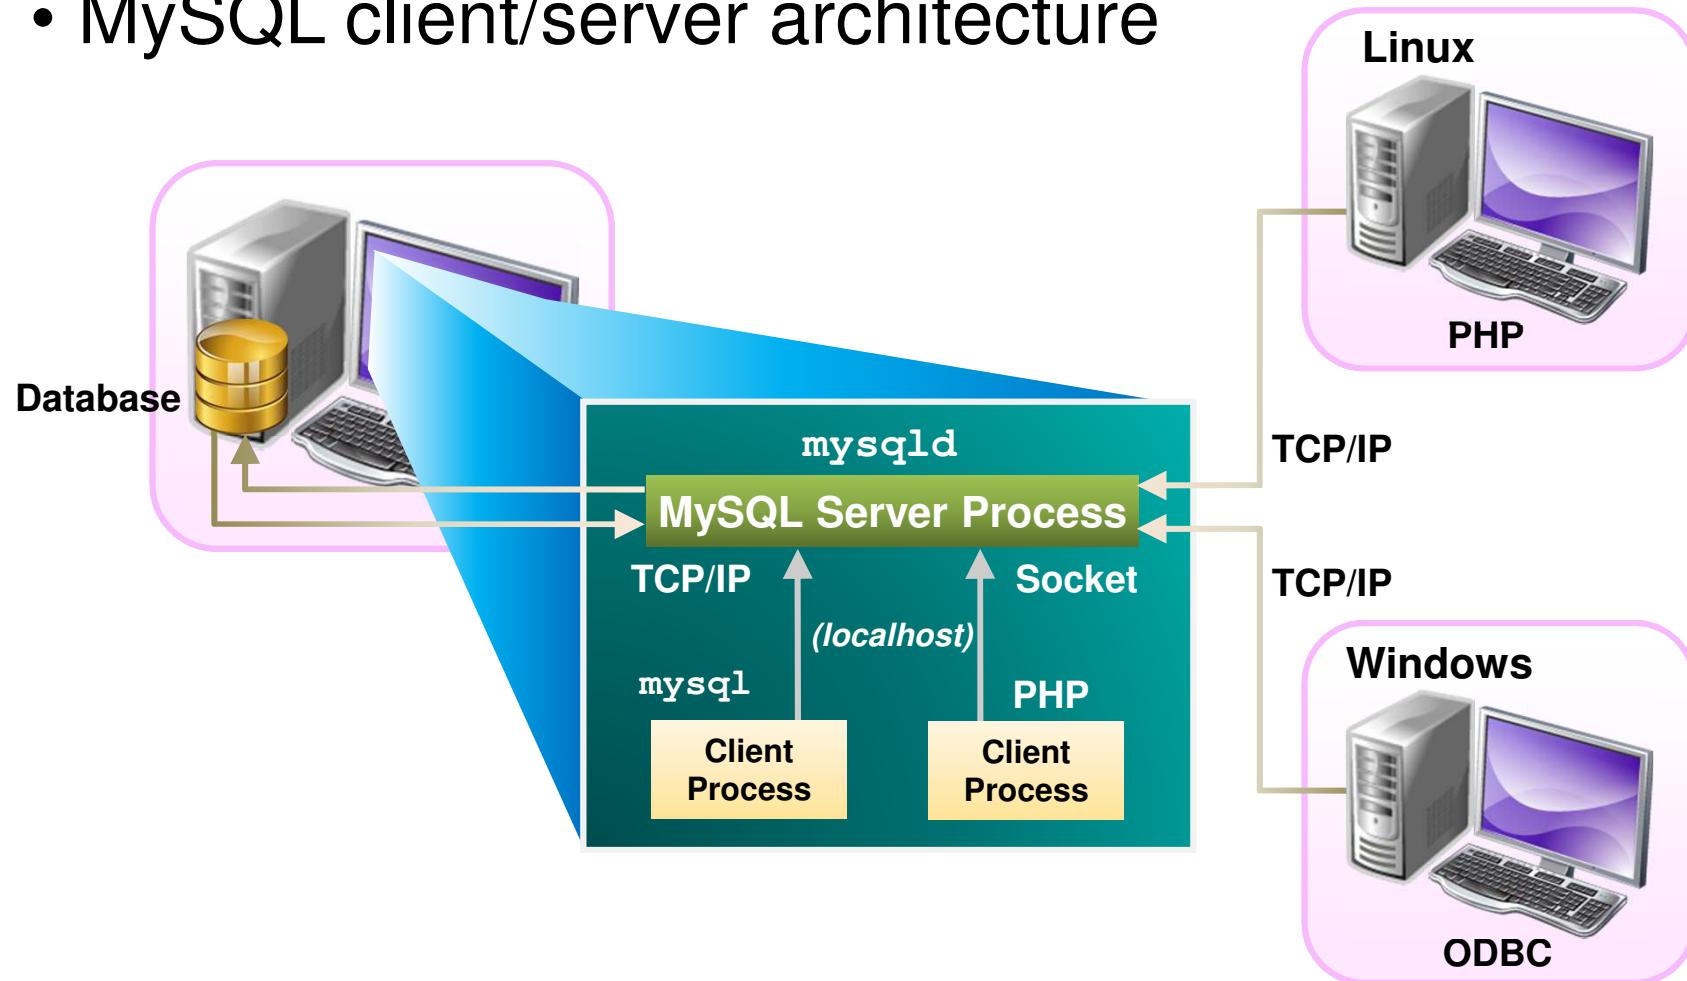


# MySQL Server and Client

- Connecting to the Database

# MySQL Client/Server Model

- MySQL client/server architecture



# MySQL Connectors

---

- Clients are applications that connect to a MySQL server and issue database statements.
- Connectors (drivers) enable clients to access data on the server.
  - These are available for multiple operating systems.
- The official MySQL-supported connectors include:
  - Connector/ODBC
  - Connector/J
  - Connector/.NET
  - Connector/Python
  - Connector/C and Connector C++
- The embedded MySQL server library is `libmysql.dll`.
- Multiple PHP extensions and APIs are supported.



# Installing MySQL

---

- The download locations for MySQL include:
  - Community Edition (GPL): <http://dev.mysql.com/downloads>
  - Commercial editions: <https://edelivery.oracle.com/>
- MySQL is available for multiple platforms, including Windows, Linux, and Mac OS X.
  - This course uses the Linux operating system.
  - Install MySQL on Linux by using Yum with RPM packages.
  - Additional RPM packages are available for MySQL management tools.
- Example databases for testing MySQL features are available at <http://dev.mysql.com/doc/index-other.html>.
  - This course uses the `world` database for examples and practices.



# Starting and Stopping MySQL

---

- You can start the MySQL server on Linux in the following ways:
  - From the Linux command line
  - Automatically on boot
  - With MySQL Workbench
- In Windows, the MySQL service starts automatically at system startup.
  - You can manage the MySQL service from the Windows Management Console.
  - You can start or stop MySQL server from MySQL Workbench.



# Starting and Stopping the MySQL Server on Linux

- On Linux, start the MySQL server and verify the status with the following commands:

```
# service mysqld start
Starting mysqld:                                     [ OK ]
# service mysqld status
mysqld (pid  22918) is running...
#
```

- On Linux, stop the MySQL server and verify the status with the following command:

```
# service mysqld stop
Stopping mysqld:                                     [ OK ]
# service mysqld status
mysqld is stopped
#
```

# Starting the MySQL Server Automatically

- On Linux, you can set the MySQL server to start automatically on system startup and check the status of automatic startup with the `chkconfig` command:

```
# chkconfig mysqld on
# chkconfig --list mysqld
mysqld           0:off    1:off    2:on 3:on 4:on 5:on
                  6:off
#
```

- The numbers followed by "on" or "off" indicate runlevels that control the modes that the server uses when starting. This example, with runlevels 2, 3, 4, and 5 so on, indicates that the MySQL server is using Multi-User Mode with support for MySQL Workbench.

# Using the MySQL Command-Line Tool

---

- In a production environment, applications access the database to query data and update the database.
- The `mysql` client is a useful command-line tool for some of the following tasks:
  - **User administration:** Setting up users and permissions
  - **Database administration:** Importing and exporting data, setting configuration options
  - **Database creation:** Creating databases and checking their structure
  - **Query validation:** Testing and optimizing queries



# Starting the MySQL Command-Line Tool

- Start the mysql command-line client with the following command:

```
# mysql -u root -p
Enter password: <password>
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 1
Server version: 5.7.12 MySQL Community Server (GPL)

...
Type 'help;' or '\h' for help. Type '\c' to clear the
current input statement.

mysql>
```

# Keyboard Editing

---

- To edit mysql client commands, you can take advantage of the following:
  - Arrow keys
  - Backspace key to delete the character to the left of the cursor
  - The command history preserved during the session
  - Character sequence (\c) to cancel a command being entered
  - Enforced quote closure



# mysql Client Commands

- Some useful mysql client commands include the following:
  - PROMPT: Changes the prompt
  - HELP: Displays a list of MySQL commands
  - EXIT or QUIT: Exits the mysql client

```
mysql> PROMPT (\u@\\h) [\d]>\_
PROMPT set to '(\u@\\h) [\d]>\_'
(root@localhost) [(none)]> PROMPT
Returning to default PROMPT of mysql>
mysql> HELP
<list of all MySQL commands>
mysql> EXIT
Bye
# service mysqld status
mysqld (pid 22918) is running...
```

# tee File Logging

---

- The `tee` file is a text file that logs all `mysql` client statements and their output.
- Everything you see displayed on the screen is appended into a file that you specify.
- You can start logging by entering the `tee` command interactively within the `mysql` client:

```
mysql> tee session_tee_filename.txt
```

- You can stop `tee` file logging with the `notee` command.
- Executing `tee` again re-enables logging and appends the new output to the existing file.
- View the `tee` file contents in your favorite text editor.



# Setting the Initial Password for the root Account

---

- The process for setting the initial password for the MySQL root login is the following:
  - The first time you start `mysqld` after installation, it generates a random password and stores it in `/var/log/mysqld.log`.
  - Copy the random password from that file and use it to log in to the `mysql` client as `root`.
  - Set the password for the `root` account to one that you select.



# Accessing the Initial Random Password for root

- Start the mysqld server with the following command:

```
# service mysqld start
Initializing MySQL database:2015-08-05T01:00:36.565655Z0
Starting mysqld:                                     [ OK ]
#
```

- View the random password with the following command:

```
# head /var/log/mysqld.log
2015-08-05T01:00:38.043563Z 0 [Warning] InnoDB: New log
files created, LSN=45790
2015-08-05T01:00:38.431937Z 0 [Warning] InnoDB: Creating
foreign key constraint system tables.
2015-08-05T01:00:40.581515Z 1 [Warning] A temporary
password is generated for root@localhost: be1(BwQLh;Ia
150805 01:01:01 mysqld_safe Starting mysqld daemon
```

# Logging In and Setting the root Password

- Log in to MySQL with the random password copied from the log:

```
# mysql -u root -p  
Enter password:  
Welcome to the MySQL monitor. Commands end with ; or \g.  
...  
Type 'help;' or '\h' for help. Type '\c' to clear the  
current input statement.
```

Enter the following statement to set the root password:

```
mysql> ALTER USER USER()  
      -> IDENTIFIED BY 'oracle';  
Query OK, 0 rows affected (0.00 sec)  
mysql>
```

# Resetting a Forgotten root Password

---

- If you forget the password for the `root` account, you can follow these steps to reset the forgotten `root` password:
  - Start the server with the following options:
    - **--skip-grant-tables**: Enables anyone to connect without a password
    - **--skip-networking**: Prevents remote clients from connecting
  - Log in to the `mysql` client as `root` with no password.
  - Use the `FLUSH PRIVILEGES` statement to load grant tables.
  - Set the password for the `root` account.
  - Stop the server and restart it normally.
  - Log in to the `mysql` client with the new `root` password.



# Statements to Reset a Forgotten root Password

- Reset a forgotten `root` password with the following statements:

```
# service mysqld start --skip-grant-tables  
--skip-networking  
Starting mysqld: [  
    OK ]  
# mysql -u root  
Welcome to the MySQL monitor. Commands end with ;  
or \g.  
...  
Type 'help;' or '\h' for help. Type '\c' to clear  
the current input statement.  
mysql> FLUSH PRIVILEGES;  
Query OK, 0 rows affected (0.00 sec)  
mysql> ALTER USER USER()  
    -> IDENTIFIED BY 'oracle';  
Query OK, 0 rows affected (0.00 sec)  
mysql> exit  
# service mysqld stop
```

# MySQL Workbench

---

- A unified (GUI) tool that helps database architects, developers, and DBAs manage the MySQL server and data.
  - Use MySQL Workbench for:
    - SQL development
    - Data modeling
    - Server administration
    - Database migration
  - There are three editions of MySQL Workbench:
    - Community Edition (GPL)
    - Standard Edition (Commercial)
    - Enterprise Edition (Commercial)
  - MySQL Workbench Runs on Linux/UNIX, Windows, and Mac OS X.



# MySQL Workbench Capabilities

---

- SQL Development
  - Edit and execute SQL queries and scripts.
  - Create or alter database objects.
  - Edit table data.
- Data Modeling
  - Extended entity-relationship (EER) modeling
  - Edit and execute SQL queries and scripts.
  - Design, generate, and manage databases.
- Server Administration
  - Start and stop the server.
  - Edit database server configuration.
  - Manage users.
  - Import and export data.



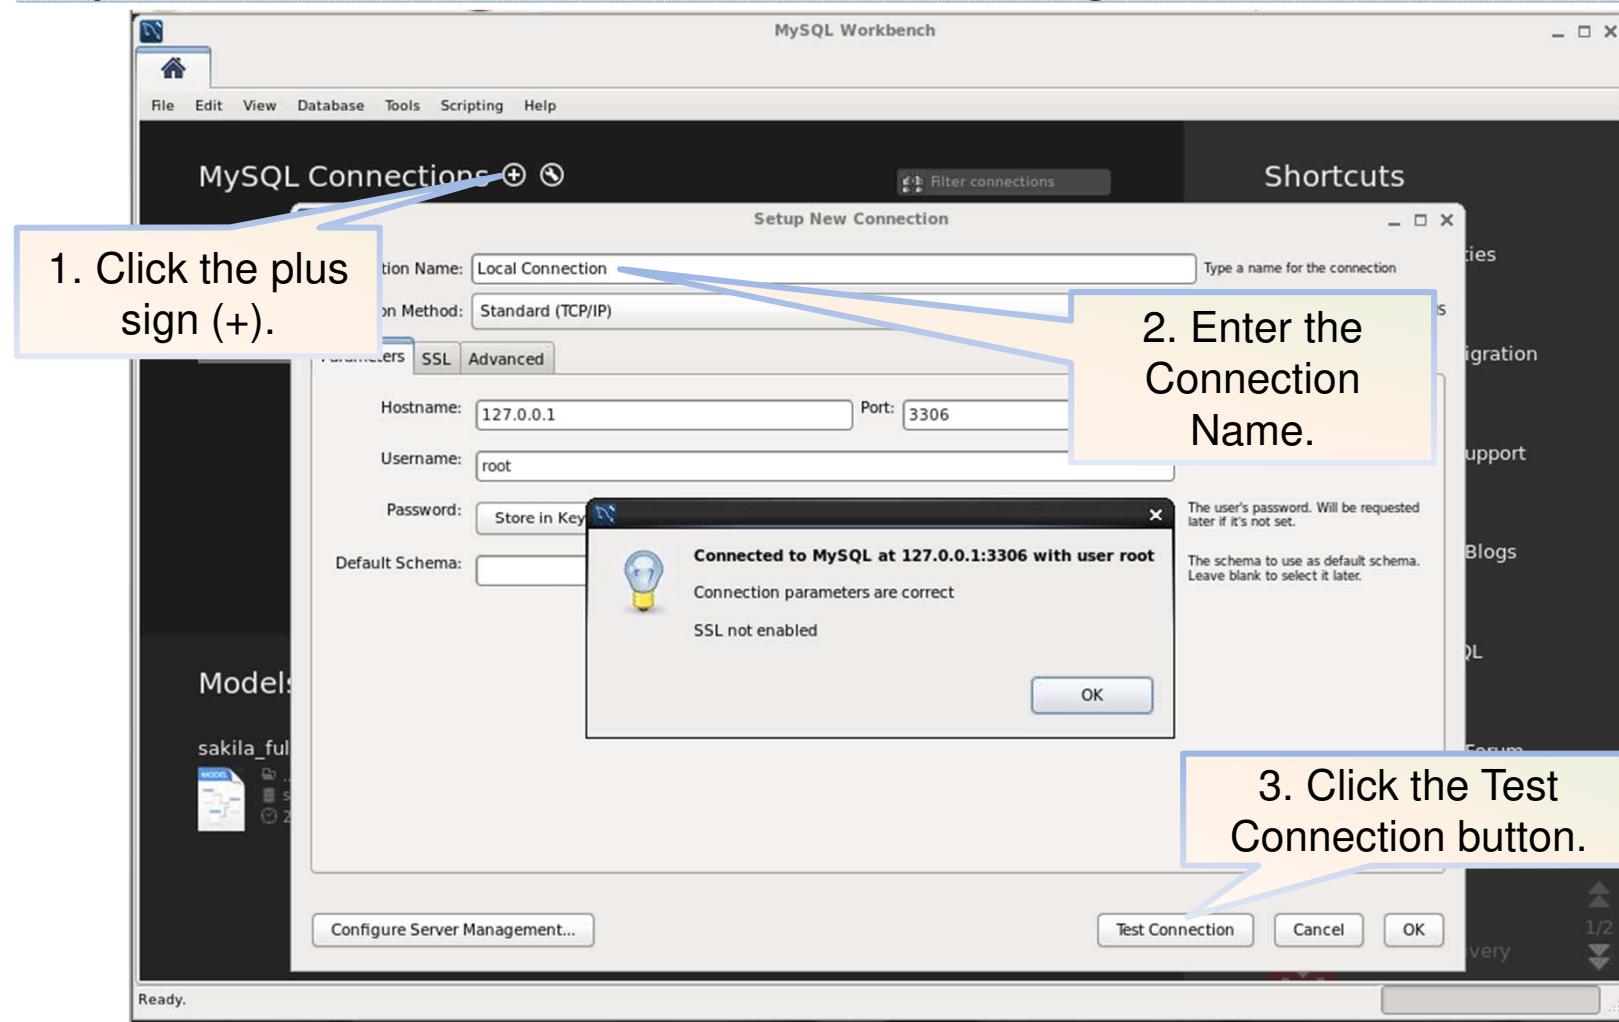
# MySQL Workbench Additional Features and Utilities

---

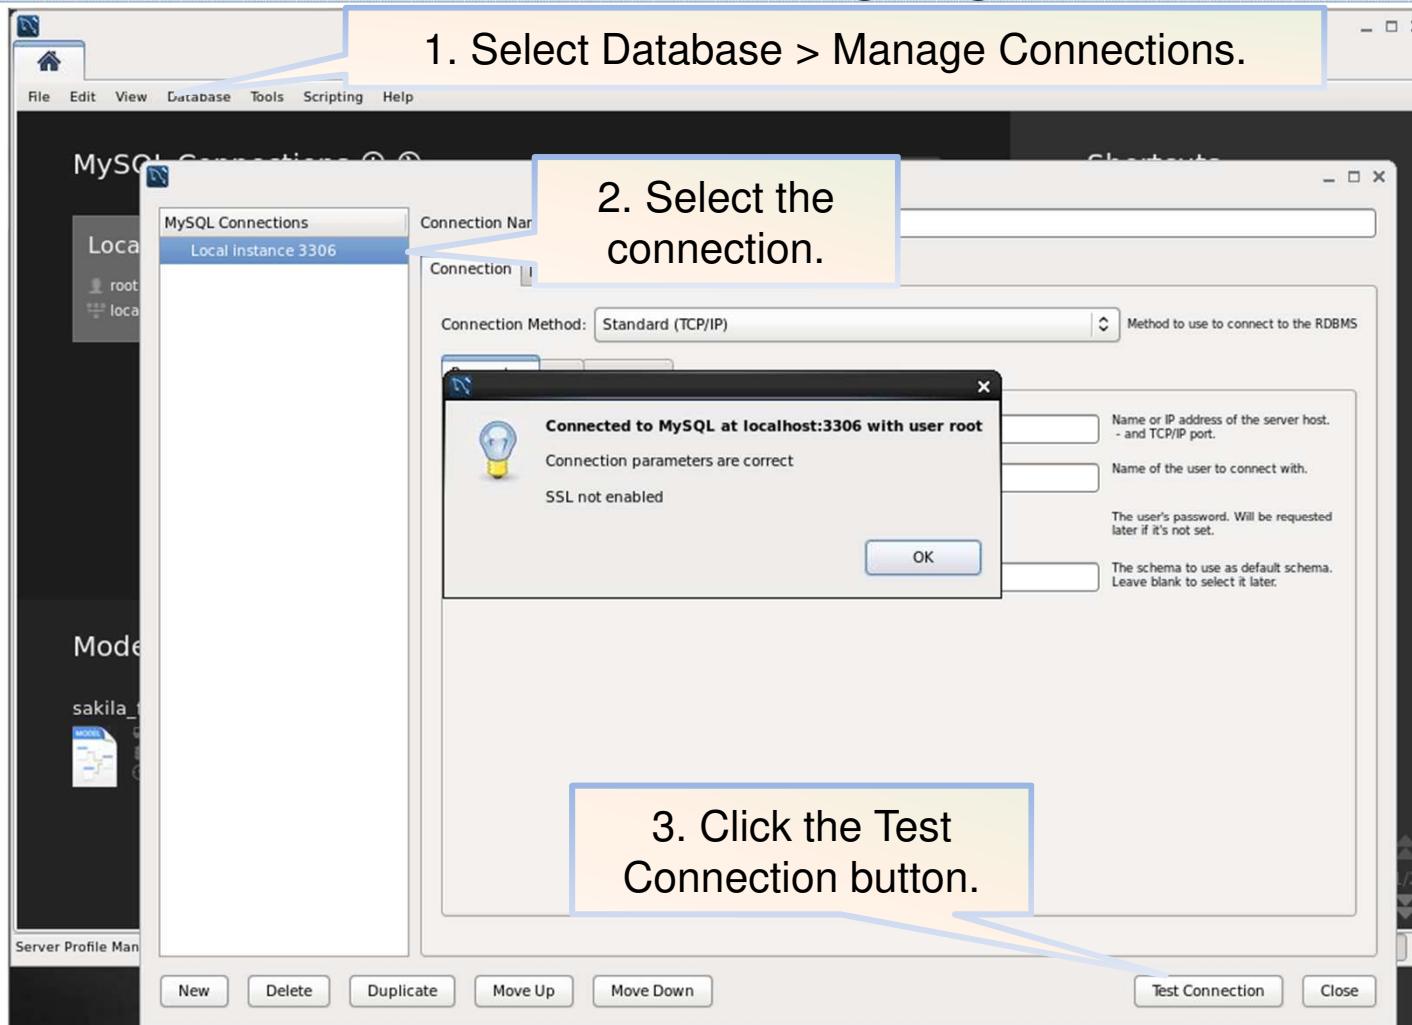
- Database Migration
  - Migrate tables and objects from other RDBMSs to MySQL.
  - Convert existing applications to run on MySQL on Windows and other platforms.
  - Upgrade to later versions of MySQL.
- Access to MySQL Utilities
  - MySQL Utilities are Python tools for working with MySQL Server.
  - MySQL Workbench provides a GUI for working with these tools.



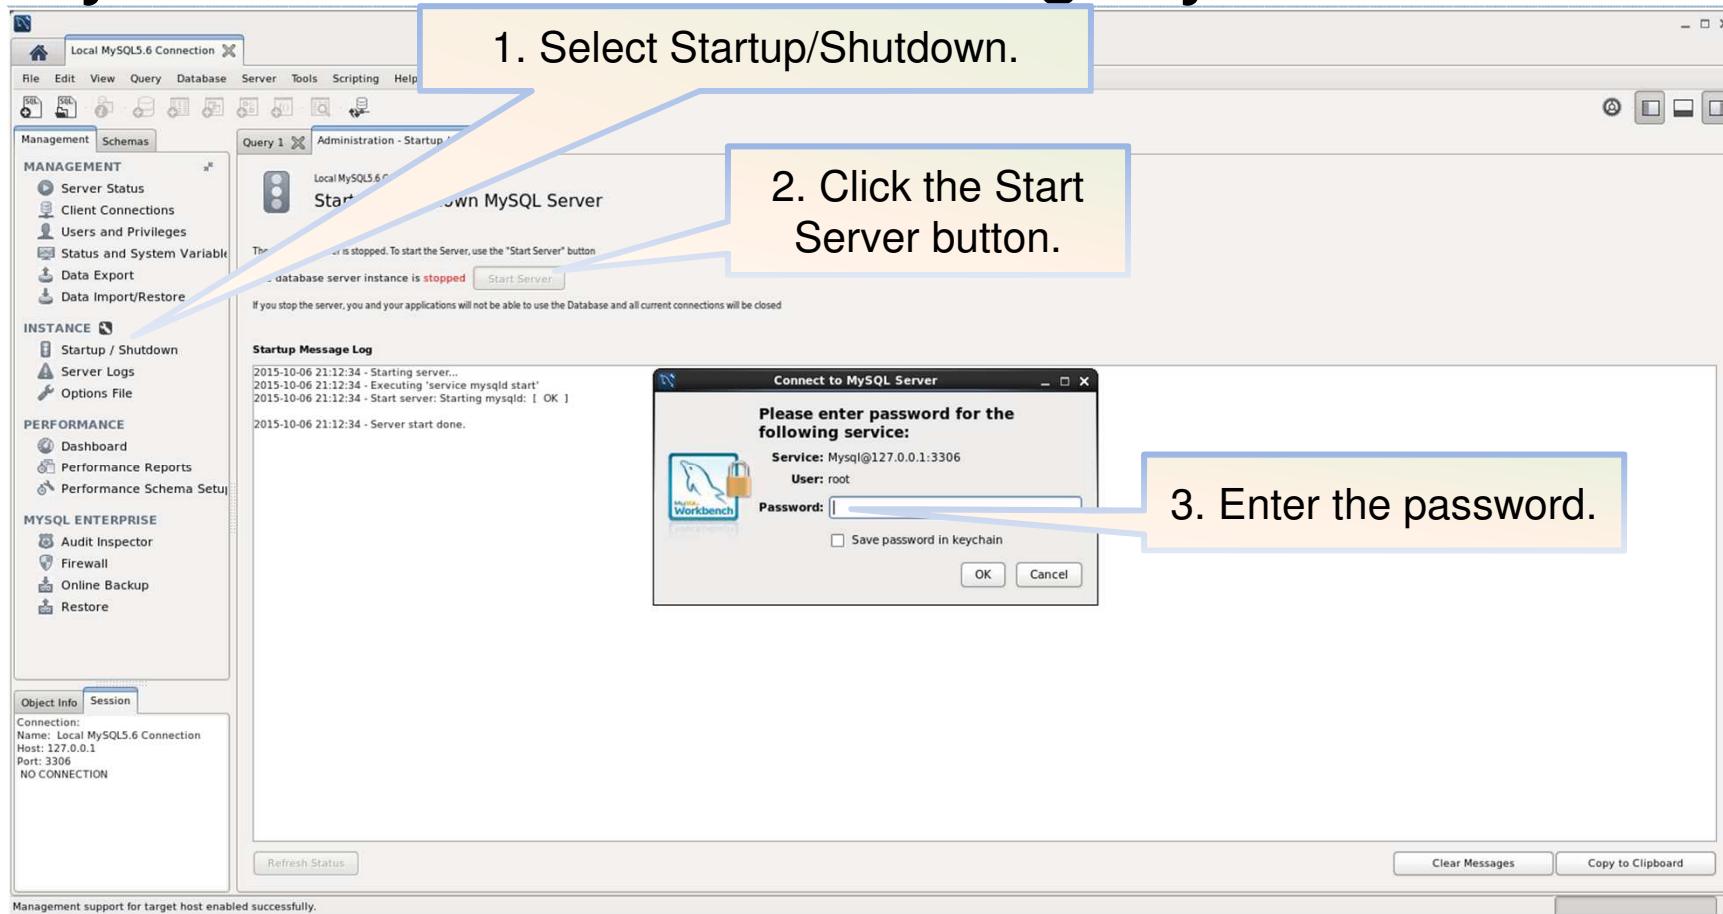
# MySQL Workbench: Creating a Connection



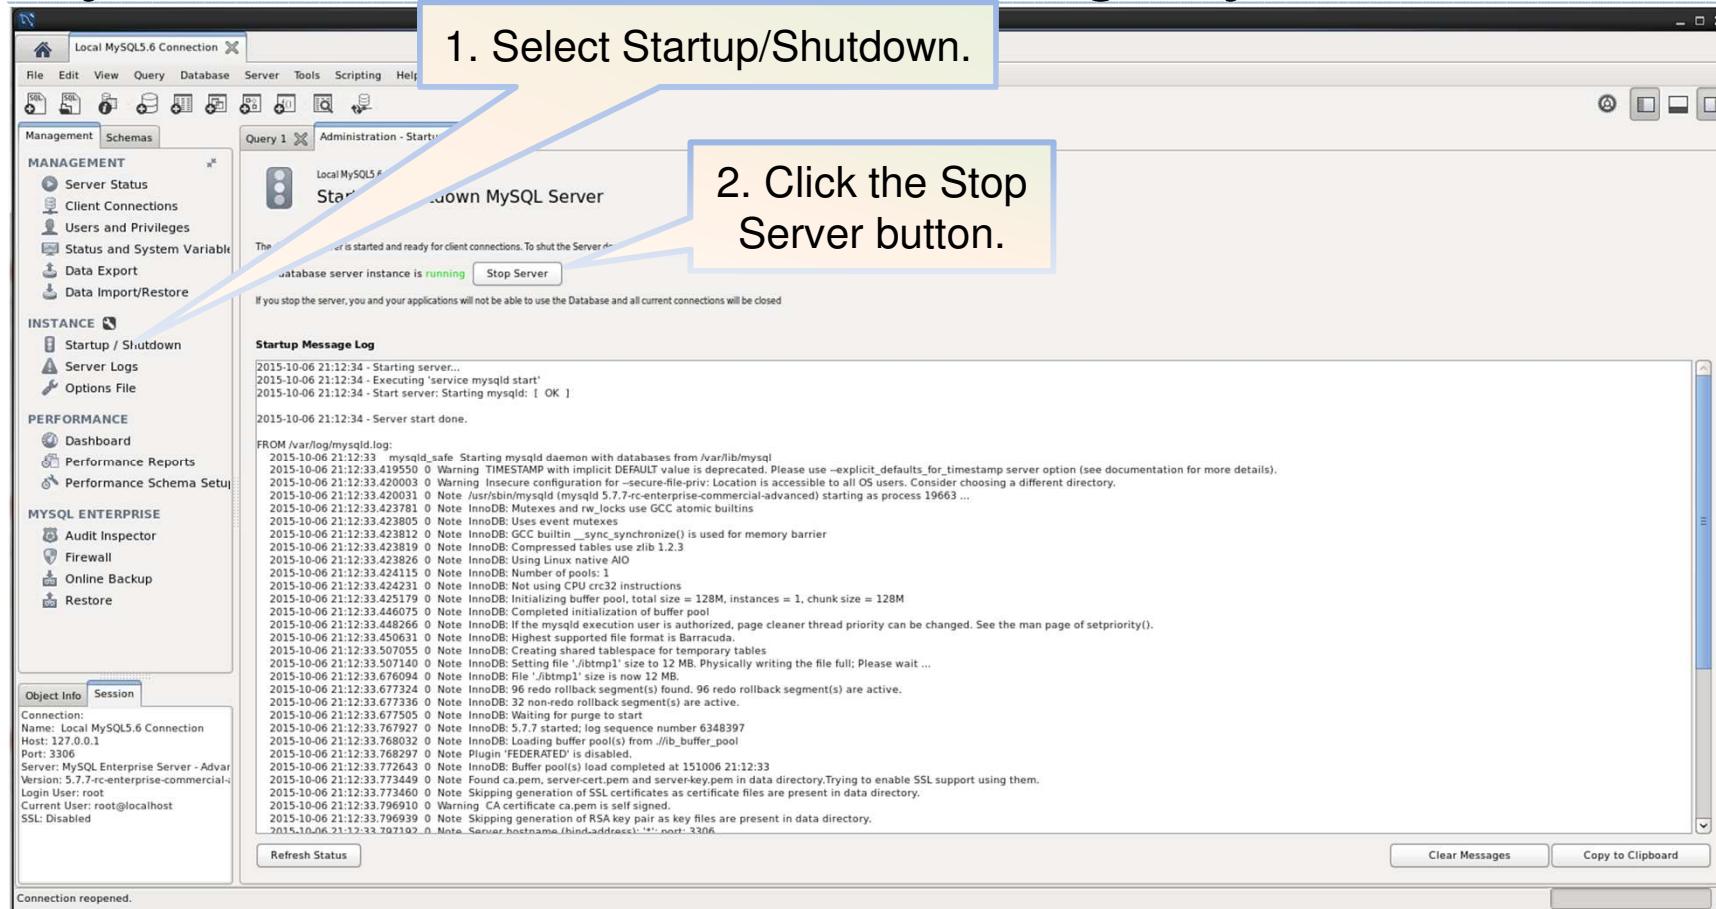
# MySQL Workbench: Managing Connections



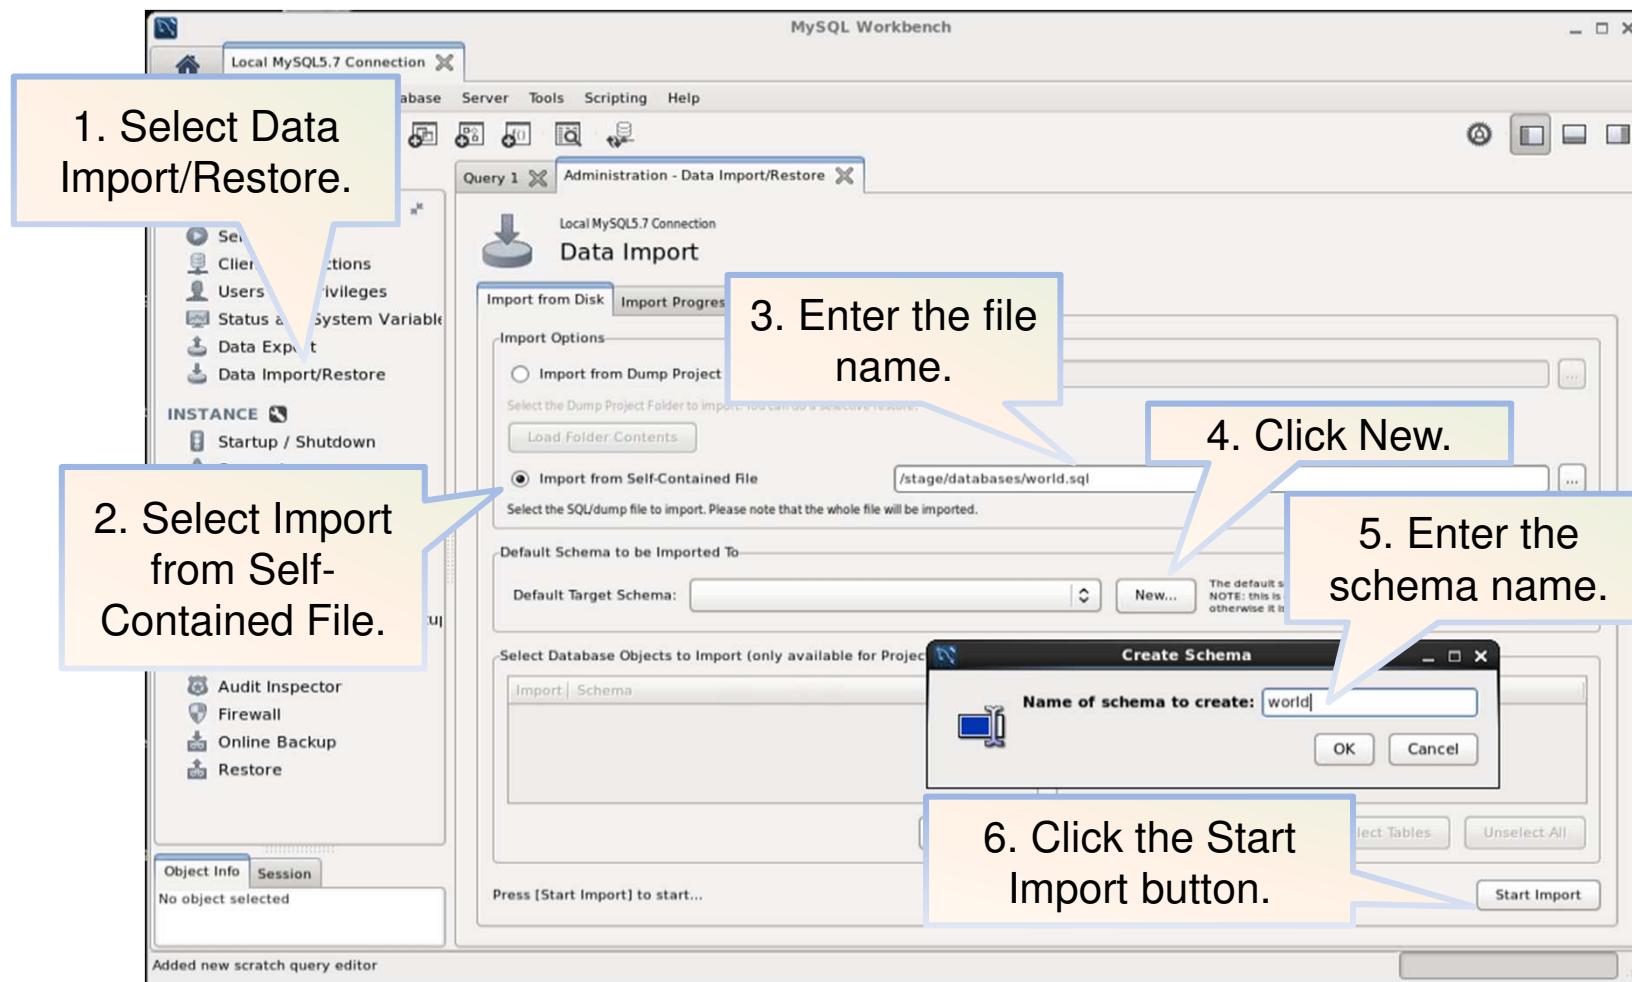
# MySQL Workbench: Starting MySQL Server



# MySQL Workbench: Stopping MySQL Server



# MySQL Workbench: Importing Data



# Querying Table Data

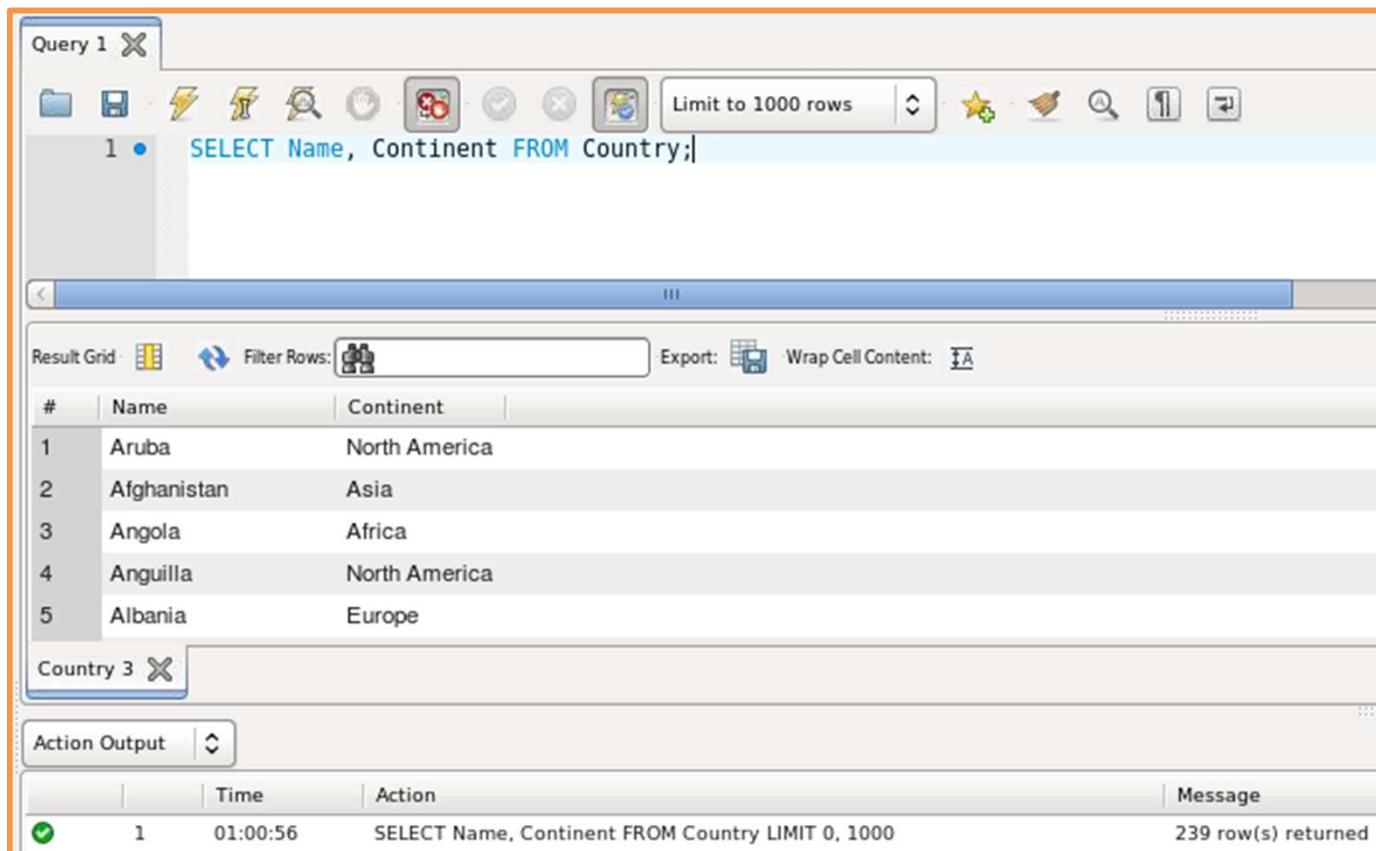
# SELECT Statement

- SELECT is the most commonly used DML command for queries:
  - Retrieves rows from tables in a database
  - Returns rows as a “result set” in the form of a table
- General syntax:

```
SELECT [clause options] column_list  
[FROM table_name] [other clauses]
```

- *column\_list* is a list of column names that make up the result set.
  - Separate the items in the list with a comma separator ( , ).
  - Use an \* to include all columns rather than list them separately.

# SELECT Statement Example



The screenshot shows the MySQL Workbench interface with a query editor and results grid.

**Query Editor:**

```
Query 1
1 •  SELECT Name, Continent FROM Country;
```

**Result Grid:**

#	Name	Continent
1	Aruba	North America
2	Afghanistan	Asia
3	Angola	Africa
4	Anguilla	North America
5	Albania	Europe

**Action Output:**

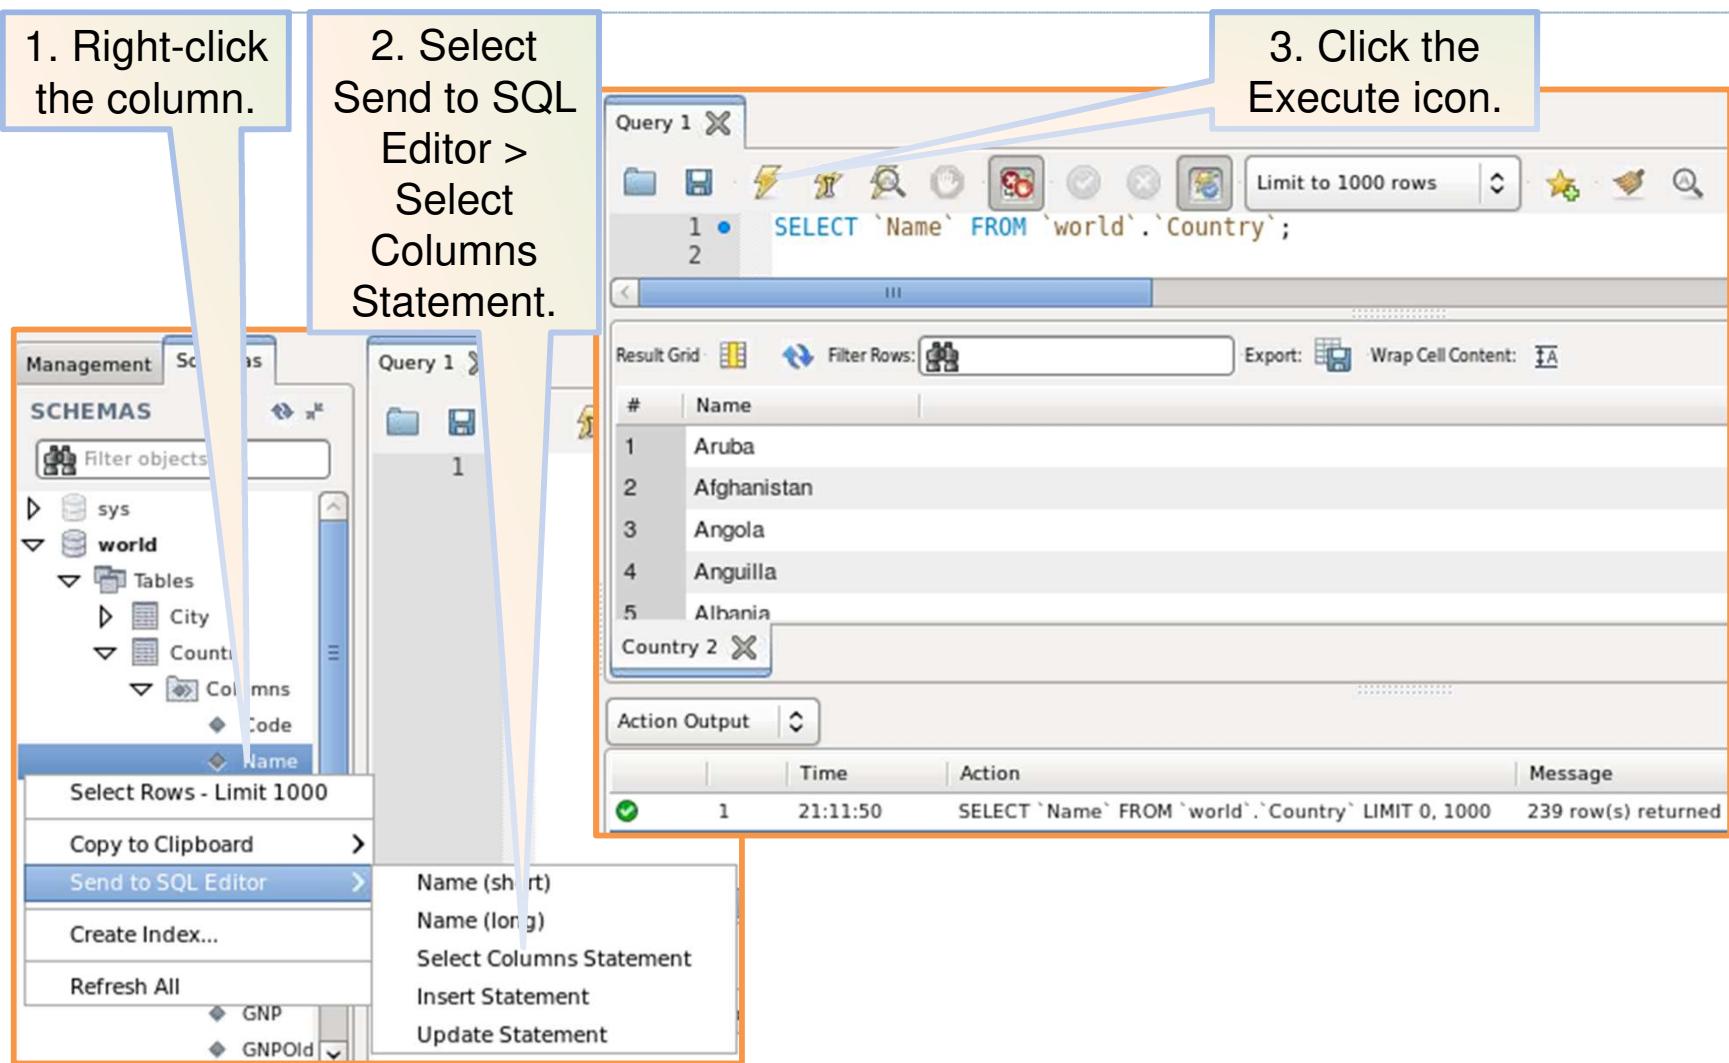
	Time	Action	Message
1	01:00:56	SELECT Name, Continent FROM Country LIMIT 0, 1000	239 row(s) returned

# SELECT Statement Example in MySQL Workbench

1. Right-click the column.

2. Select Send to SQL Editor > Select Columns Statement.

3. Click the Execute icon.



# SELECT Statement Example with \*

The screenshot shows the MySQL Workbench interface with a query editor and a results grid.

**Query Editor:**

```
Query 1 X
1 •  SELECT * FROM CountryLanguage;
```

**Result Grid:**

#	CountryCode	Language	IsOfficial	Percentage
1	ABW	Dutch	T	5.3
2	ABW	English	F	9.5
3	ABW	Papiamento	F	76.7
4	ABW	Spanish	F	7.4
5	AFG	Balochi	F	0.9
6	AFG	Dari	T	32.1

**Action Output:**

	Time	Action	Message
1	02:45:44	SELECT * FROM CountryLanguage LIMIT 0, 1000	984 row(s) returned

# SELECT All Statement in MySQL Workbench

The image shows two screenshots of MySQL Workbench. The left screenshot displays the 'Object Navigator' pane with the 'CountryLanguage' table selected. A context menu is open over the table, with the 'Send to SQL Editor' option highlighted. The right screenshot shows the 'Query Editor' pane with a SELECT statement for the 'CountryLanguage' table. The 'Execute' icon (a red circle with a white play button) is highlighted with a blue box. The 'Result Grid' shows three rows of data from the query.

1. Right-click the table.

2. Select Send to SQL Editor > Select All Statement.

3. Click the Execute icon.

Query 1

```
1 • SELECT `CountryLanguage`.`CountryCode`,  
       `CountryLanguage`.`Language`,  
       `CountryLanguage`.`IsOfficial`,  
       `CountryLanguage`.`Percentage`  
  FROM `world`.`CountryLanguage`;
```

Result Grid

#	CountryCode	Language	IsOfficial	Percentage
1	ABW	Dutch	T	5.3
2	ABW	English	F	9.5
3	ABW	Papiamento	F	76.7

Action Output

	Time	Action	Message
1	21:36:36	SELECT `CountryLanguage` ...	984 row(s) returned

# Using SELECT Clauses

- Use optional clauses alone (or in combination) to generate specific query results.
  - Types of clauses:
    - **DISTINCT**: Eliminates duplicate rows from the result set
    - **FROM**: Specifies which tables to retrieve data from
    - **WHERE**: Filters rows according to specified criteria
    - **ORDER BY**: Sorts rows in a specified order
    - **LIMIT**: Limits the maximum number of rows in the result set
  - The clauses must be used in the correct order:

```
SELECT DISTINCT values_to_return  
    FROM table_name  
    WHERE condition  
    ORDER BY how_to_sort  
    LIMIT row_count;
```

# Using SELECT with DISTINCT

- Removes duplicate rows, so every result set row is unique
- Difference between **SELECT** with and without **DISTINCT**:



Query 1

```
1 • SELECT |Continent
           FROM Country;
```

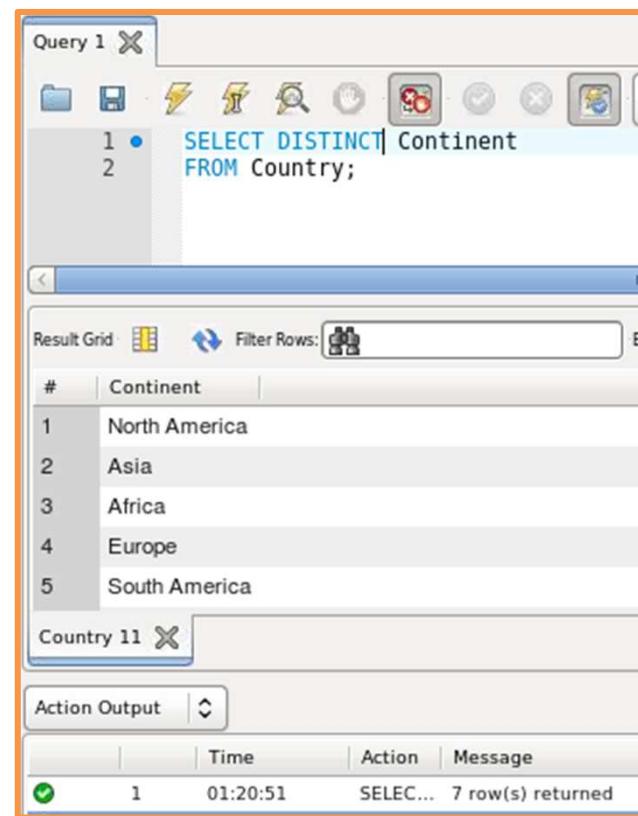
Result Grid

#	Continent
1	North America
2	Asia
3	Africa
4	North America
5	Europe

Action Output

	Time	Action	Message
1	01:19:10	SELEC...	239 row(s) returned

Compared to



Query 1

```
1 • SELECT DISTINCT Continent
           FROM Country;
```

Result Grid

#	Continent
1	North America
2	Asia
3	Africa
4	Europe
5	South America

Action Output

	Time	Action	Message
1	01:20:51	SELEC...	7 row(s) returned

# Using SELECT with DISTINCT for Multiple Columns

Example of the difference between SELECT without and with DISTINCT for multiple columns:

The screenshot shows a database query interface with two panes. The top pane, titled 'Query 1', contains the SQL code:

```
1 • SELECT Continent, Region  
2   FROM Country;
```

The bottom pane, titled 'Result Grid', displays the following data:

#	Continent	Region
1	North America	Caribbean
2	Asia	Southern and Central Asia
3	Africa	Central Africa
4	North America	Caribbean
5	Europe	Southern Europe

The bottom right corner of the interface shows an 'Action Output' table with one row:

	Time	Action	Message
1	01:25:27	SELEC...	239 row(s) returned

Compared to

The screenshot shows a database query interface with two panes. The top pane, titled 'Query 1', contains the SQL code:

```
1 • SELECT DISTINCT Continent, Region  
2   FROM Country;
```

The bottom pane, titled 'Result Grid', displays the following data:

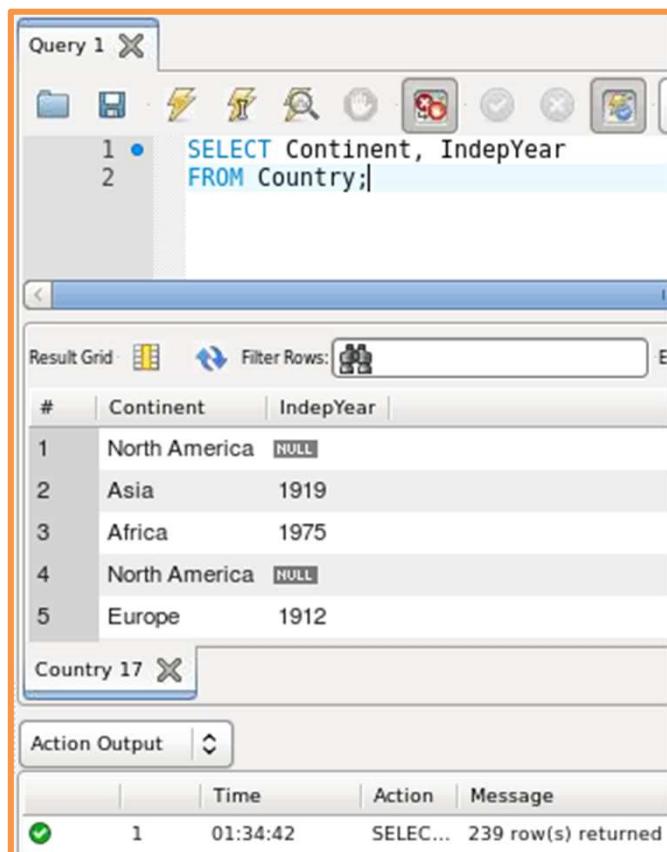
#	Continent	Region
1	North America	Caribbean
2	Asia	Southern and Central Asia
3	Africa	Central Africa
4	Europe	Southern Europe
5	Asia	Middle East

The bottom right corner of the interface shows an 'Action Output' table with one row:

	Time	Action	Message
1	01:27:18	SELEC...	25 row(s) returned

# SELECT DISTINCT with NULL Values

- Treats all NULLs as the same value
- Retrieval of NULL values with and without DISTINCT:



Query 1

```
1 • SELECT Continent, IndepYear  
2   FROM Country;
```

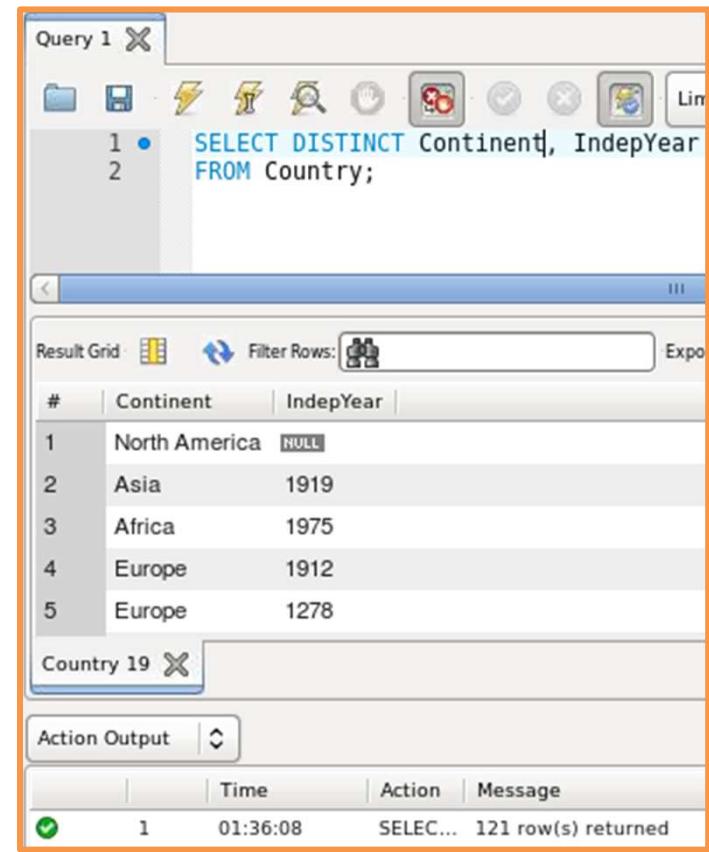
Result Grid

#	Continent	IndepYear
1	North America	NULL
2	Asia	1919
3	Africa	1975
4	North America	NULL
5	Europe	1912

Action Output

	Time	Action	Message
1	01:34:42	SELEC...	239 row(s) returned

Compared to



Query 1

```
1 • SELECT DISTINCT Continent, IndepYear  
2   FROM Country;
```

Result Grid

#	Continent	IndepYear
1	North America	NULL
2	Asia	1919
3	Africa	1975
4	Europe	1912
5	Europe	1278

Action Output

	Time	Action	Message
1	01:36:08	SELEC...	121 row(s) returned

# SELECT with WHERE

---

- Expressions in a **WHERE** clause can use the following types of operators:
  - Arithmetic: +, -, \*, /, DIV, %, MOD
  - Comparison: <, <=, =, <=>, <> or !=, >=, >, **BETWEEN**
  - Logical: **AND**, **OR**, **XOR**, **NOT**
  - Additional options: **IN**, **IS NULL**, **LIKE**, **()**



# Examples of Numeric Expressions with WHERE

- Examples of numeric comparisons:

```
•SELECT Name FROM Country WHERE IndepYear = 1919;  
•SELECT Name, Population FROM Country  
    •WHERE Population > 1000000000;  
•SELECT Name, LifeExpectancy FROM Country  
    •WHERE LifeExpectancy < 40;
```

# Examples of Character Comparisons with WHERE

- Examples of character comparisons:

```
•SELECT Name FROM Country  
    •WHERE GovernmentForm = 'Socialistic Republic';  
  
•SELECT Name, Continent FROM Country  
    •WHERE Continent != 'North America';
```

# Examples of Character Comparisons by Using LIKE with WHERE

- Examples of character comparisons by using LIKE:

```
•SELECT Name, GovernmentForm FROM Country  
    •WHERE GovernmentForm LIKE 'Social%';  
  
•SELECT Name FROM Country  
    •WHERE Name LIKE '%stan';  
  
•SELECT Code, Name FROM Country  
    •WHERE Code LIKE 'CO_';
```

# SELECT with WHERE ... IN

Example of **WHERE** with **IN**:

The screenshot shows the MySQL Workbench interface with an orange border around the central workspace.

**Query Editor:** The title bar says "Query 1". The SQL code is:

```
1 •  SELECT ID, Name, District FROM City
2   WHERE Name IN ('New York', 'Rochester',
3     'Syracuse');
4
```

**Result Grid:** The results of the query are displayed in a table:

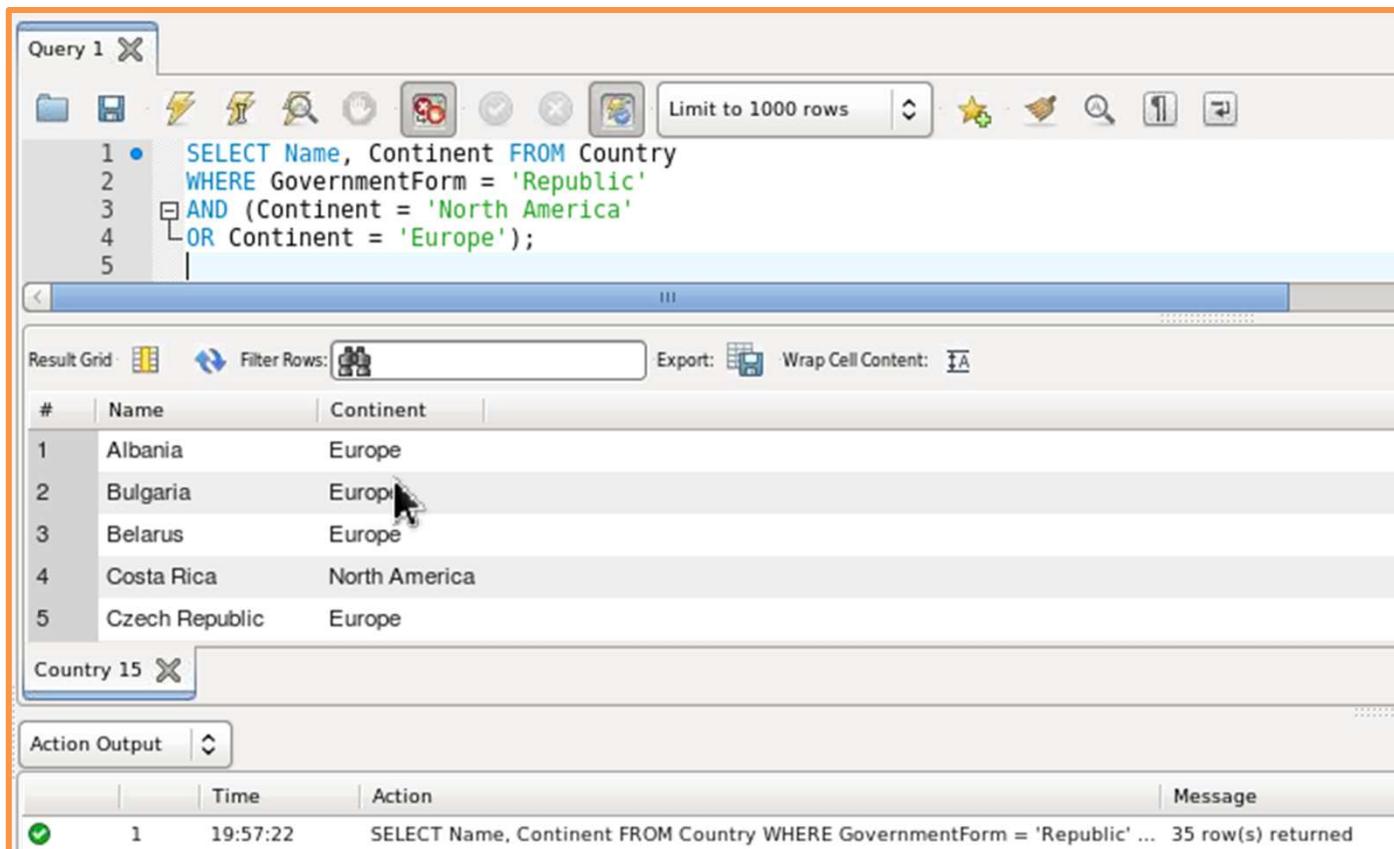
#	ID	Name	District
1	3793	New York	New York
2	3871	Rochester	New York
3	3935	Syracuse	New York
*	NULL	NULL	NULL

**Action Output:** The log shows the executed query and the number of rows returned:

Action	Time	Message
1	19:35:13	SELECT ID, Name, District FROM City WHERE Name IN ('New York', 'Rochester', ... 3 row(s) returned

# SELECT with WHERE ... AND ... OR

Example of **WHERE** with **AND** and **OR**:



The screenshot shows the MySQL Workbench interface with an orange border around the central workspace.

**Query Editor (Top Panel):**

```
Query 1 X
1 •  SELECT Name, Continent FROM Country
2   WHERE GovernmentForm = 'Republic'
3   AND (Continent = 'North America'
4       OR Continent = 'Europe');
5   |
```

**Result Grid (Bottom Panel):**

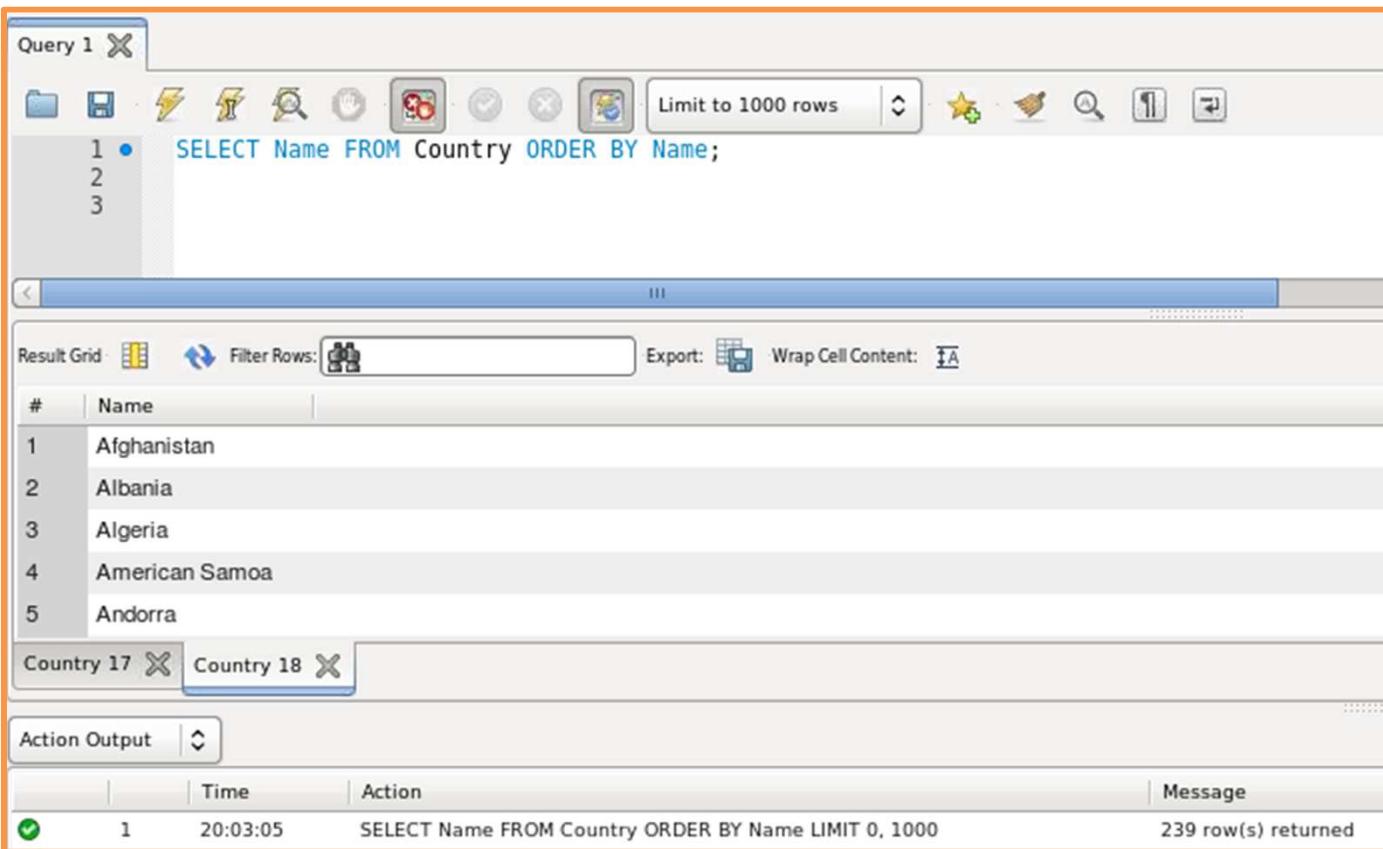
#	Name	Continent
1	Albania	Europe
2	Bulgaria	Europe
3	Belarus	Europe
4	Costa Rica	North America
5	Czech Republic	Europe

**Action Output (Bottom Panel):**

	Time	Action	Message
1	19:57:22	SELECT Name, Continent FROM Country WHERE GovernmentForm = 'Republic' ...	35 row(s) returned

# SELECT with ORDER BY

- Returns output rows in a specific order
- Example of ORDER BY:



The screenshot shows the MySQL Workbench interface with a query editor and a results grid.

**Query Editor (Query 1):**

```
1 •  SELECT Name FROM Country ORDER BY Name;
```

**Result Grid:**

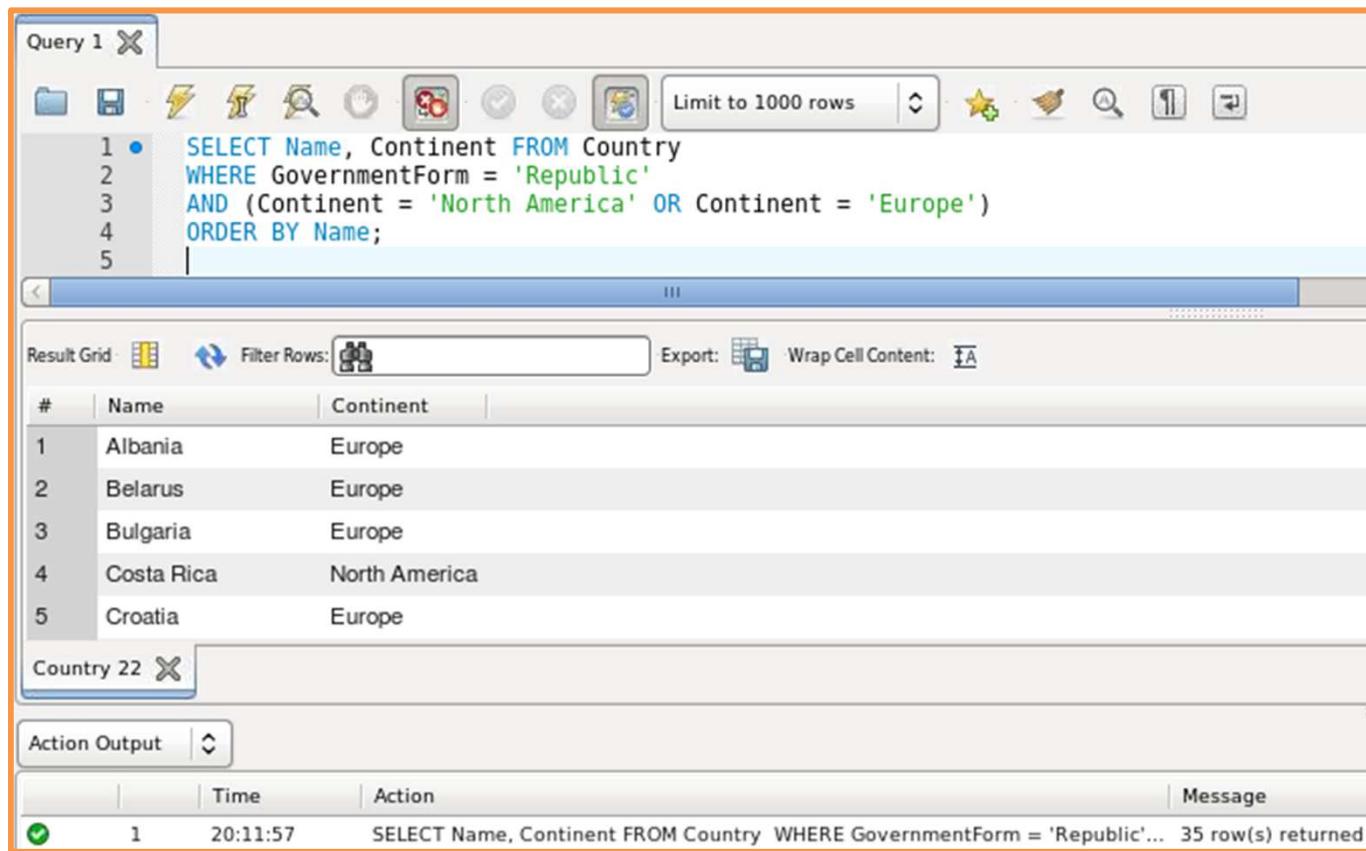
#	Name
1	Afghanistan
2	Albania
3	Algeria
4	American Samoa
5	Andorra

**Action Output:**

	Time	Action	Message
1	20:03:05	SELECT Name FROM Country ORDER BY Name LIMIT 0, 1000	239 row(s) returned

# SELECT with ORDER BY: Another Example

The following example uses ORDER BY with a more complex WHERE clause:



A screenshot of the MySQL Workbench interface. The top window is titled "Query 1" and contains the following SQL code:

```
1 •  SELECT Name, Continent FROM Country
2   WHERE GovernmentForm = 'Republic'
3     AND (Continent = 'North America' OR Continent = 'Europe')
4
5 ORDER BY Name;
```

The "Result Grid" below shows the query results:

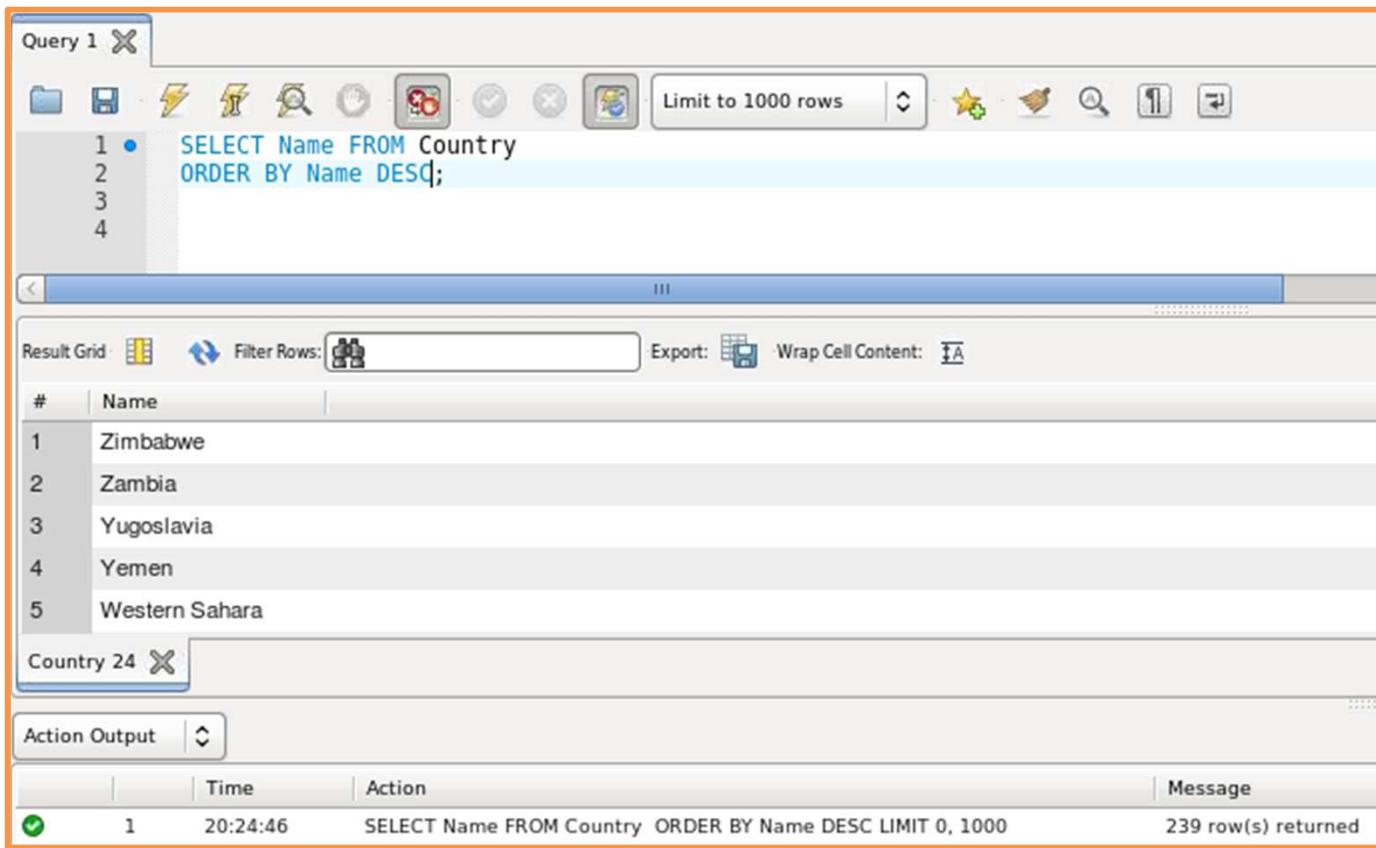
#	Name	Continent
1	Albania	Europe
2	Belarus	Europe
3	Bulgaria	Europe
4	Costa Rica	North America
5	Croatia	Europe

The bottom section shows the "Action Output" log:

	Time	Action	Message
1	20:11:57	SELECT Name, Continent FROM Country WHERE GovernmentForm = 'Republic'...	35 row(s) returned

# SELECT with ORDER BY with ASC and DESC

- **ASC:** Ascending order (default)
- **DESC:** Descending order



The screenshot shows the MySQL Workbench interface with a query editor and a results grid.

**Query Editor (Query 1):**

```
1 • SELECT Name FROM Country
2   ORDER BY Name DESC;
3
4
```

**Result Grid:**

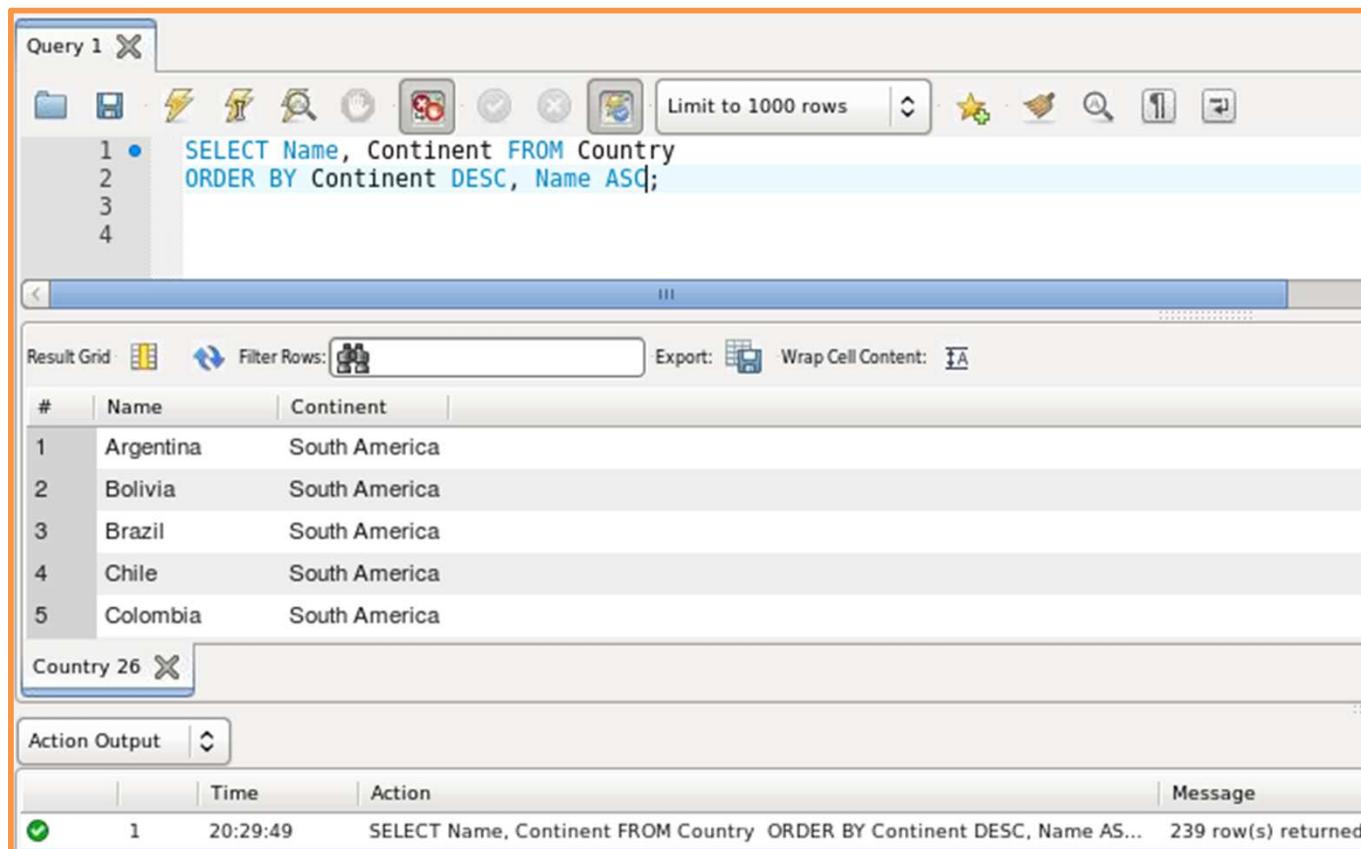
#	Name
1	Zimbabwe
2	Zambia
3	Yugoslavia
4	Yemen
5	Western Sahara

**Action Output:**

	Time	Action	Message
1	20:24:46	SELECT Name FROM Country ORDER BY Name DESC LIMIT 0, 1000	239 row(s) returned

# SELECT WITH ORDER BY WITH ASC AND DESC: Multiple Columns

- You can sort multiple columns simultaneously.
- Example of **ORDER BY** with both **ASC** and **DESC**:



The screenshot shows the MySQL Workbench interface with a query editor and a results grid. The query editor contains the following SQL code:

```
1 •  SELECT Name, Continent FROM Country
2   ORDER BY Continent DESC, Name ASC;
```

The results grid displays the following data:

#	Name	Continent
1	Argentina	South America
2	Bolivia	South America
3	Brazil	South America
4	Chile	South America
5	Colombia	South America

The Action Output pane at the bottom shows the executed query and the number of rows returned:

Action	Time	Message
1	20:29:49	SELECT Name, Continent FROM Country ORDER BY Continent DESC, Name AS... 239 row(s) returned

# SELECT with LIMIT

- Specifies the number of rows to output in a result set
- Example of **LIMIT**:

The screenshot shows the MySQL Workbench interface with a query window titled "Query 1". The query is:

```
1 •  SELECT Name FROM Country LIMIT 4;
```

The results grid displays four rows of data:

#	Name
1	Aruba
2	Afghanistan
3	Angola
4	Anguilla

Below the results grid is another window titled "Country 2" which contains an "Action Output" table:

Action	Time	Message
SELECT Name FROM Country LIMIT 4	01:08:32	4 row(s) returned

# SELECT with LIMIT in MySQL Workbench

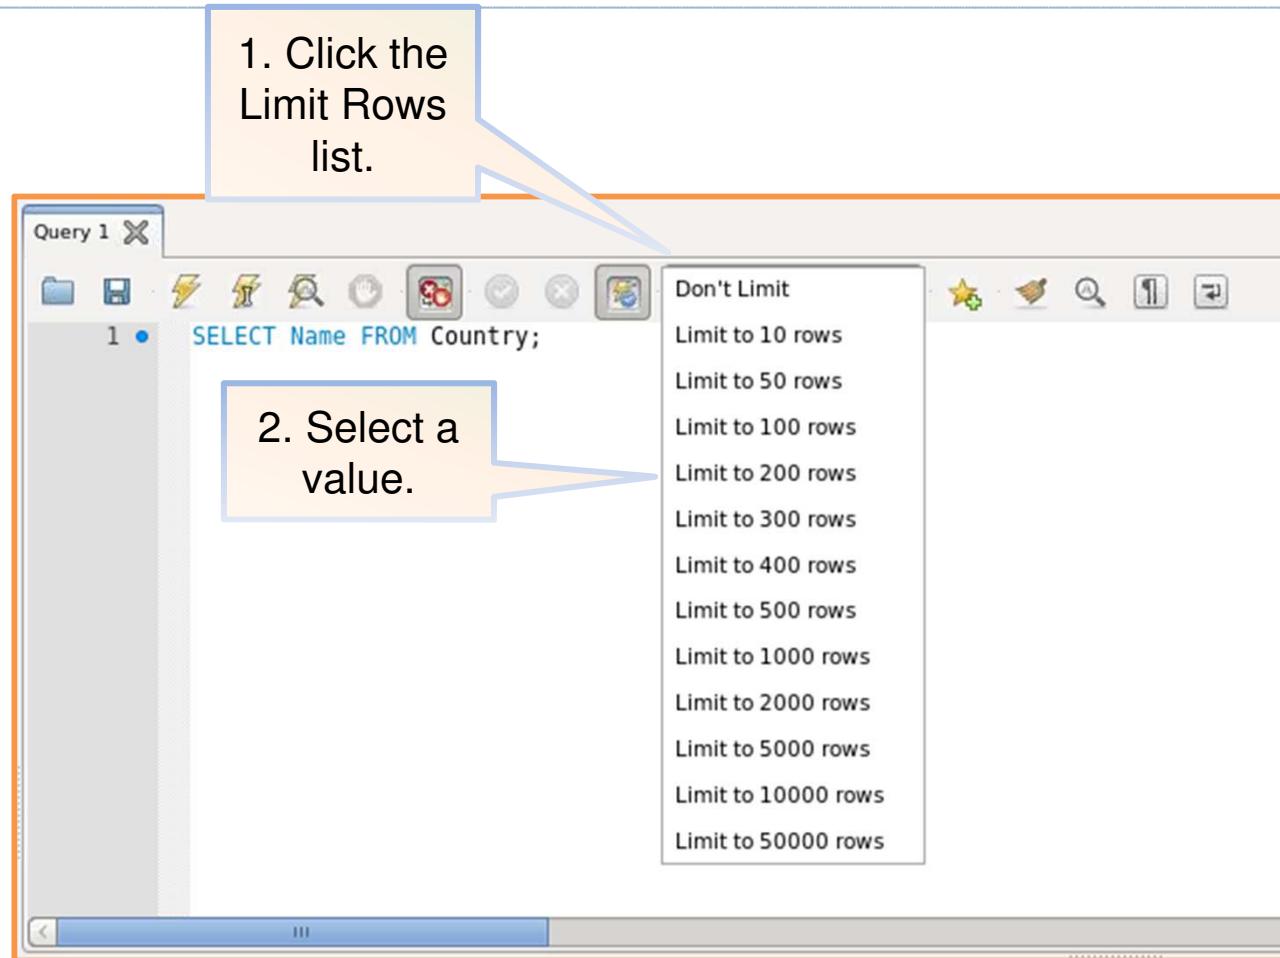
1. Right-click the table.
2. Select Select Rows - Limit 1000.
3. Click the Execute icon.

The screenshot shows the MySQL Workbench interface. On the left, the Object Navigator displays the 'world' schema with its tables. The 'City' table is selected and highlighted with a blue border. A context menu is open over the 'City' table, with the 'Select Rows - Limit 1000' option highlighted with a yellow box and connected by a blue line to the second step. The main workspace contains two tabs: 'Query 1' and 'City'. The 'Query 1' tab shows the SQL command: 'SELECT \* FROM world.City;'. The 'City' tab shows a result grid with six rows of city data from Afghanistan and the Netherlands. Below the grid is an 'Action Output' table with one entry: a green checkmark, the number 1, the time 21:25:04, the action 'SELECT \* FROM world.City LIMIT 0, 1000', and the message '1000 row(s) returned'. The entire right-hand area is enclosed in an orange box.

#	ID	Name	CountryCode	District	Population
1	1	Kabul	AFG	Kabol	3678000
2	2	Qandahar	AFG	Qandahar	491500
3	3	Herat	AFG	Herat	436300
4	4	Mazar-e-Sharif	AFG	Balkh	693000
5	5	Amsterdam	NLD	Noord-Holland	731200
6	6	Rotterdam	NLD	Zuid-Holland	593321

	Time	Action	Message
1	21:25:04	SELECT * FROM world.City LIMIT 0, 1000	1000 row(s) returned

# Setting the Limit for SELECT with LIMIT in MySQL Workbench



# SELECT with LIMIT and Skip Count

- Specifies how many rows to skip before starting the result set
- Example of skipped **LIMIT**:

The screenshot shows the MySQL Workbench interface with a query window titled "Query 1". The query is:

```
1 •  SELECT Name, Population FROM Country LIMIT 20,4;
```

The results grid displays the following data:

#	Name	Population
1	Burkina Faso	11937000
2	Bangladesh	129155000
3	Bulgaria	8190900
4	Bahrain	617000

The "Action Output" section at the bottom shows the query and the message "4 row(s) returned".

# SELECT with LIMIT and ORDER BY

- Returns a specific number of rows in a specific order
- Example:

The screenshot shows the MySQL Workbench interface with two main panes. The top pane is titled "Query 1" and contains the following SQL code:

```
1 •  SELECT Name FROM Country
2   ORDER BY Name LIMIT 4 ;
3
```

The bottom pane is titled "Result Grid" and displays the following table:

#	Name
1	Afghanistan
2	Albania
3	Algeria
4	American Samoa

Below the Result Grid is a smaller window titled "Country 36" which contains the text "Action Output". This window has a table with the following data:

	Time	Action	Message
1	20:46:38	SELECT Name FROM Country ORDER BY Name LIMIT 4	4 row(s) returned

# SELECT Usage Tips

- Treat table and database names as case-sensitive.
- From the mysql command-line client:
  - Use \c to abort a SQL statement
  - Use \G in place of ; to see results vertically with each column in a row, instead of as a table, for example:

```
mysql> SELECT Name, Continent FROM Country\G
***** 1. row *****
      Name: Afghanistan
    Continent: Asia
***** 2. row *****
      Name: Angola
    Continent: Africa
...
```

- Use \* (all columns) with caution; it can waste resources and give unpredictable results due to changes in columns.

# Querying Multiple Tables

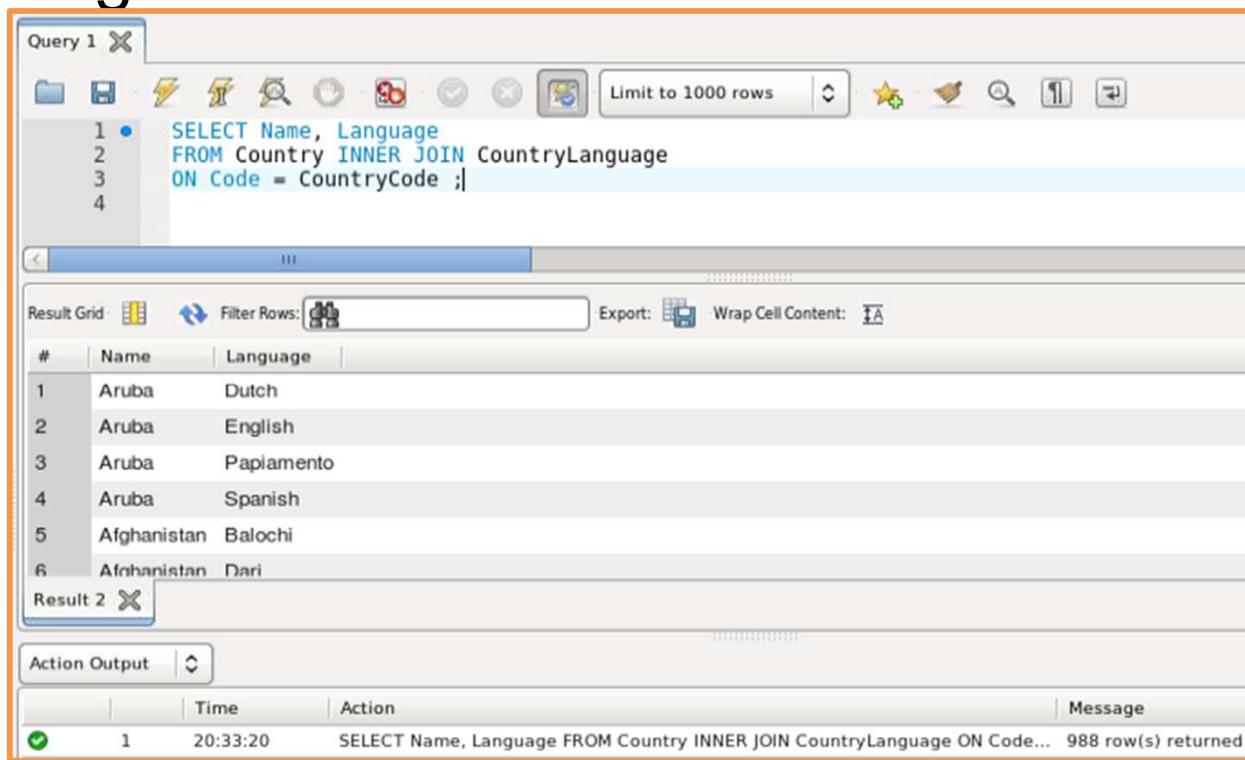
---

- Some queries need data from multiple tables.
- You use a join operation to combine data from different tables.
- A join creates a temporary result set that contains combined data.
- To join tables, there must be a relationship between certain columns in these tables.



# Combining Multiple Tables: INNER JOIN Keyword

- Join tables by including them in the FROM clause of the SELECT statement, separated by the INNER JOIN keyword. Indicate the relationship to use for joining the tables in the ON clause.



The screenshot shows the MySQL Workbench interface with a query editor and a results grid. The query editor at the top contains the following SQL code:

```
Query 1
1 • 1. SELECT Name, Language
2   FROM Country INNER JOIN CountryLanguage
3     ON Code = CountryCode ;|
```

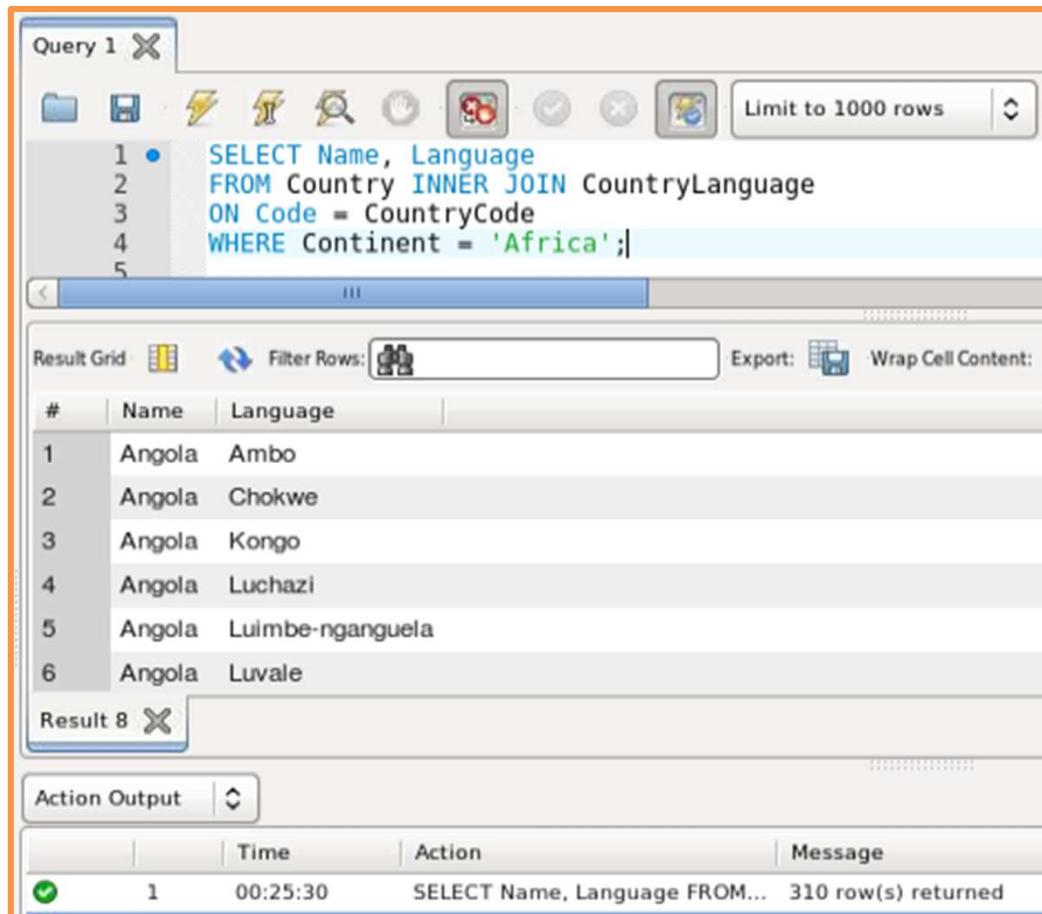
The results grid below displays the following data:

#	Name	Language
1	Aruba	Dutch
2	Aruba	English
3	Aruba	Papiamento
4	Aruba	Spanish
5	Afghanistan	Balochi
6	Afghanistan	Dari

The status bar at the bottom shows the message: "SELECT Name, Language FROM Country INNER JOIN CountryLanguage ON Code... 988 row(s) returned".

# Limiting the Results of a Join by Using a WHERE Clause

- Include a WHERE clause to limit the results of an INNER JOIN.



The screenshot shows the MySQL Workbench interface with a query window titled "Query 1". The query is:

```
1 • SELECT Name, Language
2   FROM Country INNER JOIN CountryLanguage
3     ON Code = CountryCode
4   WHERE Continent = 'Africa';
```

The result grid displays the following data:

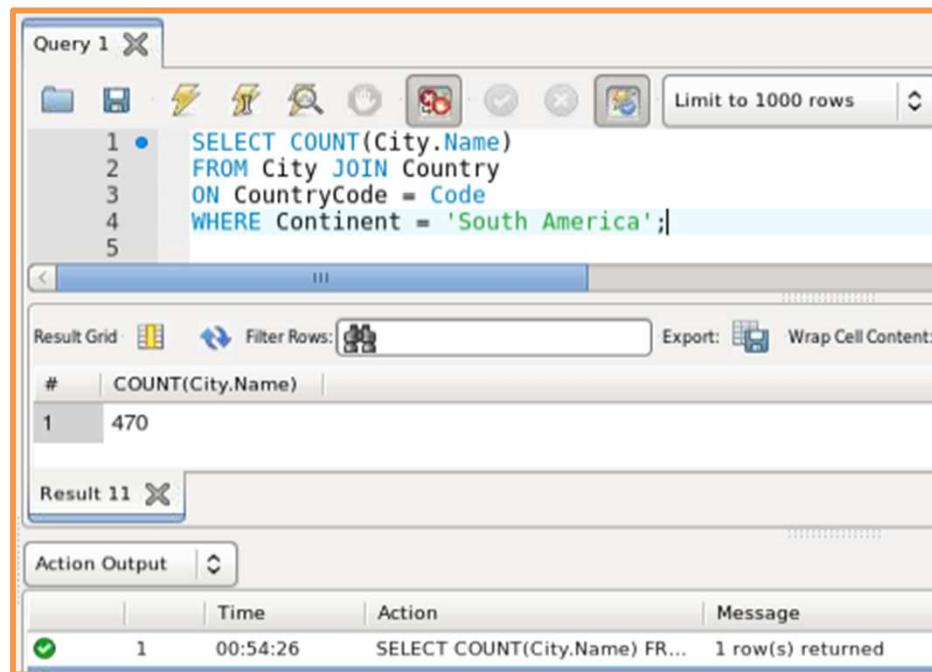
#	Name	Language
1	Angola	Ambo
2	Angola	Chokwe
3	Angola	Kongo
4	Angola	Luchazi
5	Angola	Luimbe-nganguela
6	Angola	Luvale

The action output log at the bottom shows:

Action	Time	Message
SELECT Name, Language FROM...	00:25:30	310 row(s) returned

# JOIN Keyword

- The JOIN keyword is equivalent to INNER JOIN. Use an ON clause to indicate the relationship for the join and a WHERE clause to limit results.



The screenshot shows the MySQL Workbench interface with a query editor and a results grid. The query editor contains the following SQL code:

```
Query 1 X
1 • SELECT COUNT(City.Name)
2   FROM City JOIN Country
3     ON CountryCode = Code
4     WHERE Continent = 'South America';
5
```

The results grid displays the output of the query:

#	COUNT(City.Name)
1	470

The status bar at the bottom shows the message: "SELECT COUNT(City.Name) FR... 1 row(s) returned".

# Joins: Columns That Are Not Related

- Columns from different tables can be matched for a join even if the columns are not related.

The screenshot shows the MySQL Workbench interface with a query editor and a results grid. The query editor contains the following SQL code:

```
1 • SELECT *
 2   FROM City JOIN CountryLanguage
 3     ON City.CountryCode = CountryLanguage.CountryCode
 4   WHERE City.Name IN ('Paris', 'Istanbul', 'Moscow', 'New York')
 5     AND IsOfficial = 'T';
 6
```

The results grid displays four rows of data:

#	ID	Name	CountryCode	District	Population	CountryCode	Language	IsOfficial	Percentage
1	2974	Paris	FRA	Île-de-France	2125246	FRA	French	T	93.6
2	3357	Istanbul	TUR	Istanbul	8787958	TUR	Turkish	T	87.6
3	3580	Moscow	RUS	Moscow (City)	8389200	RUS	Russian	T	86.6
4	3793	New York	USA	New York	8008278	USA	English	T	86.2

The action output shows the query executed and the number of rows returned:

Action Output	Time	Action	Message
1	21:01:00	SELECT * FROM City JOIN CountryLanguage ON City.CountryCode = CountryLa...	4 row(s) returned

# USING Keyword

- If the columns in the ON clause have the same names in both tables, you can use the USING keyword.

The screenshot shows the MySQL Workbench interface with a query editor and results grid. The query is:

```
1 • SELECT *
  2 FROM City JOIN CountryLanguage
  3 USING (CountryCode)
  4 WHERE City.Name IN ('Paris', 'Istanbul', 'Moscow', 'New York')
  5 AND IsOfficial = 'T';
  6
```

The results grid displays the following data:

#	CountryCode	ID	Name	District	Population	Language	IsOfficial	Percentage
1	FRA	2974	Paris	Île-de-France	2125246	French	T	93.6
2	TUR	3357	Istanbul	Istanbul	8787958	Turkish	T	87.6
3	RUS	3580	Moscow	Moscow (City)	8389200	Russian	T	86.6
4	USA	3793	New York	New York	8008278	English	T	86.2

The Action Output pane shows the query executed at 22:03:38, returning 4 rows.

# NATURAL JOIN Keyword

- The NATURAL JOIN keyword can replace a USING clause that lists all the column names in common between the tables.

The screenshot shows the MySQL Workbench interface with a query editor and a results grid. The query editor contains the following SQL code:

```
Query 1
SELECT *
FROM City NATURAL JOIN CountryLanguage
WHERE City.Name IN ('Paris', 'Istanbul', 'Moscow', 'New York')
AND IsOfficial = 'T';
```

The results grid displays the following data:

#	CountryCode	ID	Name	District	Population	Language	IsOfficial	Percentage
1	FRA	2974	Paris	Île-de-France	2125246	French	T	93.6
2	TUR	3357	Istanbul	Istanbul	8787958	Turkish	T	87.6
3	RUS	3580	Moscow	Moscow (City)	8389200	Russian	T	86.6
4	USA	3793	New York	New York	8008278	English	T	86.2

The results grid has a header bar with "Result Grid", "Filter Rows:", "Export:", and "Wrap Cell Content:" buttons. Below the grid, there is an "Action Output" section with a table showing the execution details:

Action	Time	Message
1	22:12:08	SELECT * FROM City NATURAL JOIN CountryLanguage WHERE City.Name IN ('P... 4 row(s) returned

# Outer Joins

---

- Outer joins include LEFT OUTER JOIN and RIGHT OUTER JOIN, or LEFT JOIN and RIGHT JOIN. The results of an outer join include all the rows from one table that meet the selection criteria in the WHERE clause, but only those rows from the other table that match with an ON or USING clause.
    - LEFT JOIN returns all rows from the left table but only matching rows from the right table.
    - RIGHT JOIN returns all rows from the right table but only matching rows from the left table.
    - The left table is listed first or to the left of the LEFT JOIN or RIGHT JOIN keyword. The right table is listed after or to the right of the keyword.
-

# LEFT JOIN Example

The screenshot shows the MySQL Workbench interface with a query editor and a results grid.

**Query Editor:**

```
1 • SELECT Country.Name, Country.Population, City.Name, CountryCode
2   FROM Country LEFT JOIN City
3     ON Code = CountryCode
4   WHERE Country.Population < 1000
5   ORDER BY Country.Population DESC;
6
```

**Result Grid:**

#	Name	Population	Name	CountryCode
1	Cocos (Keeling) Islands	600	Bantam	CCK
2	Cocos (Keeling) Islands	600	West Island	CCK
3	Pitcairn	50	Adamstown	PCN
4	Antarctica	0	NULL	NULL
5	French Southern territories	0	NULL	NULL

**Action Output:**

Action	Time	Action	Message
1	23:34:30	SELECT Country.Name, Country.Population, City.Name, CountryCode FROM Co...	10 row(s) returned

# RIGHT JOIN Example

The screenshot shows the MySQL Workbench interface with a query editor and results grid.

**Query Editor:**

```
1 • SELECT Country.Name, Country.Population, City.Name, CountryCode
2   FROM Country RIGHT JOIN City
3     ON Code = CountryCode
4   WHERE Country.Population < 1000
5   ORDER BY Population DESC;
```

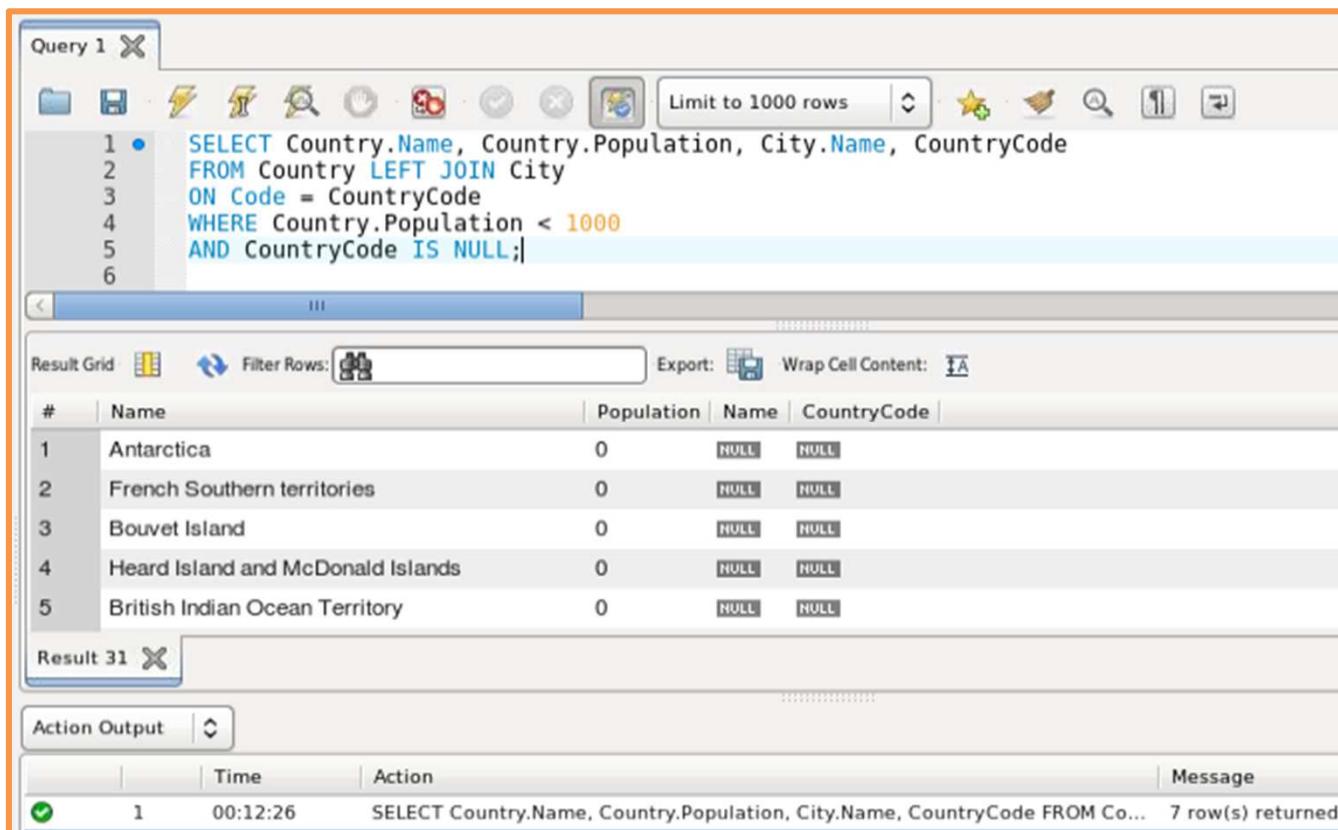
**Result Grid:**

#	Name	Population	Name	CountryCode
1	Cocos (Keeling) Islands	600	Bantam	CCK
2	Cocos (Keeling) Islands	600	West Island	CCK
3	Pitcairn	50	Adamstown	PCN

**Action Output:**

	Time	Action	Message
1	01:11:01	SELECT Country.Name, Country.Population, City.Name, CountryCode FROM Co...	3 row(s) returned

# Using OUTER JOIN with IS NULL



The screenshot shows a MySQL Workbench interface with an orange border around the main content area. At the top is a toolbar with various icons. Below it is a query editor window titled "Query 1" containing the following SQL code:

```
1 • SELECT Country.Name, Country.Population, City.Name, CountryCode
2   FROM Country LEFT JOIN City
3     ON Code = CountryCode
4   WHERE Country.Population < 1000
5   AND CountryCode IS NULL;
```

Below the query editor is a "Result Grid" table with the following data:

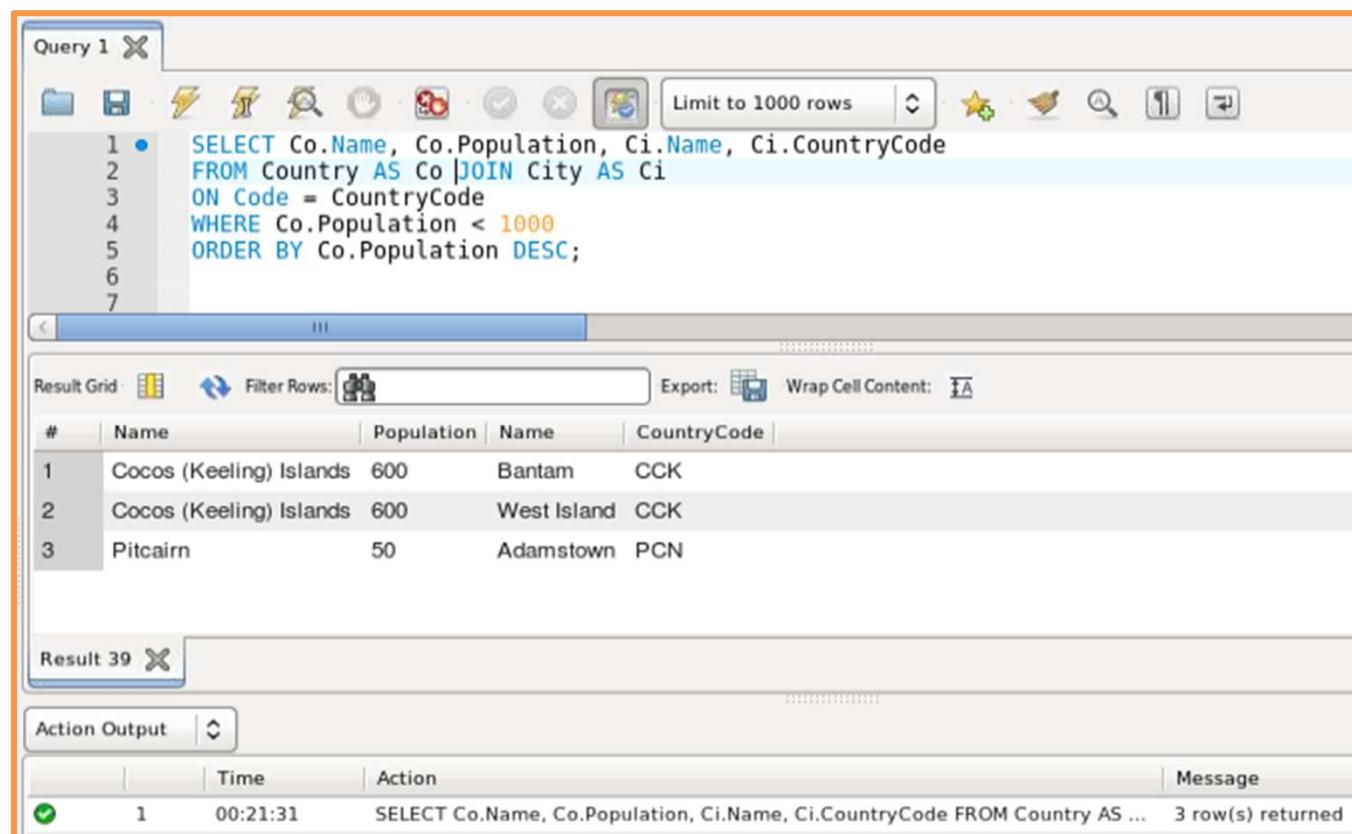
#	Name	Population	Name	CountryCode
1	Antarctica	0	NULL	NULL
2	French Southern territories	0	NULL	NULL
3	Bouvet Island	0	NULL	NULL
4	Heard Island and McDonald Islands	0	NULL	NULL
5	British Indian Ocean Territory	0	NULL	NULL

At the bottom of the interface is an "Action Output" section with a table:

Action	Time	Message
1	00:12:26	SELECT Country.Name, Country.Population, City.Name, CountryCode FROM Co... 7 row(s) returned

# Table Name Aliases

- Table names can be aliased by using the AS keyword.



The screenshot shows the MySQL Workbench interface with an orange border around the central workspace. At the top is the Query Editor titled "Query 1" containing the following SQL code:

```
1 •  SELECT Co.Name, Co.Population, Ci.Name, Ci.CountryCode
2   FROM Country AS Co JOIN City AS Ci
3     ON Co.Code = Ci.CountryCode
4   WHERE Co.Population < 1000
5   ORDER BY Co.Population DESC;
```

Below the editor is the "Result Grid" which displays the following data:

#	Name	Population	Name	CountryCode
1	Cocos (Keeling) Islands	600	Bantam	CCK
2	Cocos (Keeling) Islands	600	West Island	CCK
3	Pitcairn	50	Adamstown	PCN

At the bottom is the "Action Output" pane showing the execution details:

Action	Time	Message
1	00:21:31	SELECT Co.Name, Co.Population, Ci.Name, Ci.CountryCode FROM Country AS ... 3 row(s) returned

# Unions

---

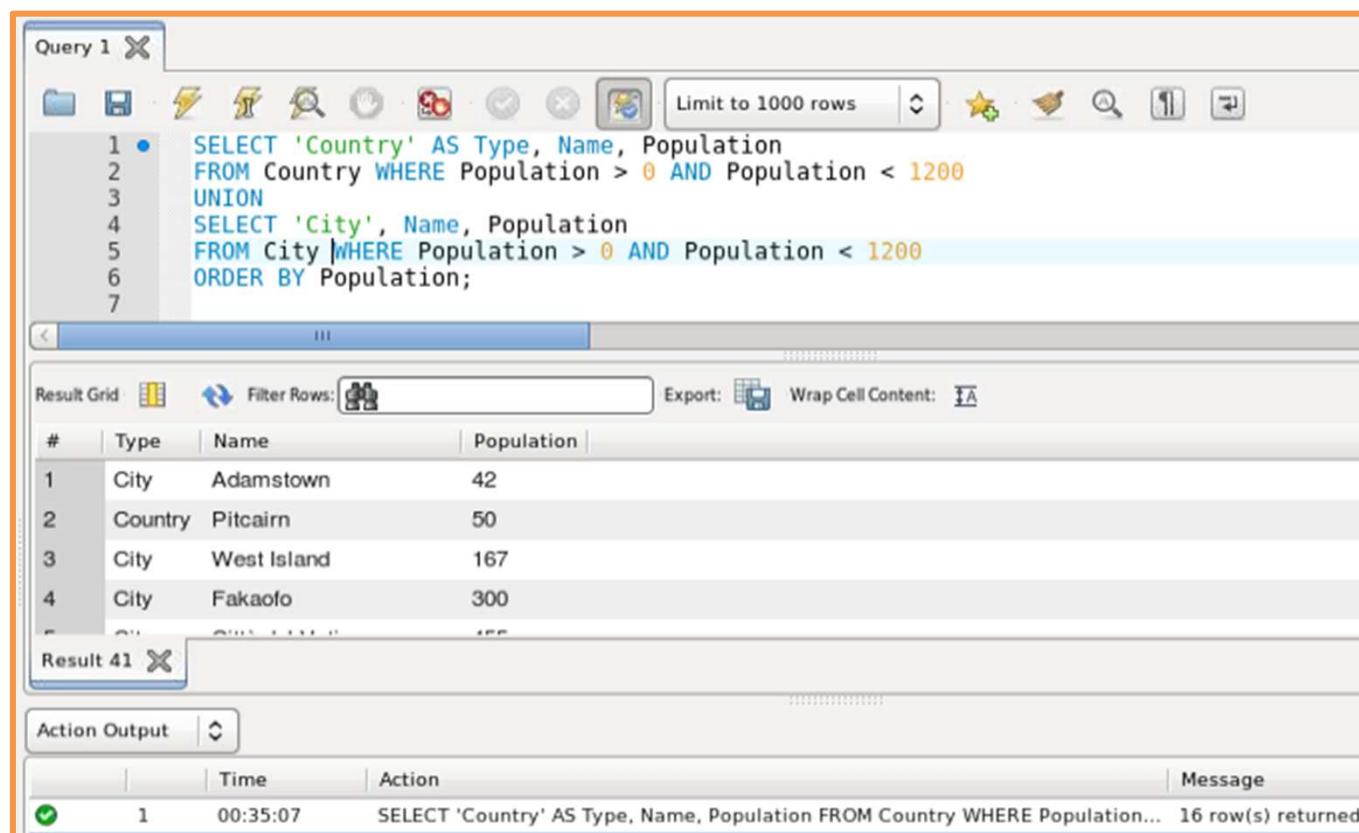
- Use the UNION keyword to combine results from multiple SELECT statements in a single result set.
- Syntax:

```
SELECT_statement
UNION [ALL | DISTINCT]
SELECT_statement
```

- The number of columns in each *SELECT\_statement* must be the same.
- The columns in the same position for each *SELECT\_statement* must be compatible data types.
- Default is UNION DISTINCT, which removes duplicate rows.
- Use UNION ALL to retain duplicate rows.



# Union Example



The screenshot shows the MySQL Workbench interface with a query editor and results grid.

**Query Editor:**

```
1 •  SELECT 'Country' AS Type, Name, Population
2   FROM Country WHERE Population > 0 AND Population < 1200
3
4  UNION
5  SELECT 'City', Name, Population
6   FROM City WHERE Population > 0 AND Population < 1200
7
8  ORDER BY Population;
```

**Result Grid:**

#	Type	Name	Population
1	City	Adamstown	42
2	Country	Pitcairn	50
3	City	West Island	167
4	City	Fakaofo	300

**Action Output:**

	Time	Action	Message
1	00:35:07	SELECT 'Country' AS Type, Name, Population FROM Country WHERE Population...	16 row(s) returned

# Subqueries

# Subqueries

---

- A *subquery* is a query that appears between parentheses as part of another SQL statement.
  - From the perspective of the query that contains the subquery, the subquery is sometimes referred to as a *nested query*.
  - From the perspective of the subquery itself, a query that (directly or indirectly) contains the subquery can be referred to as an *outer query*, a *containing query*, or an *enclosing query*.
  - An expression containing a subquery can be referred to as the *outer expression*, *containing expression*, or *enclosing expression*.



# Subquery: Example

The following query uses a subquery to list the languages spoken in Finland.

Outer Query

```
SELECT Language  
FROM CountryLanguage  
WHERE CountryCode = (
```

```
    SELECT Code  
    FROM Country  
    WHERE Name = 'Finland'
```

```
) ;
```

Subquery

Code
FIN

```
SELECT Language FROM CountryLanguage WHERE CountryCode = 'FIN';
```

# Subquery Categories

---

- Subqueries are categorized according to the type of results they return to the containing expression.
  - **Scalar subqueries:** The subquery is treated as a single *value*.
  - **Row subqueries:** The subquery is treated as a single *row*.
  - **Table subqueries:** The subquery is treated as a read-only *table*.



# Scalar Subqueries

- A *scalar subquery* evaluates to a single-value expression (a single row containing a single column).
  - This applies even if the subquery fails to retrieve a row.
    - If the row is retrieved, its column value is treated as a scalar value.
    - If no row is retrieved, the subquery evaluates to NULL.
  - You can use the value returned from a scalar subquery in the outer expression just like any other value.
  - If a scalar subquery returns more than one row or

```
mysql> SELECT 'Finland' = (SELECT Name FROM Country);
ERROR 1242 (21000): Subquery returns more than 1 row
```

```
mysql> SELECT 'Fin' = (SELECT * FROM Country);
ERROR 1241 (21000): Operand should contain 1 column(s)
```

# Example 1: Scalar Subquery

Scalar queries are useful for aggregating data. For example:

- List each country and its population as a

```
mysql> SELECT Country.Name,
->    100 * Country.Population /
->   (SELECT SUM(Population) FROM Country)
->   AS pct_of_world_pop
->   FROM Country;
+-----+-----+
| Name          | pct_of_world_pop |
+-----+-----+
| Aruba         |      0.0017     |
| Afghanistan  |      0.3738     |
| Angola        |      0.2119     |
| Anguilla      |      0.0001     |
| Albania       |      0.0560     |
| Andorra       |      0.0013     |
...
239 rows in set (#.## sec)
```

## Example 2: Scalar Subquery

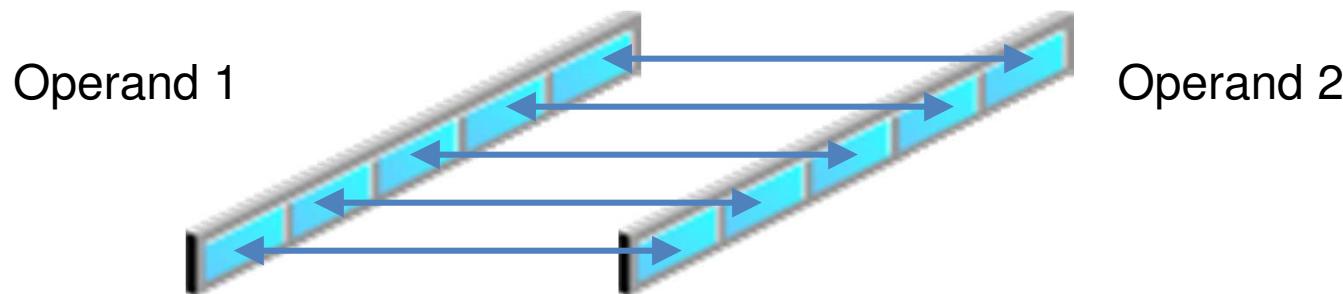
- List the number of cities and the number of languages for each country:

```
mysql> SELECT Name,
->   (SELECT COUNT(*) FROM City
->     WHERE CountryCode = Code) AS Cities,
->   (SELECT COUNT(*) FROM CountryLanguage
->     WHERE CountryCode = Code) AS Languages
->   FROM Country;
+-----+-----+-----+
| Name           | Cities | Languages |
+-----+-----+-----+
| Aruba          |      1 |        4 |
| Afghanistan    |      4 |        5 |
| Angola         |      5 |        9 |
| Anguilla        |      2 |        1 |
| Albania         |      1 |        3 |
| Andorra         |      1 |        4 |
| Netherlands Antilles | 1 |        3 |
| United Arab Emirates | 5 |        2 |
...
239 rows in set (#.## sec)
```

# Row Subqueries

---

- *Row subqueries* return a single row containing at least two columns.
  - If the subquery does not retrieve any rows, the expression in which you use it is FALSE.
- You can use row subqueries (and other row constructors) as operands for the comparison operators only if both operands are row constructors with the same number of columns.
  - MySQL evaluates these operators by performing a pairwise comparison of the individual (scalar) column values.



# Row Subquery Comparison

---

- For the *equality* operators (= and <= >=):
  - *All* pairwise comparisons need to evaluate to true in order for the expression to be true.
    - Two rows are equal only if all corresponding column values are equal.
- For the *inequality* operators (!= and <>):
  - *Only one* pairwise comparison must be true in order for the expression to be true.
    - If even one column pair is not equal, the rows are not equal.
- For the *comparison* operators (>, <, >= and <=):
  - The column *values* and the *order* of the columns is significant.
    - For example, (SELECT 2, 1) is larger than (SELECT 1, 100)



# Row Subqueries: Examples

- In the following example, the subquery is equal to the row constructor literal ('London', 'GBR'):

```
mysql> SELECT ('London', 'GBR') =
    -> (SELECT Name, CountryCode FROM City
    -> WHERE ID = 456) AS IsLondon;
+-----+
| IsLondon |
+-----+
| 1        |
+-----+
```

- The following example compares two subqueries for equality.

```
mysql> SELECT (SELECT ID, Name, CountryCode FROM City WHERE ID = 456) =
    -> (SELECT ID, Name, CountryCode FROM City WHERE CountryCode = 'GBR'
    -> AND Name = 'London') AS IsEqual;
+-----+
| IsEqual |
+-----+
| 1        |
+-----+
```

# Invalid Row Subqueries

- You get a runtime error if:
  - A row subquery retrieves more than one row:

```
mysql> SELECT (null, null) <=>  
      -> (SELECT GNP, Capital FROM Country LIMIT 2);  
ERROR 1242 (21000): Subquery returns more than 1 row
```

- The subquery returns two rows.
- You attempt to compare two row constructors with a different number of columns:

```
mysql> SELECT (null, null, null) <=>  
      -> (SELECT GNP, Capital FROM Country LIMIT 0);  
ERROR 1241 (21000): Operand should contain 3 column(s)
```

- The row constructor (null, null, null) contains two columns.
- The row subquery returns a row with three columns.

# Table Subqueries

---

- Return a result set containing zero or more rows with one or more columns
  - The result set behaves like a read-only table.
  - The number of rows or columns in the result set is irrelevant. It is still treated as a table.
- Can appear in two different contexts:
  - In the `FROM` clause of an enclosing query
  - As right operands to the following operators to return a Boolean result:
    - The logical operators `IN` and `EXISTS`
    - The regular comparison operators (`=`, `!=`, `<>`, `<`, `>`, `<=` or `>=`) quantified with `ALL`, `ANY`, and `SOME`



# Subqueries in the FROM Clause

- The result set of a subquery in the FROM clause is treated in the same way as a base table or a view.

```
SELECT *
FROM (SELECT Code, Name
      FROM Country
     WHERE IndepYear IS NOT NULL) AS IndependentCountry;
```

- The subquery executes first and produces a temporary result set
- You must specify a table alias for the result set or you

```
mysql> SELECT * FROM (SELECT Code, Name FROM Country WHERE IndepYear IS NOT NULL);
ERROR 1248 (42000): Every derived table must have its own alias
```

- The subquery result is discarded after the full statement is executed

# Calculating Aggregates of Aggregates

- Subqueries in the `FROM` clause are especially useful for calculating *aggregates of aggregates*. For example, to display the average population of each continent:

```
mysql> SELECT AVG(cont_sum)
-> FROM (
->   SELECT Continent, SUM(Population) AS cont_sum
->   FROM Country
->   GROUP BY Continent
-> ) AS t;
+-----+
| AVG(cont_sum) |
+-----+
| 868392778.5714 |
+-----+
1 row in set (#.## sec)
```

- This statement evaluates the table subquery first, calculating the total of all country populations for each continent.
- The outer query then uses the `AVG` function to aggregate the resulting rows (one for each continent).

# Subqueries Using the IN Operator

---

- You can use a table subquery as the right operand for the `IN` operator.
- The `IN` operator examines a table subquery for a value equal to its left operand, and returns:
  - `TRUE` if at least one such value exists
  - `FALSE` if there is no such value
- If the table subquery returns only a single column, the left operand must be a scalar value expression.



# Subquery Using the IN Operator: Example

```
mysql> SELECT * FROM City WHERE CountryCode IN(
    ->     SELECT Code FROM Country WHERE Continent = 'Asia');
+----+-----+-----+-----+-----+
| ID | Name          | CountryCode | District      | Population |
+----+-----+-----+-----+-----+
| 1  | Kabul          | AFG         | Kabul          | 1780000    |
| 2  | Qandahar       | AFG         | Qandahar       | 237500     |
| 3  | Herat          | AFG         | Herat          | 186800     |
| 4  | Mazar-e-Sharif | AFG         | Balkh          | 127800     |
| 64 | Dubai          | ARE         | Dubai          | 669181     |
| 65 | Abu Dhabi       | ARE         | Abu Dhabi       | 398695     |
| 66 | Sharja          | ARE         | Sharja          | 320095     |
| 67 | al-Ayn          | ARE         | Abu Dhabi       | 225970     |
| 68 | Ajman           | ARE         | Ajman           | 114395     |
| 126 | Yerevan        | ARM         | Yerevan        | 1248700    |
| 127 | Gjumri          | ARM         | Širak          | 211700     |
| 128 | Vanadzor        | ARM         | Lori           | 172700     |
| 144 | Baku            | AZE         | Baki           | 1787800    |
| 145 | Gänçä           | AZE         | Gänçä          | 299300     |
| 146 | Sumqayit        | AZE         | Sumqayit       | 283000     |
| 147 | Mingäçevir      | AZE         | Mingäçevir     | 93900      |
| 150 | Dhaka           | BGD         | Dhaka          | 3612850    |
| 151 | Chittagong      | BGD         | Chittagong     | 1392860    |
| 152 | Khulna          | BGD         | Khulna         | 663340     |
...
1766 rows in set (#.## sec)
```

# Pairwise IN Operation

- The following example illustrates *pairwise column comparison* to list Asian cities with the same name as their respective countries:

```
mysql> SELECT * FROM City WHERE (CountryCode, Name) IN (
    ->     SELECT Code, Name FROM Country WHERE Continent = 'Asia');
+----+-----+-----+-----+-----+
| ID | Name      | CountryCode | District | Population |
+----+-----+-----+-----+-----+
| 2429 | Kuwait    | KWT        | al-Asima | 28859   |
| 2454 | Macao     | MAC        | Macau    | 437500  |
| 3208 | Singapore | SGP        | -         | 4017733 |
+----+-----+-----+-----+-----+
3 rows in set (#.## sec)
```

- The subquery retrieves the codes and names of all Asian countries.
- The outer query retrieves cities and compares the city's name and country code to each of the Asian country code and name rows.

# Using NOT IN

---

- Negate the `IN` operator by prefixing the `IN` keyword with the `NOT` keyword.
  - `NOT IN` evaluates to `TRUE` if the result set does not contain an entry that is equal to the left operand.
  - If the result set does contain an entry equal to the left operand, `NOT IN` evaluates to `FALSE`.



# The IN Operator and NULL

---

- If it is impossible to determine whether the right operand contains a value equal to the left operand, `IN` evaluates to `NULL`.
- There are two different scenarios in which `IN` can evaluate to `NULL`:
  - If the left operand is `NULL` and the result set formed by the right operand is not empty
  - If the result set contains at least one row for which at least one of the columns is `NULL`, and no row matches the left operand



# EXISTS Operator

- The EXISTS operator accepts a single right argument, which must be a table subquery.
  - If the subquery result set contains at least one row, then the EXISTS expression evaluates to TRUE.
  - Otherwise, it returns FALSE.
- For example, the following query retrieves all cities that are the capital of some country:

```
mysql> SELECT * FROM City WHERE EXISTS (
    ->     SELECT NULL FROM Country WHERE Capital = ID);
+----+-----+-----+-----+-----+
| ID | Name          | CountryCode | District      | Population |
+----+-----+-----+-----+-----+
|   1 | Kabul         | AFG         | Kabol        | 1780000   |
|   5 | Amsterdam     | NLD         | Noord-Holland | 731200    |
|  33 | Willemstad   | ANT         | Curaçao      | 2345      |
...
232 rows in set (#.# sec)
```

# Using NOT EXISTS

- Negate EXISTS by placing the NOT keyword before the EXISTS keyword.
  - NOT EXISTS returns TRUE only if its table subquery argument is empty, and FALSE otherwise.
- The following query illustrates how to use NOT EXISTS to find all the countries where English is not spoken:

```
mysql> SELECT Code, Name FROM Country WHERE NOT EXISTS (
    ->     SELECT NULL FROM CountryLanguage WHERE CountryCode = Code
    ->     AND Language = 'English');
+----+-----+
| Code | Name          |
+----+-----+
| AFG  | Afghanistan   |
| AGO  | Angola         |
| ALB  | Albania        |
| AND  | Andorra        |
| ARE  | United Arab Emirates
...
179 rows in set (#.## sec)
```

# Table Subquery Quantifiers

- ALL, ANY, and SOME are quantifiers. They differ from operators in that they:
  - Are used in *conjunction* with the regular comparison operators (=, !=, <>, <, >, <=, and >=).
  - Specify how to apply an operator to repeatedly compare a single operand on the left side of an operator to the multiple values returned by a subquery on the right side of the operator.
- The following is an example using the ANY quantifier:

```
mysql> SELECT 'Finland' = ANY (
    ->     SELECT Name FROM Country) AS IsCountry;
+-----+
| IsCountry |
+-----+
|          1 |
+-----+
1 row in set (#.## sec)
```

# ALL, ANY, and SOME

---

- ALL: The comparison is TRUE if the operator is applied to all items in the subquery result set and returns TRUE in all cases.
  - For example, to list countries when their population exceeds that of all cities:

```
SELECT * FROM Country WHERE Population > ALL (
    SELECT Population FROM City);
```

- ANY: The operator is applied to the items in the subquery result set until at least one comparison returns TRUE. It returns FALSE if none of the comparisons returned TRUE.
- SOME: This is an alias of ANY, and has exactly the same effect. However, you might find the statement easier to read when using the SOME keyword instead of the ANY keyword.



# Alternatives to ANY and ALL

- ANY and ALL can sometimes be confusing to read. The following list shows more readable alternatives for several quantified operator expressions.

Quantified operator expression	Alternative
<code>scalar &gt; ANY (SELECT column...)</code>	<code>scalar &gt; (SELECT MIN(column) ...)</code>
<code>scalar &lt; ANY (SELECT column...)</code>	<code>scalar &lt; (SELECT MAX(column) ...)</code>
<code>scalar &gt; ALL (SELECT column...)</code>	<code>scalar &gt; (SELECT MAX(column) ...)</code>
<code>scalar &lt; ALL (SELECT column...)</code>	<code>scalar &lt; (SELECT MIN(column) ...)</code>
<code>scalar = ANY (SELECT column...)</code>	<code>scalar IN (SELECT column...)</code>
<code>scalar &lt;&gt; ALL (SELECT column...)</code>	<code>scalar NOT IN (SELECT column...)</code>

# Correlated Subqueries

A *correlated subquery*:

- Contains one or more expressions that are derived from the outer query and is, therefore, dependent on the outer query. For example:

```
SELECT *
FROM Country
WHERE NOT EXISTS (
    SELECT NULL
    FROM City
    WHERE CountryCode = Code
);
```

- The Code column referenced by the subquery depends on the Country table in the outer query.
- Can appear in any context where a subquery is valid, except in the FROM clause

# Correlated Subquery Scope

- Column references in correlated subqueries resolve using the nearest possible scope. Consider this

```
SELECT *
FROM City
WHERE CountryCode IN (
    SELECT Code
    FROM Country
    WHERE Name = 'Belgium'
);
```

Both tables contain  
Name columns

- The subquery is considered first and the nearest possible source for the Name column is the Country table in its FROM clause
- The subquery uses the Name column in the Country table
- The outer query uses the Name column in the City table

# Modifying Data with Subqueries

- You can use subqueries with data modification statements.

```
• DELETE FROM City
  WHERE CountryCode IN (
    SELECT Code FROM Country
    WHERE LifeExpectancy < 70.0
  );
```

```
UPDATE Country
SET Population = (
  SELECT SUM(Population) FROM City
  WHERE CountryCode = Code
);
```

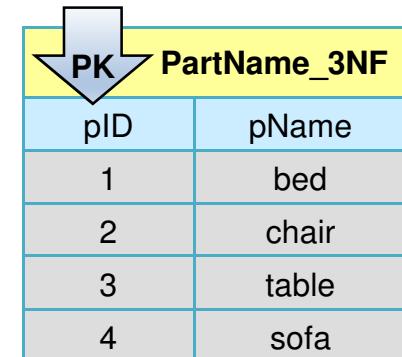


# Data Types

# Data Types as Part of Database Design

---

- A well-designed database uses the most suitable data type for each data item.
  - Think about the data each column needs to store.
  - Assign column data types accordingly:
    - Example: PartName\_3NF in the Inventory database. The pName column contains character data and the pID column contains numeric data.
  - Follow the ABCs of data types:
    - **A**ppropriate
    - **B**rief
    - **C**omplete



PartName_3NF	
pID	pName
1	bed
2	chair
3	table
4	sofa

# Data Types: Overview

---

- Numeric data types
  - Integer, decimal, and floating point data
- Temporal data types
  - Date and time data
- String data types
  - Character strings and binary strings



# Numeric Data Types

---

- Numeric data types:
  - **Integer:** Whole numbers
  - **Fixed-point:** Exact-value fractional numbers
  - **Floating-point:** Approximate-value fractional numbers
  - **Bit:** Bit-field values
- Factors to consider with numeric data types:
  - Range of values to be stored
  - Amount of storage space required
  - Column precision and scale (floating-point and fixed-point)



# Numeric: Integer Data Types

---

- Data types:
  - **TINYINT**: Very small integer data type
  - **SMALLINT**: Small integer data type
  - **MEDIUMINT**: Medium-sized integer data type
  - **INT, INTEGER**: Normal (average) size integer data type
  - **BIGINT**: Large integer data type
- Example:
  - world database, City table, Population column:  
**Population INT(11)**
  - Largest value stored: 10500000 (uses 8; 11 allowed)
- Attributes:
  - **UNSIGNED**: Has no negative numbers
  - **ZEROFILL**: Replaces padding spaces with zeroes



# Numeric: Integer Data Types Comparison

---

Type	Storage Required	Signed Range	Unsigned Range
<b>TINYINT</b>	1 byte	–128 to 127	0 to 255
<b>SMALLINT</b>	2 bytes	–32,768 to 32,767	0 to 65,535
<b>MEDIUMINT</b>	3 bytes	–8,388,608 to 8,388,607	0 to 16,777,215
<b>INT, INTEGER</b>	4 bytes	–2,147,683,648 to 2,147,483,647	0 to 4,294,967,295
<b>BIGINT</b>	8 bytes	–9,223,372,036,854,775,808 to 9,223,372,036,854,775,807	0 to 18,446,744,073,709,551,615



# Numeric: Fixed-Point Data Types

---

- Used for exact-value numbers: Integer, fractional, or both
- Data types:
  - **DECIMAL** (or **NUMERIC**):
    - Fixed-decimal, binary storage format
    - Preserve exact precision
- Example:
  - To represent currency values, such as dollars and cents:  
`cost DECIMAL(10, 2)`
  - Example value output:  
650.88



# Numeric: Floating-Point Data Types

---

- Used for approximate-value numbers: Integer, fractional, or both
- Data types:
  - **FLOAT**: Stores single-precision values using four bytes
  - **DOUBLE**: Stores double-precision values using eight bytes
- Can declare with precision and scale
- Example:
  - world database, Country table, GNP column:  
**GNP FLOAT(10, 2)**
  - Largest value stored in the GNP column:  
8510700.00 (up to 10 digits available, 9 displayed, 2 decimal digits)



# Numeric: Floating-Point Data Types Comparison

---

Type	Storage Required	Signed Range	Unsigned Range
FLOAT	4 bytes	-3.402823466E+38 to -1.175494351E-38 to 1.175494351E-38 to 3.402823466E+38	0 and 1.175494351E-38 to 3.402823466E+38
DOUBLE	8 bytes	-1.7976931348623157E+308 to -2.2250738585072014E-308 and 2.2250738585072014E-308 to 1.7976931348623157E+308	0 and 2.2250738585072014E-308 to 1.7976931348623157E+308



# Numeric: Bit Data Types

---

- Data type:
  - **BIT (M)**
- Bit-field values:
  - Column Width (**M**): Number of bits per value
  - 1 to 64 bits
- Example:
  - Storing 4 bits and 20 bits:  
**bit\_col1 BIT(4)**  
**bit\_col2 BIT(20)**
  - Can assign values by using bit-field notation
    - b'1111'** or **0b1111**
    - b'10000000'** or **0b10000000**



# Temporal Data Types

---

- **DATE:**
  - YYYY-MM-DD as in 2016-01-04
- **TIME:**
  - HH:MM:SS as in 12:59:02
- **DATETIME:**
  - YYYY-MM-DD HH:MM:SS as in 2016-01-04 12:59:02
- **TIMESTAMP:**
  - Current date and time as in 2016-01-04 12:59:02
- **YEAR:**
  - Four digits (YYYY) as in 1969
- TIME, DATETIME, and TIMESTAMP values can accept fractional seconds with up to microseconds (six digits of precision).



# Temporal Data Types Comparison

---

Type	Storage Required	Range
DATE	3 bytes	1000-01-01 to 9999-12-31
TIME	3 bytes	-838:59:59 to 838:59:59 Or -838:59:59.000000 to 838:59:59.000000
DATETIME	8 bytes	1000-01-01 00:00:00 to 9999-12-31 23:59:59 Or 1000-01-01 00:00:00.000000 to 9999-12-31 23:59:59.999999
TIMESTAMP	4 bytes	1970-01-01 00:00:01 to 2038-01-19 03:14:07 Or 1970-01-01 00:00:01.000000 to 2038-01-19 03:14:07.999999
YEAR	1 byte	1901 to 2155



# String Types

---

- CHAR and VARCHAR types
- BINARY and VARBINARY types
- BLOB and TEXT types
- ENUM type
- SET type



# String: CHAR and VARCHAR Types

---

- Data types:
  - **CHAR**: Fixed-length (static) character strings
  - **VARCHAR**: Variable-length (dynamic) character strings
- CHAR example:
  - world database, CountryLanguage table, Language column defined as:  
**Language CHAR (30)**
  - Largest value (needs 25):  
Southern Slavic Languages
- VARCHAR example:  
**mfrName VARCHAR (45)**



# String: BINARY and VARBINARY Types

---

- Data types:
  - **BINARY**: Fixed-length (static) binary byte strings
  - **VARBINARY**: Variable-length (dynamic) binary byte strings
- BINARY example:  
`bin_val BINARY(6)`
  - Stores 6 bytes of binary data, right-padded with zeros
- VARBINARY example:  
`varbin_val VARBINARY(6)`
  - Stores up to 6 bytes of binary data with no padding



# String: BLOB and TEXT Types

---

- **BLOB:**
  - **Binary Large Object**
  - Variable-length (dynamic) binary byte strings
  - Includes TINYBLOB, BLOB, MEDIUMBLOB, and LONGBLOB
- **TEXT:**
  - Variable-length (dynamic) character strings
  - Includes TINYTEXT, TEXT, MEDIUMTEXT, and LONGTEXT
- BLOB and TEXT data types can store extremely long values, such as media files.



# ENUM and SET Data Types

---

- Data types:
  - **ENUM**: Single string value chosen from a list of enumerated string values
  - **SET**: Zero or more string values chosen from a list of string values
- **ENUM example:**
  - Country table, Continent column:  
`Continent ENUM('Asia', 'Europe', 'North America', 'Africa', 'Oceana', 'Antarctica', 'South America')`
- **SET example:**  
`Symptom SET('sneezing', 'runny nose', 'stuffy head', 'red eyes')`



# String Data Types Comparison

Type	Storage Required	Maximum Length
<b>CHAR (M)</b>	$M$ bytes	255 characters
<b>BINARY (M)</b>	$M$ bytes	255 bytes
<b>VARCHAR (M), VARBINARY (M)</b>	$L + 1$ or 2 bytes	65,535 bytes
<b>TINYBLOB, TINYTEXT</b>	$L + 1$ bytes	255 bytes
<b>BLOB, TEXT</b>	$L + 2$ bytes	65,535 bytes
<b>MEDIUMBLOB, MEDIUMTEXT</b>	$L + 3$ bytes	16,777,215 bytes
<b>LONGBLOB, LONGTEXT</b>	$L + 4$ bytes	4,294,967,295 bytes
<b>ENUM</b>	1 or 2 bytes	65,535 values
<b>SET</b>	1, 2, 3, 4, or 8 bytes	64 members



# Character Set and Collation Support

---

- **Character set:** A named, encoded group of characters
  - All character strings belong to a specific character set.
- **Collation:** A named collating sequence for a character set
  - Defines character sort order
  - Affects comparison of characters and strings
- Column definitions for string data types can specify a character set or collation for the column.
- Character set and collation example:

```
CREATE TABLE t
( c1 VARCHAR(20) CHARACTER SET utf8,
  c2 TEXT CHARACTER SET latin1
    COLLATE latin1_general_cs );
```

# Choosing Data Types

---

- When selecting a data type, consider:
  - What sort of data it will store. For example, use a numeric data type for a number (instead of a character string).
  - The size of the data it will store:
    - Max size of integers, for example:
      - Age: **TINYINT UNSIGNED** (0 to 255)
      - A country's population: **INT UNSIGNED** (0 to 4,294,967,295)
    - Max length of character strings, for example:
      - Student grade: **CHAR (2)** (A+, B-)
      - Last name: **VARCHAR (64)**



# Setting Data Types to NULL

---

- **NULL** is a SQL keyword used to define data types that allow missing values. It means one of two things:
  - Unknown: There is a value, but the precise value is not known at this time.
  - Not applicable: If a value is specified, it would not be accurately representative.
- Use **NULL**:
  - In place of a real value in several situations
    - For example: “no value,” “unknown value,” “missing value,” “out of range,” “not applicable,” and so on
    - To represent an empty query result
  - Plan for null values during database design.



# When to Use NULL

---

- During database design, in cases where column information is not available, determine whether null values are allowed.
- For example, the `world` database's `Country` table contains countries with no life expectancy data and should allow nulls.
- Nulls are allowed by default.
- Use `NOT NULL` if a column must not allow nulls, to ensure data integrity.
- You can apply `NULL` and `NOT NULL` to existing column definitions.



# Working with Strings

# SQL Expressions

---

- A SQL expression is a combination of one or more values and SQL functions (“terms”) that, combined with operators and rules of precedence and association, evaluates to a single value. An expression in MySQL:
  - Can appear in many places as part of SQL statements, such as the `SELECT` list, the `WHERE` clause, the `ORDER BY` clause, and the `GROUP BY` clause to identify which rows to affect
  - Generally assumes the data type of its components
  - Can include many different terms, including:
    - Literals (numbers, strings, dates, and times)
    - Built-in constants such as `NULL`, `TRUE`, and `FALSE`
    - Column references
    - Function calls



# String Expressions

---

- When using string expressions you
  - Must quote literal string values
    - MySQL allows either single (' ) or double ( ' ' ) quotation marks.
    - The `ANSI_QUOTES` SQL mode interprets double-quoted characters as identifiers.
    - Single quotation marks are preferred for portability
  - Can use the following string data types: `CHAR`, `VARCHAR`, `TEXT`, and `BLOB`
  - Can use string expressions in comparison operations
    - Comparison operators (such as `=`, `<>`, `<`, and `BETWEEN...AND`) can be applied to string values.
  - Perform most string operations by using functions



# String Functions

---

- Perform operations on strings, such as:
  - Calculating string lengths
  - Extracting pieces of strings
  - Searching for substrings or replacing them
  - Performing case conversion
- String functions are divided into two categories:
  - **Numeric:** Return numbers
  - **String:** Return strings



# Concatenating Strings Using MySQL Functions

- Use the `CONCAT()` function

```
mysql> SELECT CONCAT('abc', 'def', REPEAT('X', 3));
+-----+
| CONCAT('abc', 'def', REPEAT('X', 3)) |
+-----+
| abcdefXXX
+-----+
1 row in set (0.00 sec)
```

- Use the `CONCAT_WS()` function

```
mysql> SELECT CONCAT_WS(' ', 'abc', 'def', REPEAT('X', 3));
+-----+
| CONCAT_WS(' ', 'abc', 'def', REPEAT('X', 3)) |
+-----+
| abc def XXX
+-----+
1 row in set (0.00 sec)
```

# Concatenating NULL Values

- Any NULL argument passed to CONCAT() causes it to return NULL.
- CONCAT\_WS() ignores NULL arguments.
- Comparing CONCAT() and CONCAT\_WS():

```
mysql> SELECT CONCAT('a', 'b'),
->   CONCAT('a', NULL, 'b');
+-----+-----+
| CONCAT('a', 'b') | CONCAT('a', NULL, 'b') |
+-----+-----+
| ab           | NULL          |
+-----+-----+
mysql> SELECT CONCAT_WS('/', 'a', 'b'),
->   CONCAT_WS('/', 'a', NULL, 'b');
+-----+-----+
| CONCAT_WS('/', 'a', 'b') | CONCAT_WS('/', 'a', NULL, 'b') |
+-----+-----+
| a/b           | a/b           |
+-----+-----+
```

# Concatenating Strings Using the OR (||) Operator

1. Set the PIPES\_AS\_CONCAT SQL mode

```
mysql> SET sql_mode='PIPES_AS_CONCAT';
Query OK, 0 rows affected, 1 warning (0.00 sec)
```

2. Use the OR (||) operator to concatenate strings

```
mysql> SELECT 'abc' || 'def';
+-----+
| 'abc' || 'def' |
+-----+
| abcdef         |
+-----+
1 row in set (0.0 sec)
```

# Retrieving the Left and Right Portions of a String

- Use `LEFT(str, len)` to retrieve the `len` left-most characters of the string `str`

- Use  $RIGHT(str, len)$  to retrieve the  $len$  right-most characters of the string  $str$

# Retrieving Part of a String

- Use `SUBSTRING(str, pos, len)` to return `len` characters of the string, starting at `pos`

```
mysql> SELECT SUBSTRING('Alice and Bob', 1, 5);
+-----+
| SUBSTRING('Alice and Bob', 1, 5) |
+-----+
| Alice                                |
+-----+
1 row in set (0.00 sec)
```

Return a part of the first string starting from the offset:

```
mysql> SELECT SUBSTRING_INDEX(
    -> 'Alice and Bob', 'and', 1)
    -> AS Result;
+-----+
| Result |
+-----+
| Alice  |
+-----+
1 row in set (0.00 sec)
```

Return a part of the second string starting from the offset:

```
mysql> SELECT SUBSTRING_INDEX(
    -> 'Alice and Bob', 'and', -1)
    -> AS Result;
+-----+
| Result |
+-----+
| Bob   |
+-----+
1 row in set (0.00 sec)
```

# Trimming a String

- Use LTRIM(), RTRIM(), or TRIM() to remove space characters appearing at the left, right, or both extremities of string, respectively

```
mysql> SELECT CONCAT('<', LTRIM(' Alice '), '>'),
-> CONCAT('<', RTRIM(' Alice '), '>'),
-> CONCAT('<', TRIM(' Alice '), '>') \G
*****
1. row ****
CONCAT('<', LTRIM(' Alice '), '>'): <Alice >
CONCAT('<', RTRIM(' Alice '), '>'): < Alice>
CONCAT('<', TRIM(' Alice '), '>'): <Alice>
```

- Remove spaces or specific characters from the beginning (LEADING), end (TRAILING) or both (BOTH) sides of a string. For example:

```
mysql> SELECT TRIM(LEADING 'Cha' FROM 'ChaChaChalice');
+-----+
| TRIM(LEADING 'Cha' FROM 'ChaChaChalice') |
+-----+
| lice                                         |
+-----+
```

# Inserting Into and Replacing Portions of a String

- Use `REPLACE(str, substr, newstring)` to replace `substr` in `str` with `newstring`

```
mysql> SELECT REPLACE('Alice & Bob', '&', 'and');  
+-----+  
| REPLACE('Alice & Bob', '&', 'and') |  
+-----+  
| Alice and Bob |  
+-----+
```

- Use `INSERT(str, pos, len, newstring)` to place `newstring` at `pos` in `str`, overwriting `len` characters

```
mysql> SELECT INSERT('Alice and Bob', 6, 5, ', Carol & ');  
+-----+  
| INSERT('Alice and Bob', 6, 5, ', Carol & ') |  
+-----+  
| Alice, Carol & Bob |  
+-----+
```

# Determining String Length

- Use the LENGTH() and CHAR\_LENGTH() to return the string lengths in byte and character

```
mysql> SELECT LENGTH('MySQL'),  
-> CHAR_LENGTH('MySQL');  
+-----+-----+  
| LENGTH('MySQL') | CHAR_LENGTH('MySQL') |  
+-----+-----+  
| 5 | 5 |  
+-----+-----+  
  
mysql> SELECT LENGTH(CONVERT('MySQL' USING ucs2))  
-> AS length,  
-> CHAR_LENGTH(CONVERT('MySQL' USING ucs2))  
-> AS c_length;  
+-----+-----+  
| length | c_length |  
+-----+-----+  
| 10 | 5 |  
+-----+-----+
```

# Comparing Strings

- Use `STRCMP(str1, str2)` to compare `str1` and `str2`.
  - Check return value:  $0 = str1$  and  $str2$  are the same,  $-1 = str1$  is smaller than  $str2$ ,  $1 = str1$  is larger than  $str2$

```
mysql> SELECT STRCMP ('abc' , 'def') ,
   -> STRCMP ('def' , 'def') ,
   -> STRCMP ('def' , 'abc') ;
+-----+-----+-----+
| STRCMP ('abc' , 'def') | STRCMP ('def' , 'def') | STRCMP ('def' , 'abc') |
+-----+-----+-----+
| -1          | 0           | 1           |
+-----+-----+-----+
```

- Use the equals operator (`=`) to test for string equality:

```
mysql> SELECT 'abc' = 'def' , 'def' = 'def' , 'def' = 'abc' ;
+-----+-----+-----+
| 'abc' = 'def' | 'def' = 'def' | 'def' = 'abc' |
+-----+-----+-----+
| 0           | 1           | 0           |
+-----+-----+-----+
```

# Finding Strings Within Strings

- Use INSTR(), LOCATE() or POSITION() to return the position of the specified string inside another string:

```
mysql> SELECT INSTR('Alice and Bob', 'and'),
-> LOCATE('and', 'Alice and Bob'),
-> POSITION('and' IN 'Alice and Bob')
-> \G
***** 1. row *****
INSTR('Alice and Bob', 'and'): 7
LOCATE('and', 'Alice and Bob'): 7
POSITION('and' IN 'Alice and Bob'): 7
```

- Use LOCATE(*searchstr, str, pos*) , where *pos* specifies the start position for the search :

```
mysql> SELECT LOCATE(' ', 'Alice and Bob', 7);
+-----+
| LOCATE(' ', 'Alice and Bob', 7) |
+-----+
| 10
+-----+
```

# String Comparison and Sorting Behavior

- String comparison and sort behavior depends on the *collation*
- Collations can be case sensitive, case insensitive, and language-specific
- The default collation in MySQL is latin1\_swedish\_ci. The \_ci suffix denotes a case insensitive collation. Example behavior:

```
mysql> CREATE TABLE t1 (col1 CHAR(3) CHARACTER SET latin1);
mysql> INSERT INTO t1 (col1) VALUES('AAAA') , ('bbbb') , ('aaaa') , ('BBBB');
mysql> SELECT col1 FROM t1 ORDER BY col1;
+----+
| c  |
+----+
| AAAA |
| aaaa |
| bbbb |
| BBBB |
+----+
```

# Character Sets and Collations

---

- A *character set* is a set of symbols and encodings. A *collation* is a set of rules for comparing characters in a character set.
- MySQL allows you to:
  - Store strings using a variety of character sets.
  - Compare strings using a variety of collations.
  - Mix strings with different character sets or collations in the same server, the same database, or even the same table.
  - Enable specification of character set and collation at any level.
  - Display the available character sets and collations using the following syntax

```
SHOW COLLATION [LIKE 'pattern' | WHERE expr]
```

```
SHOW CHARACTER SET [LIKE 'pattern' | WHERE expr]
```

# Character Set and Collation Functions

- Use `CHARSET(str)` and `COLLATION(str)` to retrieve the character set and collation of `str`, respectively

```
mysql> SELECT USER(), CHARSET(USER()), COLLATION(USER());
+-----+-----+-----+
| USER() | CHARSET(USER()) | COLLATION(USER()) |
+-----+-----+-----+
| root@localhost | utf8 | utf8_general_ci |
+-----+-----+
```

- Use `CONVERT()` or `CAST()` to change a string's character set

```
mysql> SELECT
CONVERT(_latin1'Müller'
USING utf8) AS Result;
+-----+
| Result |
+-----+
| Müller |
+-----+
1 row in set (0.00 sec)
```

```
mysql> SELECT CAST(_latin1'test'
AS CHAR CHARACTER SET utf8) AS
Result;
+-----+
| Result |
+-----+
| test |
+-----+
1 row in set (0.00 sec)
```

# Changing the Default Collation

- Use the `COLLATE` operator on the column you want to affect

```
mysql> SELECT col1 FROM t1 ORDER BY col1 COLLATE latin1_general_ci;  
+-----+  
| col1 |  
+-----+  
| AAAA |  
| aaaa |  
| BBBB |  
| bbbb |  
+-----+
```

- Specify a character set and collation for the entire database

```
CREATE DATABASE mydb  
  DEFAULT CHARACTER SET utf8  
  DEFAULT COLLATE utf8_general_ci;
```

- Set the `--character-set-server` and `--collation-server` options when you start the MySQL server

# String Pattern Matching

---

- Use the `LIKE` pattern-matching operator to perform comparisons based on similarity.
  - General syntax:

```
SELECT...<expression> LIKE '<pattern>'...
```

- Use metacharacters with `LIKE`:
  - The percent character ‘%’ matches any sequence of zero or more characters.
  - The underscore character ‘\_’ matches any single character.
- Use `NOT LIKE` for opposite comparisons.



# String Pattern Matching: LIKE Versus NOT LIKE

- Example of LIKE:

```
mysql> SELECT Name FROM Country
      -> WHERE Name LIKE 'United%';
+-----+
| Name
+-----+
| United Arab Emirates
| United Kingdom
| United States Minor Outlying Isl.
| United States
+-----+
```

- Example of logical NOT LIKE:

```
mysql> SELECT Name FROM Country
      -> WHERE Name NOT LIKE 'United%';
+-----+
| Name
+-----+
| Aruba
...
| Zimbabwe
+-----+
```

# String Pattern Matching: Regular Expressions

---

A regular expression is a powerful way of specifying a pattern for a complex search.

- Use the RLIKE or REGEXP operators.
  - General syntax:

```
SELECT... <string-expr>
      RLIKE | REGEXP '<regexp-string>'...
```

- Use for complex patterns that go beyond general wildcard characters such as '%' and '\_' in LIKE patterns.
- Use to test the following:
  - URLs
  - IP addresses
  - Email addresses
  - Postal codes



# String Pattern Matching: RLIKE Syntax

---

- The RLIKE expression is a character string with its own syntax to allow a declarative denotation of complex patterns.
    - Literal text: Match character *exactly*.
    - Period: (.) Match any single character.
    - Occurrence: ( \*, ?, +, {m, n} ) Specify the number of occurrences.
    - Choice: ( | ) Match either of two alternatives.
    - Escape: (\) Escape special characters
    - Anchors: Match special positions. (^ and \$ match start and end of input; [[:>:] ] and [[:<:] ] match word boundaries.)
    - Character classes: [spec] Match against a collection.
    - Parentheses: ( and ) Group a sequence of elements.
-

# String Pattern Matching: RLIKE Character Class

---

- Notate character class specifications in brackets ( [ and ] ).

The syntax for this notation is as follows:

- Literal text: ' [abc] ' matches either 'a' or 'b' or 'c'.
- Negation: ' [^abc] ' matches everything *except* 'a', 'b', or 'c'.
- Ranges: ' [a-z] ' matches 'a', 'b', 'c',... 'z'.
- Named range: ' [[:space:]] ' matches all whitespace characters.



# String Pattern Matching: Regular Expressions

- Simple (all matches possible):

```
mysql> SELECT Name FROM City WHERE Name RLIKE 'nat';
+-----+
| Name      |
+-----+
| Natal   |
| Cabanatuan |
| Maunnath Bhanjan |
| Minatitlβn |
| Cincinnati |
+-----+
```

- Anchors (only complete match):

```
mysql> SELECT Name FROM City WHERE Name RLIKE '^new.*rk$';
+-----+
| Name      |
+-----+
| New York |
| Newark   |
+-----+
```



# String Pattern Matching: Regular Expressions

- Alternation (choice):

```
mysql> SELECT Name FROM City
      -> WHERE Name RLIKE ' Los | Las ';
+-----+
| Name           |
+-----+
| San Nicol s de los Arroyos |
| Santiago de los Caballeros |
| Santo Domingo de los Colorados |
| Victoria de las Tunas |
| San Nicol s de los Garza |
| Chilpancingo de los Bravo |
| San Crist bal de las Casas |
| East Los Angeles |
| North Las Vegas |
+-----+
```

- Equivalent character class notation:

```
mysql> SELECT Name FROM City
      -> WHERE Name RLIKE ' L[aos]s ';
```

# String Pattern Matching: Ranges of Characters

- Character classes offer convenient features to match an entire range of characters. The following example uses a character range to match Argentian postal codes:

```
SELECT CityName, StreetName FROM Addresses  
WHERE PostalCode  
RLIKE '^ [A-HJ-NP-Z] [0-9] {4} ([A-Z] {3}) ?$';
```

- Letter (A..H; J..N; P..Z) for the province: [A-HJ-NP-Z]
- Four digits for the municipality: [0-9] {4}
- (Optional) Three letters for the side of the street block:  
( [A-Z] {3} ) ?



# String Pattern Matching: Escaping Characters

To match characters that are part of the regular expression syntax themselves, you must “escape” them.

- Requires ‘double’ escaping:

```
mysql> SELECT '\\\' RLIKE '\\\'';
ERROR 1139 (42000): Got error 'trailing backslash (\)' from regexp

mysql> SELECT '\\\' RLIKE '\\\\\'';
+-----+
| '\\\' RLIKE '\\\\\'' |
+-----+
| 1 |
+-----+
```

- No escape possible inside character class:

```
mysql> SELECT '\\\' RLIKE '[\\\\]';
+-----+
| '\\\' RLIKE '[\\\\]' |
+-----+
| 1 |
+-----+
```

# Full-Text Search: Overview

---

- Characteristics of `LIKE` and regular expression search:
  - **Poor performance:** MySQL has to scan the entire table.
  - **Little flexibility:** It is hard to implement certain searches.
  - **No relevance factor:** You cannot specify which results are most relevant.
- Characteristics of full-text search:
  - Requires the `FULLTEXT` index on searchable text columns
  - Better performance for complex queries
  - An automatically rebuilt index when data changes
  - Native SQL-like interface to search
  - Used on `CHAR`, `VARCHAR`, or `TEXT` columns
  - Works for `InnoDB` and `MyISAM` tables
  - Supports Natural Language and Boolean searches



# Full-Text Search: Natural Language Mode

- Default is IN NATURAL LANGUAGE MODE.
- Example using the `film` table in `sakila` database:
  - Add FULLTEXT index to `description` column:

```
ALTER TABLE film
ADD FULLTEXT (description);
```

- Use `MATCH()` . . . `AGAINST()` with delimited list of search keywords:

```
mysql> SELECT title, description FROM film
-> WHERE MATCH(description) AGAINST ('composer, monkey');
...
164 rows in set (0.00 sec)
```

- Results are returned in order of relevance
- Use delimiters between keywords:
  - Example: “ ” (space), “,” (comma), “.” (period)
- Use a stopword list to ignore certain words.

# Full-Text Search: Boolean Mode

- Specify `IN BOOLEAN MODE` in the `AGAINST()` function to search for *words* instead of *concepts*
- Enables very sophisticated queries using operators
- Example using the `-` (omit) operator to find all occurrences of the word “composer” in `description` where there is no instance of the word “monkey”:

```
mysql> SELECT title, description FROM film
-> WHERE MATCH(description)
-> AGAINST ('composer -monkey' IN BOOLEAN MODE );
...
83 rows in set (0.00 sec)
```

- Results are not returned in order of relevance.

# Full-Text Search: Boolean Mode Operators

Operator	Usage
no operator	The word is optional, but rows containing it are ranked higher, as IN NATURAL LANGUAGE MODE.
@distance	Tests whether two words start in a specified number of words of each other (InnoDB only)
+	This word must be present in every row returned by the query.
-	This word must not be present in every row returned by the query.
> and <	Increases (>) or decrease(<) a word's relevance
( )	Parentheses group words into subexpressions.
~	Is used to mark “noise” words, causing their contribution to the row’s relevance to be negative. Rows with negated words have a lower overall relevance but are not omitted entirely.
*	Is appended to a word as a wildcard
“ ”	Rows must match the literal contents of the string.



# Working with Numeric and Temporal Data

# Numeric Expressions

---

- Numbers can be exact-value or approximate-value literals.
  - **Exact-value literals:**
    - Used just as specified in SQL statements
    - Written as integer or decimal values, with no exponent
    - Not subject to the inexactness produced by rounding error
  - **Approximate-value literals:**
    - Not always used as specified in SQL statements
    - Written as floating-point numbers in scientific notation, with an exponent
    - Subject to rounding errors
- Nearly all numeric expressions that contain NULL return NULL.



# Numeric Expressions: Examples

- Comparing the use of exact-values literal compared to approximate-value literals in expressions

```
mysql> SELECT 1.1 + 2.2 = 3.3, 1.1E0 + 2.2E0 = 3.3E0;
+-----+-----+
| 1.1 + 2.2 = 3.3 | 1.1E0 + 2.2E0 = 3.3E0 |
+-----+-----+
| 1           | 0           |
+-----+-----+
```

- Mixing numbers with strings

```
mysql> SELECT 1 + '1', 1 = '1';
+-----+-----+
| 1 + '1' | 1 = '1' |
+-----+-----+
| 2       | 1       |
+-----+-----+
```

# Temporal Data Types

---

Data Type	Default Format
DATE	YYYY-MM-DD
TIME	HH:MI:SS
DATETIME	YYYY-MM-DD HH:MI:SS
TIMESTAMP	YYYY-MM-DD HH:MI:SS
DAY	DD
MONTH	MM
QUARTER	Q
YEAR	YYYY



# Interval Arithmetic

- To perform interval arithmetic, use the `INTERVAL` keyword and specify the unit value.
  - The order of operands in interval addition is not

```
mysql> SELECT '2012-01-01' + INTERVAL 10 DAY,  
-> INTERVAL 10 DAY + '2012-01-01';  
+-----+  
| '2012-01-01' + INTERVAL 10 DAY | INTERVAL 10 DAY + '2012-01-01' |  
+-----+  
| 2012-01-11 | 2012-01-11 |  
+-----+
```

- In interval subtraction, the interval value must be the

```
mysql> SELECT '2012-01-01' - INTERVAL 10 DAY;  
+-----+  
| '2012-01-01' - INTERVAL 10 DAY |  
+-----+  
| 2011-12-22 |  
+-----+
```



# Numeric Functions

---

- Numeric functions perform different of mathematical operations.

Function Syntax	Definition
ABS (<number>)	Returns the absolute value of number
SIGN (<number>)	Returns –1, 0, or 1 depending on whether the number is negative, zero, or positive
TRUNCATE (<number>, <decimals>)	Returns number truncated to decimals
FLOOR (<number>)	Rounds number down to the closest lower integer
CEILING (<number>)	Rounds number up to the closest higher integer
ROUND (<number>)	Rounds number to the closest integer
SIN(), COS(), TAN(), PI(), DEGREES(), RADIANS()	Perform trigonometric calculations, including conversions between degrees and radians



# Using ABS() and SIGN()

- ABS () returns the absolute value of negative and positive values

```
mysql> SELECT ABS (-42) , ABS (42) ;  
+-----+-----+  
| ABS (-42) | ABS (42) |  
+-----+-----+  
|        42 |        42 |  
+-----+-----+
```

- SIGN () returns the results of sign determination  
(-1 = negative, 0 = zero, 1 = positive)

```
mysql> SELECT SIGN (-42) , SIGN (-1) , SIGN (0) , SIGN (1) , SIGN (42) ;  
+-----+-----+-----+-----+-----+  
| SIGN (-42) | SIGN (-1) | SIGN (0) | SIGN (1) | SIGN (42) |  
+-----+-----+-----+-----+-----+  
|    -1      |    -1      |    0     |    1     |    1     |  
+-----+-----+-----+-----+-----+
```

# Approximating Numbers With FLOOR() and CEILING()

- FLOOR () returns the largest integer not greater than its argument

```
mysql> SELECT FLOOR(-14.7), FLOOR(14.7);  
+-----+-----+  
| FLOOR(-14.7) | FLOOR(14.7) |  
+-----+-----+  
| -15          | 14           |  
+-----+-----+
```

- CEILING () returns the smallest integer not less than its argument

```
mysql> SELECT CEILING(-14.7), CEILING(14.7);  
+-----+-----+  
| CEILING(-14.7) | CEILING(14.7) |  
+-----+-----+  
| -14          | 15           |  
+-----+-----+
```

# Approximating Numbers With ROUND()

- Rounding up to the next integer with exact values

```
mysql> SELECT ROUND(28.5), ROUND(-28.5);  
+-----+-----+  
| ROUND(28.5) | ROUND(-28.5) |  
+-----+-----+  
| 29          | -29         |  
+-----+-----+
```

- Rounding up to the next integer with approximate values:

```
mysql> SELECT ROUND(2.85E1), ROUND(-2.85E1);  
+-----+-----+  
| ROUND(2.85E1) | ROUND(-2.85E1) |  
+-----+-----+  
| 28          | -28         |  
+-----+-----+
```

# Trigonometric Functions

- Perform trigonometric calculations, including conversions between degrees and radians.

```
mysql> SELECT SIN(0), COS(0), TAN(0);
```

SIN(0)	COS(0)	TAN(0)
0	1	0

...

```
mysql> SELECT PI(), DEGREES(PI()), RADIANS(180);
```

PI()	DEGREES(PI())	RADIANS(180)
3.141593	180	3.1415926535898

# Temporal Functions

---

- Perform operations such as:
  - Extracting parts of dates and times
  - Reformatting values
  - Converting values to seconds or days
- Use more than one format for the same information:
  - 2012-02-10 22:50:15
  - Friday, February 10, 2012
- Generate temporal data in many ways:
  - Copy existing data.
  - Execute the built-in function.
  - Build a string representation to be evaluated by the server.



# Temporal Functions: Function Types

Function Syntax	Definition
NOW()	Current time as set on the server host ( <code>TIME</code> format)
CURDATE()	Current date as set on the server host ( <code>DATE</code> format)
CURTIME()	Current time as set on the server host ( <code>TIME</code> format)
DATEDIFF()	Subtracts two dates
DATE_ADD()	Adds time values (intervals) to a date value
YEAR(<date_expression>)	Year in four-digit <code>YEAR</code> format
MONTH(<date_expression>)	Month of the year in integer format, per expression
DAYOFMONTH(<date_expression>) or DAY(<date_expression>)	Day of the month in integer format, per expression
DAYNAME(<date_expression>)	Day of the week in string format, per expression (English)
HOUR(<date_expression>)	Hour of the day in integer format (in 0–23 range), per expression
MINUTE(<date_expression>)	Minute of the day in integer format, per expression
SECOND(<date_expression>)	Second of the minute in integer format, per expression
GET_FORMAT(<date_expression>)	Returns a date format string, per values indicated for date-type and international format



# Displaying the Current Date and Time

- Display current date and time:

```
mysql> SELECT NOW();  
+-----+  
| NOW() |  
+-----+  
| 2015-11-11 10:40:43 |  
+-----+
```

- Extract current temporal data (date, time, and day of week):

```
mysql> SELECT CURDATE(), CURTIME(), DAYNAME(NOW());  
+-----+-----+-----+  
| CURDATE() | CURTIME() | DAYNAME(NOW()) |  
+-----+-----+-----+  
| 2015-11-11 | 10:41:47 | Wednesday |  
+-----+-----+-----+
```

# Determining Date and Time Format

- Use the `GET_FORMAT()` function to retrieve the format for the specified data type and standard
  - Syntax:

```
GET_FORMAT({DATE|TIME|DATETIME},  
{'EUR' | 'USA' | 'JIS' | 'ISO' | 'INTERNAL'})
```

- Example, using DATE type and EUR/USA formats

```
mysql> SELECT GET_FORMAT(DATE, 'EUR'), GET_FORMAT(DATE, 'USA');  
+-----+-----+  
| GET_FORMAT(DATE, 'EUR') | GET_FORMAT(DATE, 'USA') |  
+-----+-----+  
| %d.%m.%Y           | %m.%d.%Y          |  
+-----+-----+
```

# Customizing the Output Format

- Use the `DATE_FORMAT()` function with format

```
mysql> SELECT NOW(), DATE_FORMAT(NOW(), '%W the %D of %M');
+-----+-----+
| NOW() | DATE_FORMAT(NOW(), '%W the %D of %M') |
+-----+-----+
| 2015-11-11 11:11:04 | Wednesday the 11th of November |
+-----+-----+
```

- Other specifiers include
  - `%a`: Abbreviated weekdays name (Sun...Sat)
  - `%b`: Abbreviated month name (Jan...Dec)
  - `%d`: Numeric day of month (0, 1, 2...)
  - `%m`: Numeric month (1...12)
  - `%r`: 12 hr time (hh:mm:ss, followed by AM or PM)
  - `%T`: 24 hr time (hh:mm:ss)
  - `%Y`: Four-digit year (YYYY)

# Retrieving Components of Dates and Times

```
mysql> SELECT YEAR('2015-11-11'),
-> MONTH ('2015-11-11'),
-> DAYOFMONTH('2015-11-11');
+-----+-----+
| YEAR('2015-11-11') | MONTH ('2015-11-11') | DAYOFMONTH('2015-11-11') |
+-----+-----+
|          2015 |                 11 |                  11 |
+-----+-----+
1 row in set (0.00 sec)

mysql> SELECT DAYOFYEAR('2015-11-11');
+-----+
| DAYOFYEAR('2015-11-11') |
+-----+
|          315 |
+-----+
1 row in set (0.00 sec)

mysql> SELECT HOUR('11:28:45'),
-> MINUTE('11:28:45'),
-> SECOND('11:28:45');
+-----+-----+
| HOUR('11:28:45') | MINUTE('11:28:45') | SECOND('11:28:45') |
+-----+-----+
|          11 |                 28 |                  45 |
+-----+-----+
1 row in set (0.00 sec)
```

# Creating Date and Time Values

- Use `MAKEDATE (year, dayofyear)` to create dates

```
mysql> SELECT MAKEDATE(2015, 230);  
+-----+  
| MAKEDATE(2015, 230) |  
+-----+  
| 2015-08-18 |  
+-----+
```

- Use `MAKETIME (hour, min, sec)` to create times

```
mysql> SELECT MAKETIME(11, 20, 35);  
+-----+  
| MAKETIME(11, 20, 35) |  
+-----+  
| 11:20:35 |  
+-----+
```

# Tables

# Using the SHOW CREATE TABLE Statement

- Issue `SHOW CREATE TABLE` to understand the structure of a table. It displays the `CREATE TABLE` statement that creates the specified table. For example:

```
mysql> SHOW CREATE TABLE City\G
***** 1. row *****
      Table: City
Create Table: CREATE TABLE `City` (
    `ID` int(11) NOT NULL auto_increment,
    `Name` char(35) NOT NULL default '',
    `CountryCode` char(3) NOT NULL default '',
    `District` char(20) NOT NULL default '',
    `Population` int(11) NOT NULL default '0',
    PRIMARY KEY  (`ID`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1
1 row in set (0.00 sec)
```

# Creating New Tables from Existing Tables

---

- MySQL provides two ways to create a new table based on an existing table:
  - Use `CREATE TABLE` with `SELECT` to create a new table that stores the result set returned by the `SELECT` query.
    - You can use almost any type of `SELECT` statement and reference multiple tables with joins and unions.
  - Use `CREATE TABLE` with `LIKE` to create a table with the same structure as the one specified.
    - Does not keep foreign key assignments
    - Does not copy any data



# Creating a Table from Data in Existing Tables

---

- Use `CREATE TABLE with SELECT`. With this approach:
  - You do not have to create the table structure first and then populate it with data.
  - The statement creates the table based on the `SELECT` query results.
  - You can use multiple tables with joins, subqueries, or unions in the `SELECT` statement, or even a `SELECT` statement without any table.
  - Your new table might not retain all the features of the original:
    - No table options are copied, including indexes and constraints
    - The data types might be different, particularly if the column is based on an expression



# Examples: CREATE TABLE with SELECT

---

- The following examples create new tables based on the City table. The new tables contain:
  - All rows from the original City table:

```
CREATE TABLE CitySelect1 SELECT * FROM City;
```

- All cities with a population of over 2,000,000:

```
CREATE TABLE CitySelect2 SELECT * FROM City  
WHERE Population > 2000000;
```

- Only the Name and CountryCode columns:

```
CREATE TABLE CitySelect3 SELECT Name, CountryCode  
FROM City;
```

- An empty copy of the table:

```
CREATE TABLE CitySelect4 SELECT * FROM City LIMIT 0;
```



# Copying an Existing Table Structure

---

- Use CREATE TABLE with LIKE. With this approach:
  - Indexes and column options are preserved in the new table
  - No row data is copied

```
mysql> CREATE TABLE CityLike LIKE City;  
Query OK, 0 rows affected (0.05 sec)
```

- The new table is empty



# Comparing SELECT with LIKE Results

```
mysql> CREATE TABLE CitySelect
    SELECT * FROM City;

Query OK, 4079 rows affected (1.44 sec)
Records: 4079  Duplicates: 0  Warnings: 0

mysql> SHOW CREATE TABLE CitySelect\G
***** 1. row *****
    Table: CitySelect
Create Table: CREATE TABLE `CitySelect`(
  `ID` int(11) NOT NULL DEFAULT '0',
  `Name` char(35) NOT NULL DEFAULT '',
  `CountryCode` char(3) NOT NULL
DEFAULT '',
  `District` char(20) NOT NULL DEFAULT
 '',
  `Population` int(11) NOT NULL DEFAULT
'0'
) ENGINE=InnoDB DEFAULT CHARSET=latin1
1 row in set (0.00 sec)
```

```
mysql> CREATE TABLE CityLike
    LIKE City;

Query OK, 0 rows affected (0.49 sec)

mysql> SHOW CREATE TABLE CityLike\G
***** 1. row *****
    Table: CityLike
Create Table: CREATE TABLE `CityLike` (
  `ID` int(11) NOT NULL AUTO_INCREMENT,
  `Name` char(35) NOT NULL DEFAULT '',
  `CountryCode` char(3) NOT NULL
DEFAULT '',
  `District` char(20) NOT NULL DEFAULT
 '',
  `Population` int(11) NOT NULL DEFAULT
'0',
  PRIMARY KEY (`ID`),
  KEY `CountryCode` (`CountryCode`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1
1 row in set (0.00 sec)
```

# Using Temporary Tables

---

- Use temporary tables to improve the efficiency of queries that repeatedly access large amounts of data in multiple tables. A temporary table:
  - Is created by the CREATE TEMPORARY TABLE statement
  - Is visible only to the client that created and kept it
  - Is kept only until the client disconnects
  - Does not conflict with other clients using the same table name
  - Overrides an existing table until the temporary table is dropped
  - Does not appear in SHOW TABLES or INFORMATION\_SCHEMA.TABLES.



# Example: Creating a Temporary Table

The following SQL creates a temporary table from the existing City table. The temporary table contains the names of all the cities in the state of Texas (USA).

```
mysql> CREATE TEMPORARY TABLE Texas  
-> SELECT Name FROM City WHERE District='Texas';  
Query OK, 26 rows affected (0.01 sec)  
Records: 26  Duplicates: 0  Warnings: 0
```

# Modifying Table Data

# Inserting Data in a Table

- Populate a table with row data by using the

```
INSERT INTO table_name (<column_list>)  
VALUES (<value_list>)
```

- **table\_name**: Table to which you add row data
- **column\_list**: Column names from the specified table
- **value\_list**: Value expressions for columns
- Example:

```
INSERT INTO City (ID, Name, CountryCode)  
VALUES (NULL, 'Essaouira', 'MAR'),  
        (NULL, 'Sankt-Augustin', 'DEU');
```

# Inserting Data in a Table: INSERT with SET

- Use `INSERT` with the `SET` option to indicate column names and values. For example:

```
INSERT INTO City SET ID=NULL, Name='Essaouira',
    CountryCode='MAR';

INSERT INTO City SET ID=NULL,
    Name='Sankt-Augustin', CountryCode='DEU';
```

- Content of the new rows:

ID	Name	CountryCode	District	Population
4080	Essaouira	MAR		0
4081	Sankt-Augustin	DEU		0

# Inserting Data in a Table: INSERT with SELECT

- Use `INSERT` with the `SELECT` option to copy rows from an existing table, or (temporarily) to store a

```
INSERT INTO table_name (<column_list>)
(<query_expression>)
```

- **table\_name and column\_list:** Work the same way as they do for `INSERT...VALUES` statements
- **query\_expression:** Must produce exactly as many columns as specified by `column_list`
- Example:

```
INSERT INTO Top10Cities (ID, Name, CountryCode)
SELECT ID, Name, CountryCode FROM City
ORDER BY Population DESC LIMIT 10;
```

# Retrieving the AUTO\_INCREMENT Value of the Last Insertion

- The MySQL-specific `SELECT LAST_INSERT_ID()` statement retrieves the last AUTO\_INCREMENT value.

```
mysql> INSERT INTO City (Name, CountryCode)
-> VALUES ('Sarah City', 'USA');
```

```
Query OK, 1 row affected (0.0 sec)
```

```
mysql> SELECT LAST_INSERT_ID();
+-----+
| LAST_INSERT_ID() |
+-----+
| 4080           |
+-----+
```

# Retrieving the ID of the Last Insert by Using PHP/PDO

- Use the PDO database connection object's `lastInsertId()` method to retrieve the most recent AUTO\_INCREMENT value.

```
$dbh->exec ("INSERT INTO City (Name, CountryCode)  
VALUES ('Waterford', 'IRL')");  
  
$last_id = $dbh->lastInsertId();
```

# Retrieving the ID of the Last Insert by Using Connector/J

- Using the Statement or PreparedStatement object:
  1. Include the Statement.RETURN\_GENERATED\_KEYS argument in the call to executeUpdate()
- Example (using Statement):

```
Statement s = conn.createStatement ();  
  
s.executeUpdate ("INSERT INTO City (Name, CountryCode) "  
    + "VALUES ('Waterford', 'IRL'),"  
    Statement.RETURN_GENERATED_KEYS);
```

# Retrieving the ID of the Last Insert by Using Connector/J

2. Call `getGeneratedKeys()` to retrieve the sequence value
  - Example (using Statement):

```
long last_id;
ResultSet rs = s.getGeneratedKeys();
if (rs.next ())
{
    last_id = rs.getLong (1);
}
else
{
    throw new SQLException
        ("getGeneratedKeys() produced no value");
}
rs.close ();
s.close ();
```

# Retrieving the Last Insert ID by Using Connector/Python

- Use the `lastrowid` cursor object attribute.
- Example:

```
cursor = conn.cursor()  
  
cursor.execute('''  
  
    INSERT INTO City (Name, CountryCode)  
  
    VALUES ('Waterford', 'IRL')  
  
    ''')  
  
last_id = cursor.lastrowid
```

# DELETE Statement

- Use the DELETE statement to remove entire rows:

```
DELETE FROM table_name  
[WHERE condition]  
[ORDER BY ...]  
[LIMIT row_count]
```

- Filter the rows to be deleted by using the WHERE clause:

```
DELETE FROM CountryLanguage WHERE IsOfficial='F';
```

- Omit the WHERE clause to delete all rows from the table.
- Use DELETE with extreme caution because it does not have the “undo” feature.
  - Consider setting the --safe-updates option.



# DELETE with ORDER BY and LIMIT: Descending

- Control the order and count of DELETE by using ORDER BY and LIMIT.
- Example using rows sorted in descending order:

```
DELETE FROM CountryLanguage  
WHERE Language = 'MySQL'  
ORDER BY CountryCode DESC  
LIMIT 1;
```

- Effect on the CountryLanguage table:

CountryCode	Language	IsOfficial	Percentage	
BEL	MySQL	F	0.0	
FJI	MySQL	F	0.0	
GRL	MySQL	F	0.0	Removed

# DELETE with ORDER BY and LIMIT: Ascending

- Example using rows sorted in ascending order:

```
DELETE FROM CountryLanguage  
WHERE Language = 'MySQL'  
ORDER BY CountryCode ASC  
LIMIT 1;
```

- Effect on the CountryLanguage table:

CountryCode	Language	IsOfficial	Percentage	Removed
BEL	MySQL	F	0.0	
FJI	MySQL	F	0.0	
GRL	MySQL	F	0.0	

# UPDATE Statement

- Modifies the contents of existing rows
- Is used with the **SET** clause for column assignments
- General syntax:

```
UPDATE table_name SET column=expression  
[ , column=expression, ... ]  
[WHERE condition] [other_clauses]
```

- Example:

```
mysql> UPDATE Country  
      -> SET Population = Population * 2,  
      ->           Region = 'Dolphin Country'  
      -> WHERE Code = 'SWE';  
Query OK, 1 row affected (0.03 sec)  
Rows matched: 1    Changed: 1   Warnings: 0
```

# UPDATE Statement: Column Assignment Effects

---

- Effects are subject to column constraints:
  - When you try to update a column to a value that does not match the column definition, the value is converted or truncated by the server.
- UPDATE has no effect under the following conditions:
  - When it matches no rows for updating:
    - Due to an empty table
    - If no rows match the WHERE clause
  - When it does not actually *change* any column values:
    - When the given value is the same as the existing value
    - If there is no row in the table that contains the specified key values



# UPDATE Statement: Order of Affected Rows

- There is no guarantee of the order in which rows are updated. This can result in errors such as the following:

```
mysql> UPDATE City SET ID = ID+1;  
ERROR 1062 (23000): Duplicate entry '2' for key 'PRIMARY'
```

- Use ORDER BY to control the order. For example:

```
mysql> UPDATE City SET ID = ID+1  
      -> ORDER BY ID DESC;  
Query OK, 4079 rows affected (0.13 sec)  
Rows matched: 4079 Changed: 4079 Warnings: 0
```

- Table contents after update (ordered by ID):

ID	Name	CountryCode	District	Population
2	Kabul	AFG	Kabol	1780000
3	Qandahar	AFG	Qandahar	237500
4	Herat	AFG	Herat	186800
...				

# UPDATE with LIMIT

- Use `LIMIT` to control the number of rows updated. For example:

```
mysql> UPDATE City SET ID = ID-1 LIMIT 1;
```

- Contents of the affected row:

ID	Name	CountryCode	District	Population
1	Kabul	AFG	Kabol	1780000
3	Qandahar	AFG	Qandahar	237500
4	Herat	AFG	Herat	186800
...				

Note that only the first `ID` was affected by the update. The other IDs kept their existing values.

- Combine `ORDER BY` and `LIMIT` for more control over updates.

# REPLACE Statement

- Performs the same as `INSERT` *except* that when an old row in the table has the same value as a new row for a `PRIMARY KEY` or `UNIQUE` constraint, the old row is deleted before the new row is inserted.
- Is a MySQL extension to the SQL standard
- General syntax:

```
REPLACE INTO table_name (<column_list>)  
VALUES (<value_list>)
```

- Example:

```
REPLACE INTO CountryLanguage (CountryCode, Language,  
                           Percentage)  
VALUES ('ALB', 'Albanian', 78.1);
```

# REPLACE Results

---

- Example:

```
Query OK, 2 rows affected (0.06 sec)
```

- Returns a count to indicate the number of rows affected
  - The count is the sum of the rows deleted and inserted.
  - Count of 1: One row was inserted and no rows were deleted.
  - Count greater than 1: One or more old rows were deleted before the new row was inserted.
- Row count makes it easy to determine if REPLACE only added a row or if it also replaced rows.



# REPLACE Algorithm

---

- MySQL uses the following algorithm for REPLACE:
  - 1.Try to insert the new row in the table.
  - 2.When insertion fails because of a duplicate key error (for a primary key or unique constraint), delete the conflicting row that has the duplicate key value from the table.
  - 3.Try again to insert the new row in the table.



# Insertion of Duplicate Rows

---

- Use ON DUPLICATE KEY UPDATE with INSERT to prevent a duplicate row error.
  - The behavior of ON DUPLICATE KEY UPDATE is similar to REPLACE except for the following:
    - REPLACE
      - Discards the *existing* row
      - Adds the *new* row to the table
    - ON DUPLICATE KEY UPDATE
      - Discards the *new* row
      - Preserves the *existing* row
  - This option is specific to MySQL.



# Removing Table Rows by Truncating

- Use the TRUNCATE TABLE (or just TRUNCATE) statement to remove **all** rows from a table.
  - General syntax:

```
TRUNCATE [TABLE] <table_name>;
```

- Comparing DELETE to TRUNCATE TABLE:

DELETE	TRUNCATE TABLE
Can delete specific rows with WHERE	Deletes <i>all</i> rows; no filtering is possible
Is usually slower than TRUNCATE TABLE	Is usually faster than DELETE
Returns a true row count	Can return a row count of zero
Can be rolled back as part of a transaction	Cannot be rolled back in a transaction
Does not reset the AUTO_INCREMENT sequence	Can reset the sequence of AUTO_INCREMENT values

# Transactions

# What Is a Transaction?

---

- A set of data manipulation tasks treated as a single unit of work
- Takes the database from one consistent state to another consistent state

## Nontransactional Process

- 1 Remove \$1000 from account #10001
- 2 Write to database
- 3 Deposit \$1000 into account #10243
- 4 Write to database

## Transactional Process

- 1 Remove \$1000 from account #10001
- 2 Deposit \$1000 into account #10243
- 3 Write to database



# Handling Partial Operations

---

- All of the data manipulation steps must be carried out successfully. If any step fails, the system must:
  - Permanently retain those operations that succeeded completely
  - Disregard those operations containing steps that did not succeed



# The ACID Properties of Transactions

---

- Transactional systems are **ACID** compliant, where “ACID” stands for:
  - **Atomic:** All the statements execute successfully or are canceled as a unit. A transaction is “all-or-nothing.”
  - **Consistent:** A transaction changes the database from one consistent data state to another.
  - **Isolated:** One transaction does not affect another. Each transaction locks the results of its individual steps until the whole transaction ends.
  - **Durable:** All the changes made by a transaction that ends successfully are recorded correctly in the database. Changes are not lost.



# Transaction Control Statements

---

- Use the following statements to control transactions explicitly:
  - START TRANSACTION (or BEGIN): Commences a new transaction
  - COMMIT: Commits the current transaction, making its changes permanent
  - ROLLBACK: Rolls back the current transaction, canceling its changes
  - SET AUTOCOMMIT: Disables or enables the default autocommit mode for the current connection



# autocommit Mode

---

- autocommit mode determines how and when new transactions are started.
  - Enabled (default):
    - A single SQL statement implicitly starts a new transaction.
    - All changes to a table take effect immediately.
    - You can execute multi-statement transactions using the START TRANSACTION statement, which temporarily disables autocommit until you execute COMMIT or ROLLBACK.
  - Disabled:
    - Transactions span multiple statements by default.
    - You must explicitly COMMIT a transaction to persist its effect or ROLLBACK to cancel it.



# Controlling autocommit Mode

---

- The server variable `autocommit` specifies whether autocommit mode is enabled.
  - By default, autocommit is enabled for new sessions.
  - Disable autocommit in the current session with:

```
mysql> SET autocommit = OFF;
```

- Disable autocommit globally in an options file:

```
[mysqld]  
autocommit=0
```

- Disable autocommit globally with a mysqld startup

```
# mysqld --autocommit=0
```



# Statements Causing an Implicit COMMIT

---

- The COMMIT statement explicitly persists the effects of the current transaction to the database.
- Many statements commit transactions implicitly, including:
  - Transaction control statements such as START TRANSACTION and SET autocommit = ON;
  - Data Definition Language statements (DDL), such as ALTER, CREATE, DROP, and TRUNCATE TABLE.
  - Data access and user management statements, such as GRANT, REVOKE, SET PASSWORD
  - Locking statements, such as LOCK TABLES, UNLOCK TABLES
- Statements that cause an implicit COMMIT cannot be rolled back.



# General Approach to Coding Transactional Statements

---

- All MySQL APIs support transactions.
  - Many APIs provide an abstracted interface for working with transactions.
    - You can see the resulting SQL statements by enabling and then examining the general query log.
  - With others, you must construct the transactional SQL statements yourself.
- Within the code that executes your transaction:
  - Use your programming language's ability to handle exceptions, if it exists
    - Execute the transaction statements in a block.
    - Raise an exception and roll back if statements in the block fail.
  - Keep the amount of tests required to determine if a rollback is needed to a minimum.



# Coding Transactional Statements in PHP/PDO

```
01 try
02 {
03     $dbh->beginTransaction ();
04     $dbh->exec ("UPDATE Country SET GNP = (GNP * 1.1) WHERE Name =
05
06         'Armenia'");
07
08     $dbh->exec ("UPDATE Country SET GNP = (GNP * 1.1) WHERE Name =
09
10         'Bermuda'");
11
12     $dbh->commit ();
13 }
14 catch (Exception $e)
15 {
16     print ("Transaction failed, rolling back. Error was:\n");
17     print ($e->getMessage () . "\n");
18
19     try
20     {
21         $dbh->rollback ();
22     }
23     catch (Exception $er) {
24         print ("Unable to roll back. Error was:\n");
25         print ($er->getMessage () . "\n");
26     }
27 }
```

# Coding Transactional Statements in Connector/J

```
01 try {
02     conn.setAutoCommit (false);
03     Statement s = conn.createStatement ();
04     s.executeUpdate ("UPDATE Country SET GNP = (GNP * 1.1) WHERE Name
05
06         = 'Armenia'");
07     s.executeUpdate ("UPDATE Country SET GNP = (GNP * 1.1) WHERE Name
08
09         = 'Bermuda'");
10 }
11 catch (SQLException e) {
12     System.err.println ("Transaction failed, rolling back. Error: ");
13     printSQLErrorMessage (e); // custom function
14     try {
15         conn.rollback ();
16         conn.setAutoCommit (true);
17     }
18     catch (SQLException er) {
19         System.err.println ("Roll back failed. Error: ");
20         printSQLErrorMessage (er); // custom function
21 }
```

# Coding Transactional Statements in Connector/Python

```
01  try:
02      cursor = conn.cursor()
03      cursor.execute("UPDATE Country SET GNP = (GNP * 1.1)
04                           WHERE
05                           Name = 'Armenia'")
06      cursor.execute("UPDATE Country SET GNP = (GNP * 1.1)
07                           WHERE           Name = 'Bermuda'")
08      cursor.close()
09      conn.commit()
10  except mysql.connector.Error as e:
11      print("Transaction failed, rolling back. Error was:")
12      print(e)
13  try:
14      conn.rollback()
15  except mysql.connector.Error as er:
16      print("Unable to roll back. Error was:")
17      print(er)
```

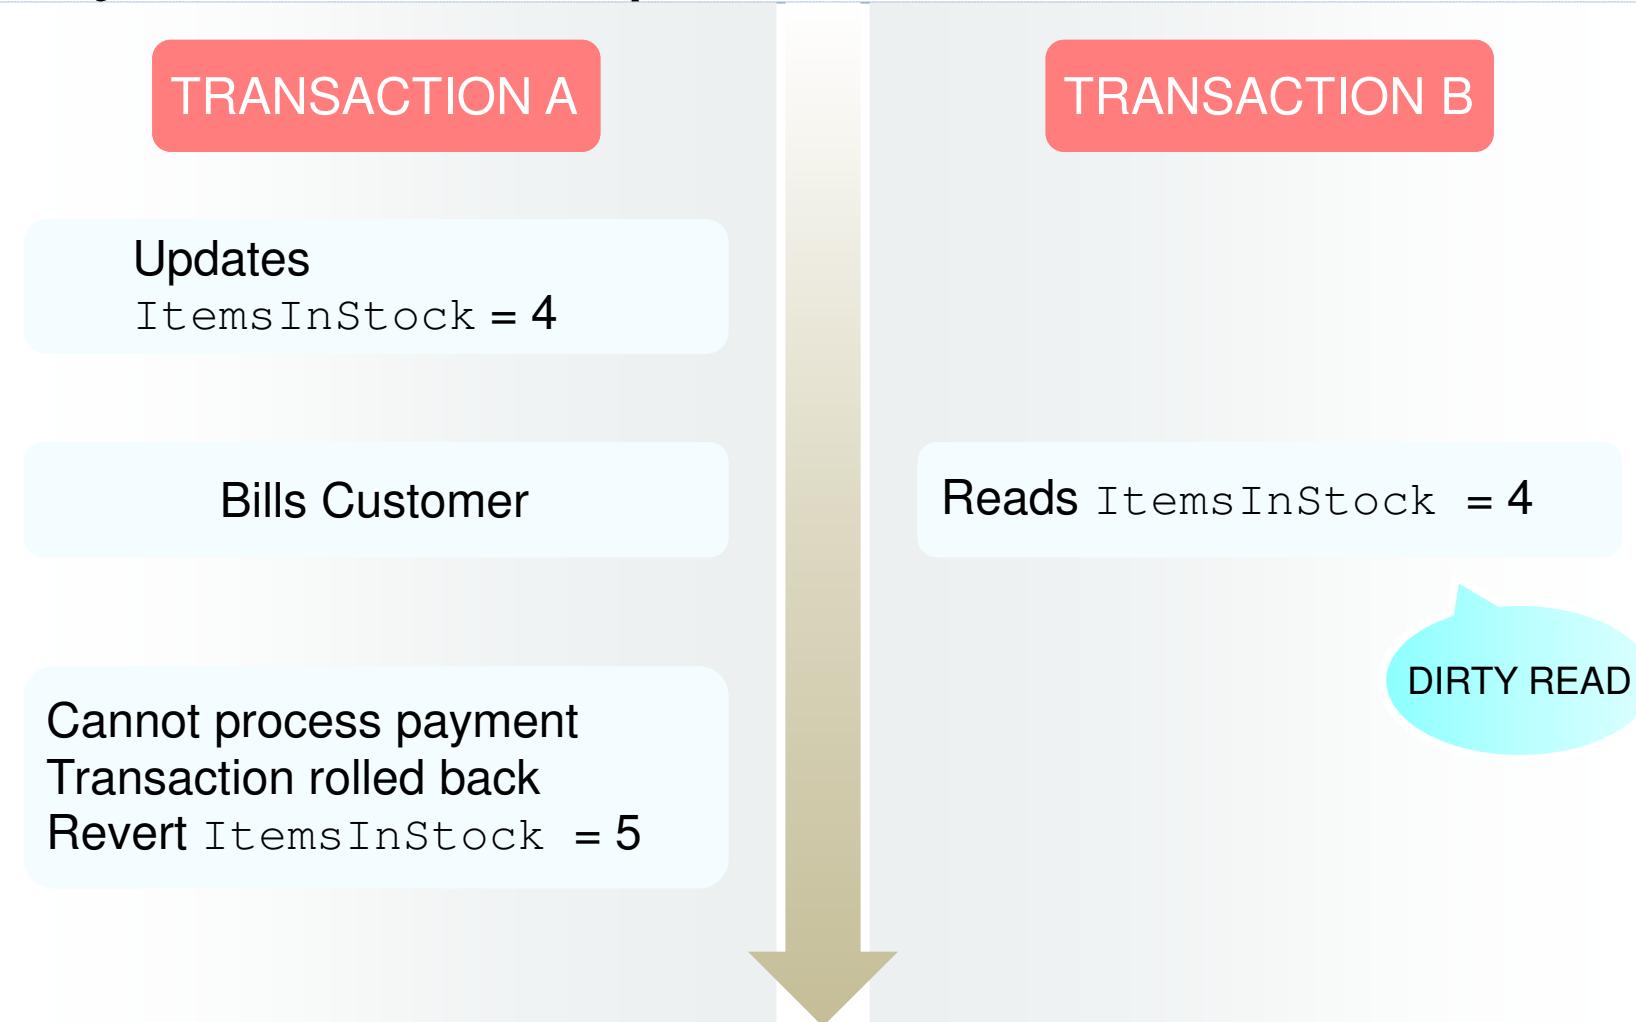
# Consistency Issues

---

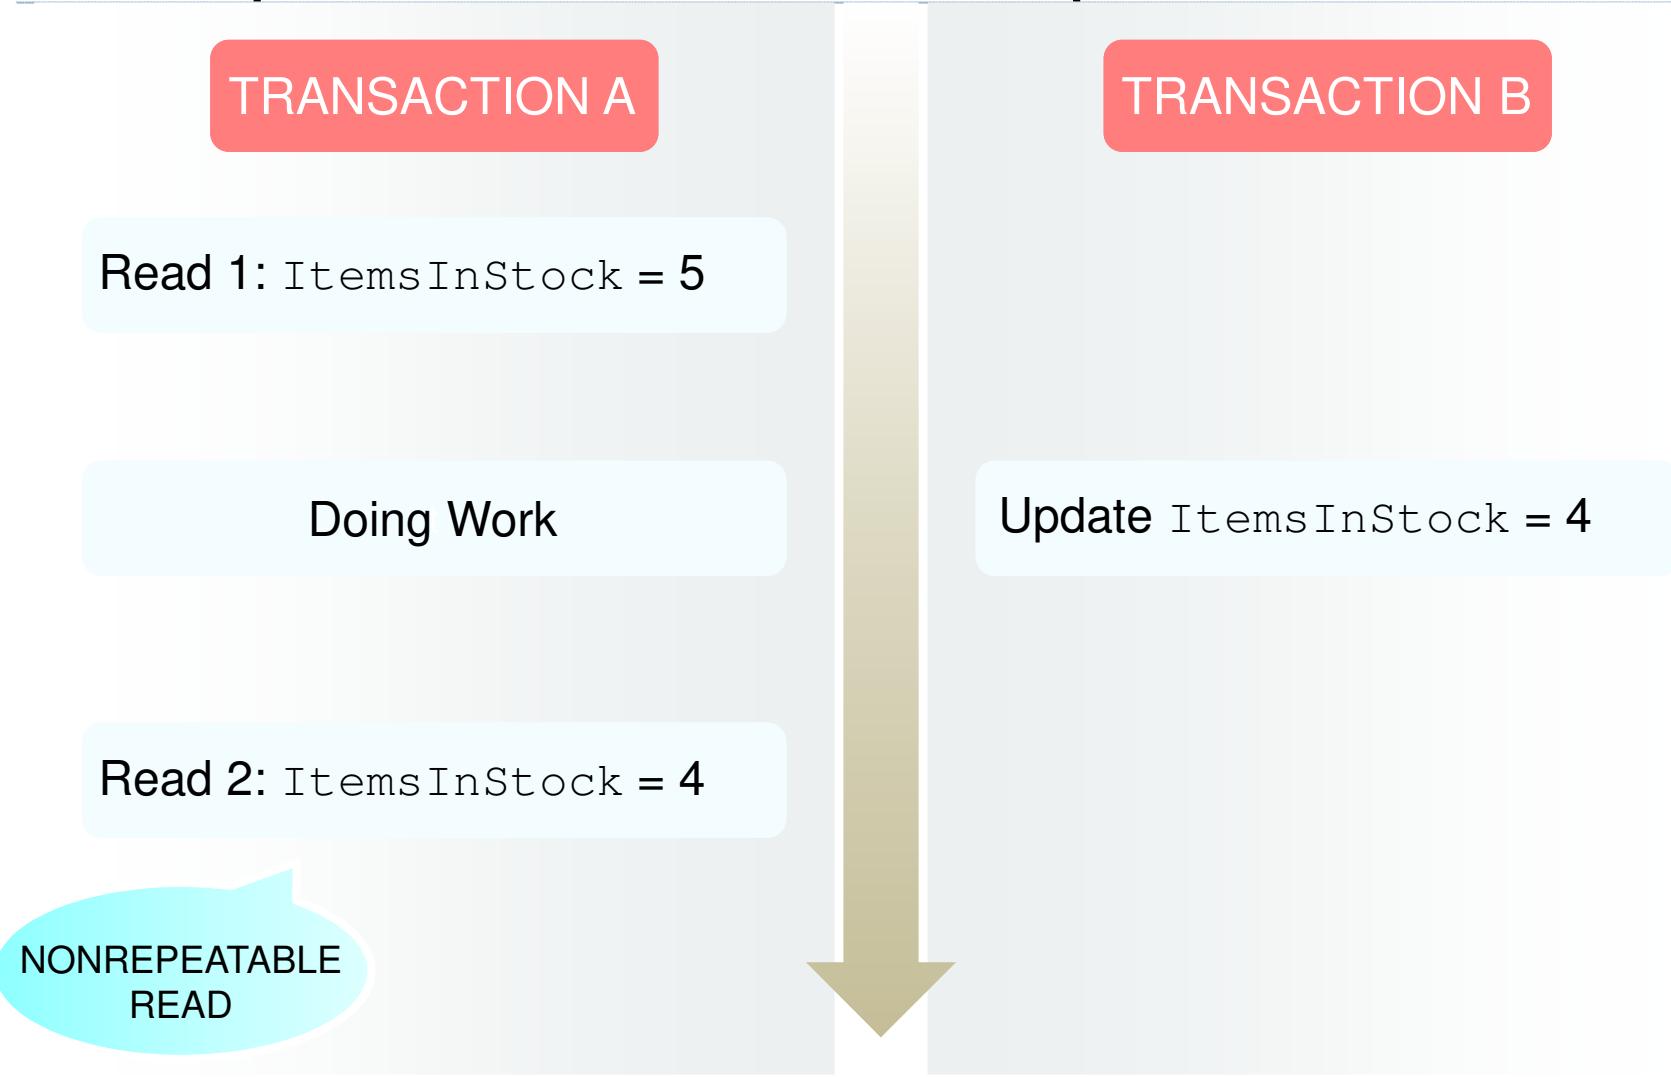
- When multiple clients concurrently access data from the same table, the following consistency issues can occur:
  - **Dirty reads:** A dirty read occurs when a transaction reads the changes made by another uncommitted transaction.
  - **Nonrepeatable reads:** A nonrepeatable read occurs when the same read operation yields different results when it is repeated at a later time within the same transaction.
  - **Phantom reads:** A phantom is a row that appears that was not previously visible within the same transaction.



# Dirty Read: Example



# Nonrepeatable Read: Example



# Phantom Read: Example

TRANSACTION A

Read 1:

```
SELECT * FROM inventory  
WHERE type = 'ABC123';
```

Output: 3 rows

Doing Work

Read 2:

```
SELECT * FROM inventory  
WHERE type = 'ABC123';
```

Output: 4 rows

PHANTOM READ

TRANSACTION B

Insert new item into  
Inventory of type 'ABC123'

# Isolation Levels

---

- **Isolation levels** control how changes made by transactions affect other transactions happening at the same time. This is made possible by *multiversioning*.
  - READ UNCOMMITTED allows a transaction to see uncommitted changes made by other transactions.
  - READ COMMITTED allows a transaction to see changes made by other transactions only if they have been committed.
  - REPEATABLE READ (the default) ensures that if a transaction issues the same SELECT query twice, it gets the same result both times, regardless of committed or uncommitted changes made by other transactions.
  - SERIALIZABLE completely isolates the effects of one transaction from another.



# Resolving Consistency Issues with Isolation Levels

---

- The following table displays which consistency issues are possible with each isolation level.

Isolation Level	Dirty reads	Nonrepeatable reads	Phantom reads
READ UNCOMMITTED	Possible	Possible	Possible
READ COMMITTED	Not possible	Possible	Possible
REPEATABLE READ*	Not possible	Not possible	Not possible
SERIALIZABLE	Not possible	Not possible	Not possible

- \* The default isolation level is REPEATABLE READ.



# Setting Isolation Levels

---

- The transaction isolation level is a server option you can set at startup, or as a session or global option.

```
[mysqld]
transaction-isolation = READ-COMMITTED
```

```
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
```

```
SET GLOBAL TRANSACTION ISOLATION LEVEL READ COMMITTED;
```

startup:



# Displaying the Current Isolation Level

- Examine the value of the `tx_isolation` server variable
  - `@@session.tx_isolation` or  
`@@session.tx_isolation`: The isolation level for the current connection
  - `@@global.tx_isolation`: The isolation level for all connections to the server

```
mysql> SELECT @@tx_isolation;
+-----+
| @@tx_isolation |
+-----+
| REPEATABLE-READ |
+-----+
1 row in set (#.## sec)
```

# Locking

---

- Locking prevents problems that can occur when two or more clients access the same data at the same time.
  - The lock allows access to data by the client that holds the lock, but limits what other clients can do with that data.
  - MySQL supports two types of lock:
    - **Shared:** If a client wants to read data, other clients that want to read the same data do not cause a conflict, and they all can read at the same time. However, another client that wants to write (modify) data must wait until the read has finished.
    - **Exclusive:** If a client wants to write data, all other clients must wait until the write has finished before reading or writing.



# Locking Reads

---

- MySQL has two locking modifiers that you can add to the end of SELECT statements:
  - LOCK IN SHARE MODE: This clause locks each selected row with a shared lock. This means that no other transactions can take exclusive locks, but other transactions can use shared locks. Because read operations do not lock rows, they are not affected by the locks.
  - FOR UPDATE: This clause locks each selected row with an exclusive lock. It prevents others from acquiring any lock on the rows, but allows reading the rows.



# SELECT...LOCK IN SHARE MODE: Example

---

- **Problem:** Adding a new Australian city in the world database:
  1. You issue a SELECT query to verify that the country code for Australia (AUS) exists in the Country table.
  2. You insert a new row in the City table with a country code of AUS.
  3. Another user deletes AUS from the country table before your new City row is inserted.
  4. Your City row is now orphaned, because its parent country no longer exists.
- **Solution:** Execute the initial SELECT of the Country table using LOCK IN SHARE MODE:

```
mysql> SELECT Code FROM Country WHERE Name='Australia'  
-> LOCK IN SHARE MODE;
```



# SELECT...FOR UPDATE: Example

- **Problem:** You are manually maintaining an integer counter column in a table called CityCodes, which you use to assign a unique identifier to each city added to the City table.
  - To calculate the next available code, you read the current code and add 1 to its value.
  - In a concurrent environment, two clients could retrieve the same value for the counter simultaneously and use that value to construct the next code in the sequence, resulting in a duplicate key error.
- **Solution:** Execute the initial SELECT of the CityCodes table using FOR UPDATE:

```
mysql> SELECT counter INTO @@counter FROM CityCodes  
-> FOR UPDATE;
```

# View

# Creating Views to Reuse Queries

---

- A view is a virtual table defined by the result of a SELECT query.
  - A custom view of your tables makes some operations simpler.
  - Views provide additional benefits compared to selecting data directly from base tables.
  - You can include regular tables or other views in your defining SELECT statement.
  - Creation syntax:

```
CREATE VIEW view_name [column_list]  
AS SELECT_expression
```



# CREATE VIEW: Basic Example

This example creates a view containing the entire contents of the Country table:

The diagram illustrates the process of creating a view. On the left, a screenshot of a MySQL Workbench Query Editor window titled "Query 1" shows the SQL command to create a view:

```
1 • CREATE VIEW Country_View  
2 AS SELECT * FROM Country;
```

An orange arrow points from this window to a larger screenshot on the right, which shows the results of running the query. The "Result Grid" displays the following data:

#	Code	Name	Continent	Region
1	ABW	Aruba	North America	Caribbean
2	AFG	Afghanistan	Asia	Southern and C
3	AGO	Angola	Africa	Central Africa
4	AIA	Anguilla	North America	Caribbean

The "Action Output" section at the bottom of the right-hand screenshot shows a log entry:

Time	Action
01:33:41	SELECT * FROM Country_View

# CREATE VIEW: More Complex Example

More complex example, using computed values:

The diagram illustrates the process of creating a view. On the left, a screenshot of a MySQL Workbench Query Editor window titled "Query 1" shows the SQL command to create a view named "per\_capita\_v". The command includes a computed column "per\_capita" defined as "GNP/Population". On the right, a screenshot of the same window after execution shows the results of the "SELECT \* FROM per\_capita\_v;" query. A teal arrow labeled "Creates" points from the creation step to the resulting data grid.

Creates

```
Query 1 X
1 • CREATE VIEW per_capita_v
AS SELECT Name, GNP, Population,
GNP/Population
AS per_capita FROM Country;
```

```
Query 1 X
1 • SELECT * FROM per_capita_v;
```

#	Name	GNP	Population	per_capita
1	Aruba	828.00	103000	0.008039
2	Afghanistan	5976.00	22720000	0.000263
3	Angola	6648.00	12878000	0.000516
4	Anguilla	63.20	8000	0.007900
5	Albania	3205.00	3401200	0.000942

# Displaying View Information

- Base table statements, like `DESCRIBE`, also work for views.
- Example using `DESCRIBE view_name`:

The screenshot shows the MySQL Workbench interface with a query editor titled "Query 1" and a result grid. The query entered is `DESCRIBE per_capita_v;`. The result grid displays the structure of the view, listing four columns: Name, GNP, Population, and per\_capita, with their respective data types and properties. Below the result grid is an "Action Output" section showing the history of the query execution.

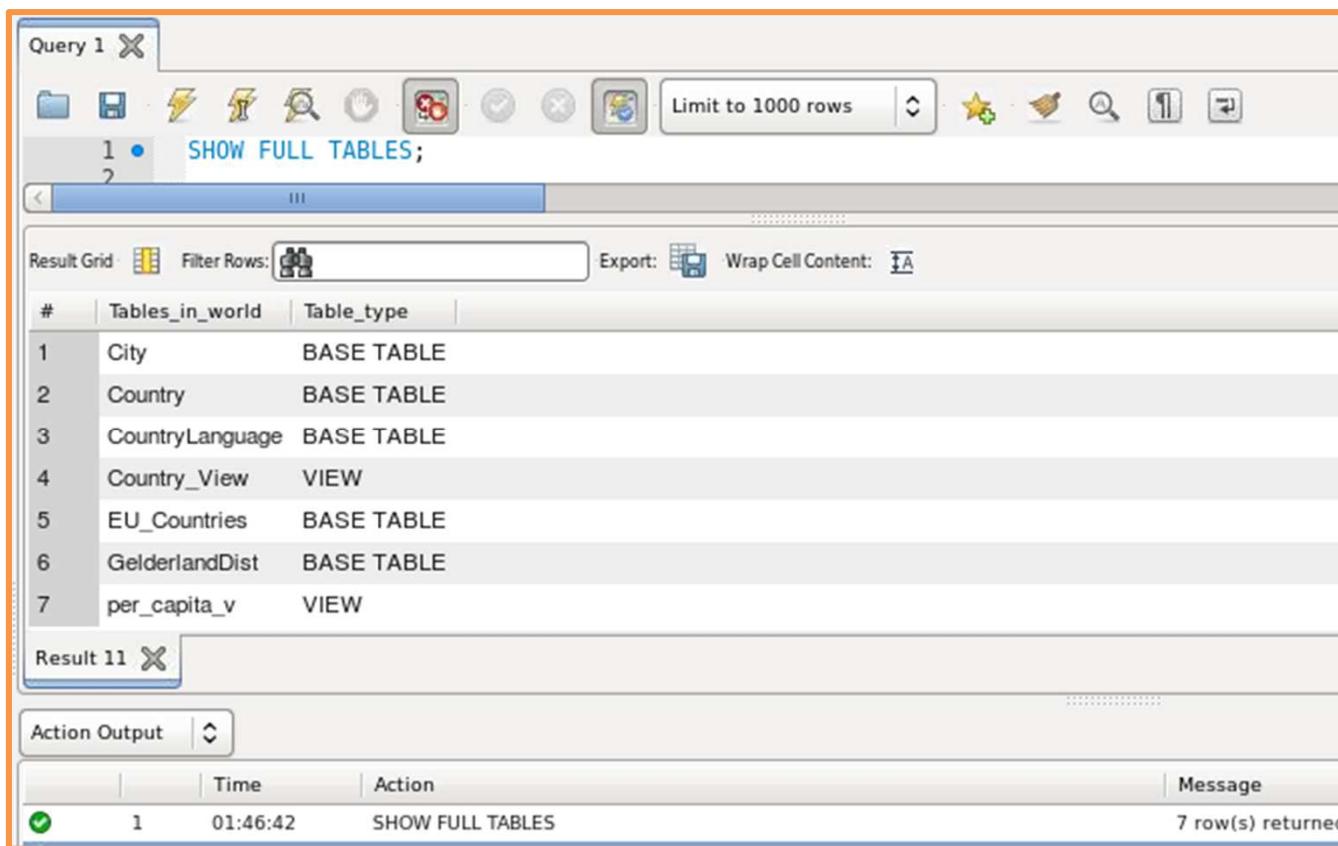
#	Field	Type	Null	Key	Default	Extra
1	Name	char(52)	NO			
2	GNP	float(10,2)	YES		NULL	
3	Population	int(11)	NO		0	
4	per_capita	double(14,6)	YES		NULL	

Action Output:

	Time	Action
1	00:35:53	DESCRIBE per_capita_v

# Views: Showing Table Types

- You can use **SHOW FULL TABLES** to differentiate base tables and views.



The screenshot shows the MySQL Workbench interface with a query window titled "Query 1". The query "SHOW FULL TABLES;" has been run, and the results are displayed in a result grid. The grid has columns for "#", "Tables\_in\_world", and "Table\_type". The results show seven entries: City, Country, CountryLanguage, Country\_View, EU\_Countries, GelderlandDist, and per\_capita\_v. All entries are listed as "BASE TABLE". Below the result grid, there is an "Action Output" section showing the query "SHOW FULL TABLES" and its execution time "01:46:42". The message "7 row(s) returned" is also present.

#	Tables_in_world	Table_type
1	City	BASE TABLE
2	Country	BASE TABLE
3	CountryLanguage	BASE TABLE
4	Country_View	VIEW
5	EU_Countries	BASE TABLE
6	GelderlandDist	BASE TABLE
7	per_capita_v	VIEW

Action	Time	Message
1	01:46:42	SHOW FULL TABLES 7 row(s) returned

# Querying Data from an Application

- Example of an application querying and displaying data:

```
# php samplephp.php
```

City	District	Population
New York	New York	8008278
Buffalo	New York	292648
Rochester	New York	219773
Yonkers	New York	196086
Syracuse	New York	147306
Albany	New York	93994

Query being executed by the application:

```
SELECT Name, District, Population FROM City  
WHERE District = 'New York'
```

# Connecting to MySQL from an Application

- PHP code to make a connection to a database:

```
$mysqli = new mysqli("localhost", "root", "MySQL5.7",
    "world");
if ($mysqli->connect_errno) {
    echo "Error connecting: ". $mysqli->connect_error;
}
```

# Executing Queries in an Application

- PHP code to execute a query:

```
$res = $mysqli->query("SELECT Name, District,  
Population FROM City WHERE District = 'New York'");  
if (! $res) {  
    echo "Query failed: " . $mysqli->connect_error;  
}
```

# Displaying Result Sets in an Application

- PHP code to display the result set:

```
while ($row = $res->fetch_assoc()) {  
    echo $row['Name'] . "\t"  
        , $row['District'], "\t"  
        . $row['Population'] . "\n";  
}
```

# Exporting Data as a Text File

---

- You can export data from a table to create a text file. Some reasons for exporting data:
  - Loading the data into another application, such as a spreadsheet, report writer, word processor, or other application that can process the data loaded from a text file
  - Copying databases from one server to another:
    - From within the same host
    - To another host
  - Testing a new MySQL release with real data
  - Transferring data from one RDBMS to another
  - Keeping a backup copy of a version of your data



# Using SELECT ... INTO OUTFILE to Create a Delimited File

- Write query results directly into a file by using **SELECT** with the **INTO OUTFILE** clause. **INTO OUTFILE** writes the results of the **SELECT** operation to the output file instead of returning them to the client.

- Syntax:

```
SELECT column_list
INTO OUTFILE '/path/filename'
[format_options]
FROM table [select_clauses];
```

- List the columns or specify \* for all columns.
- Specify the output file location on the host.
- Provide the name of the table to export.

# Exporting with a Query: INTO OUTFILE

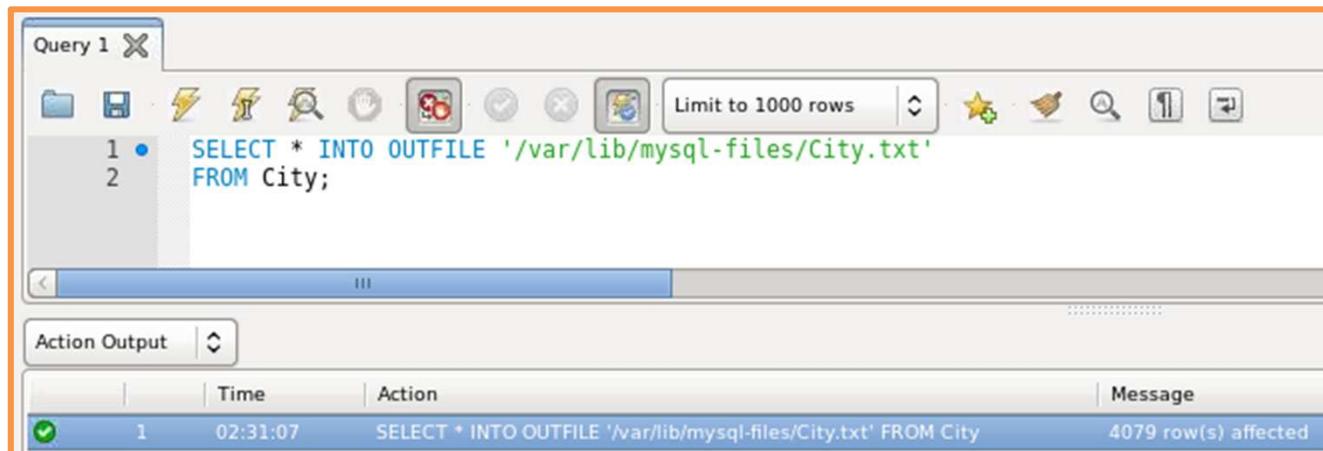
---

- The process creates a new output file. It cannot already exist.
- Each line in the output file represents one row of data.
- Default file format:
  - Columns are delimited by tabs.
  - Lines are terminated by the newline character.
  - Null values are exported as NULL.
- Change the file format with the **INTO OUTFILE** format options.
  - Column delimiters (for example, comma or semicolon)
  - Line terminators (for example, carriage return)
  - Character strings enclosed in quotation marks (either " or ')



# Example of SELECT ... INTO OUTFILE

- Example of exporting with a query:



The screenshot shows the MySQL Workbench interface with a query editor titled "Query 1". The query is:

```
1 •  SELECT * INTO OUTFILE '/var/lib/mysql-files/City.txt'  
2   FROM City;
```

The results pane below shows a single row of output:

Action Output	Time	Action	Message
1	02:31:07	SELECT * INTO OUTFILE '/var/lib/mysql-files/City.txt' FROM City	4079 row(s) affected

- Contents of the first few lines of the output file:

1	Kabul	AFG	Kabul	1780000
2	Qandahar	AFG	Qandahar	237500
3	Herat	AFG	Herat	186800
4	Mazar-e-Sharif	AFG	Balkh	
127800				
5	Amsterdam	NLD	Noord-Holland	
	731200			

# Format Options for SELECT ... INTO OUTFILE

---

- Syntax for the format options:

```
SELECT column_list
      INTO OUTFILE '/path/filename'
      FIELDS TERMINATED BY 'delimiter'
      [OPTIONALLY] ENCLOSED BY 'character'
      LINES TERMINATED BY 'terminator'
      FROM table;
```

- Replace tab with a delimiter like comma (,) or semicolon (;).
- Enclose data from each column in a character such as ' or ".
- Optionally, enclose only fields with string data types in a character.
- Change the line terminator to a different value.



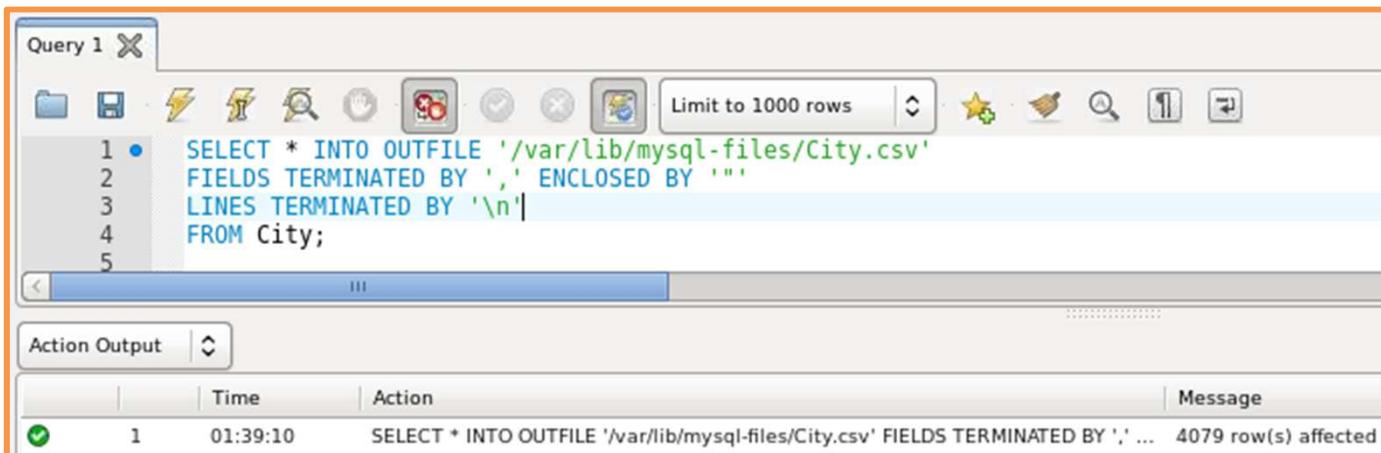
# Exporting with a Query: CSV Format

- The comma-separated values (CSV) format:
  - Separates columns with commas
  - Encloses values in double quotation marks
  - Terminates lines with newline
- Syntax with CSV format options:

```
SELECT column_list
      INTO OUTFILE '/path/filename'
      FIELDS TERMINATED BY ','
      ENCLOSED BY '"'
      LINES TERMINATED BY '\n'
      FROM table;
```

# Example of SELECT INTO OUTFILE: CSV Format

- Example of exporting in CSV Format:



The screenshot shows the MySQL Workbench interface. The top bar has tabs for 'Query 1' and 'SQL'. Below the tabs is a toolbar with various icons. The main area contains a query editor with the following SQL code:

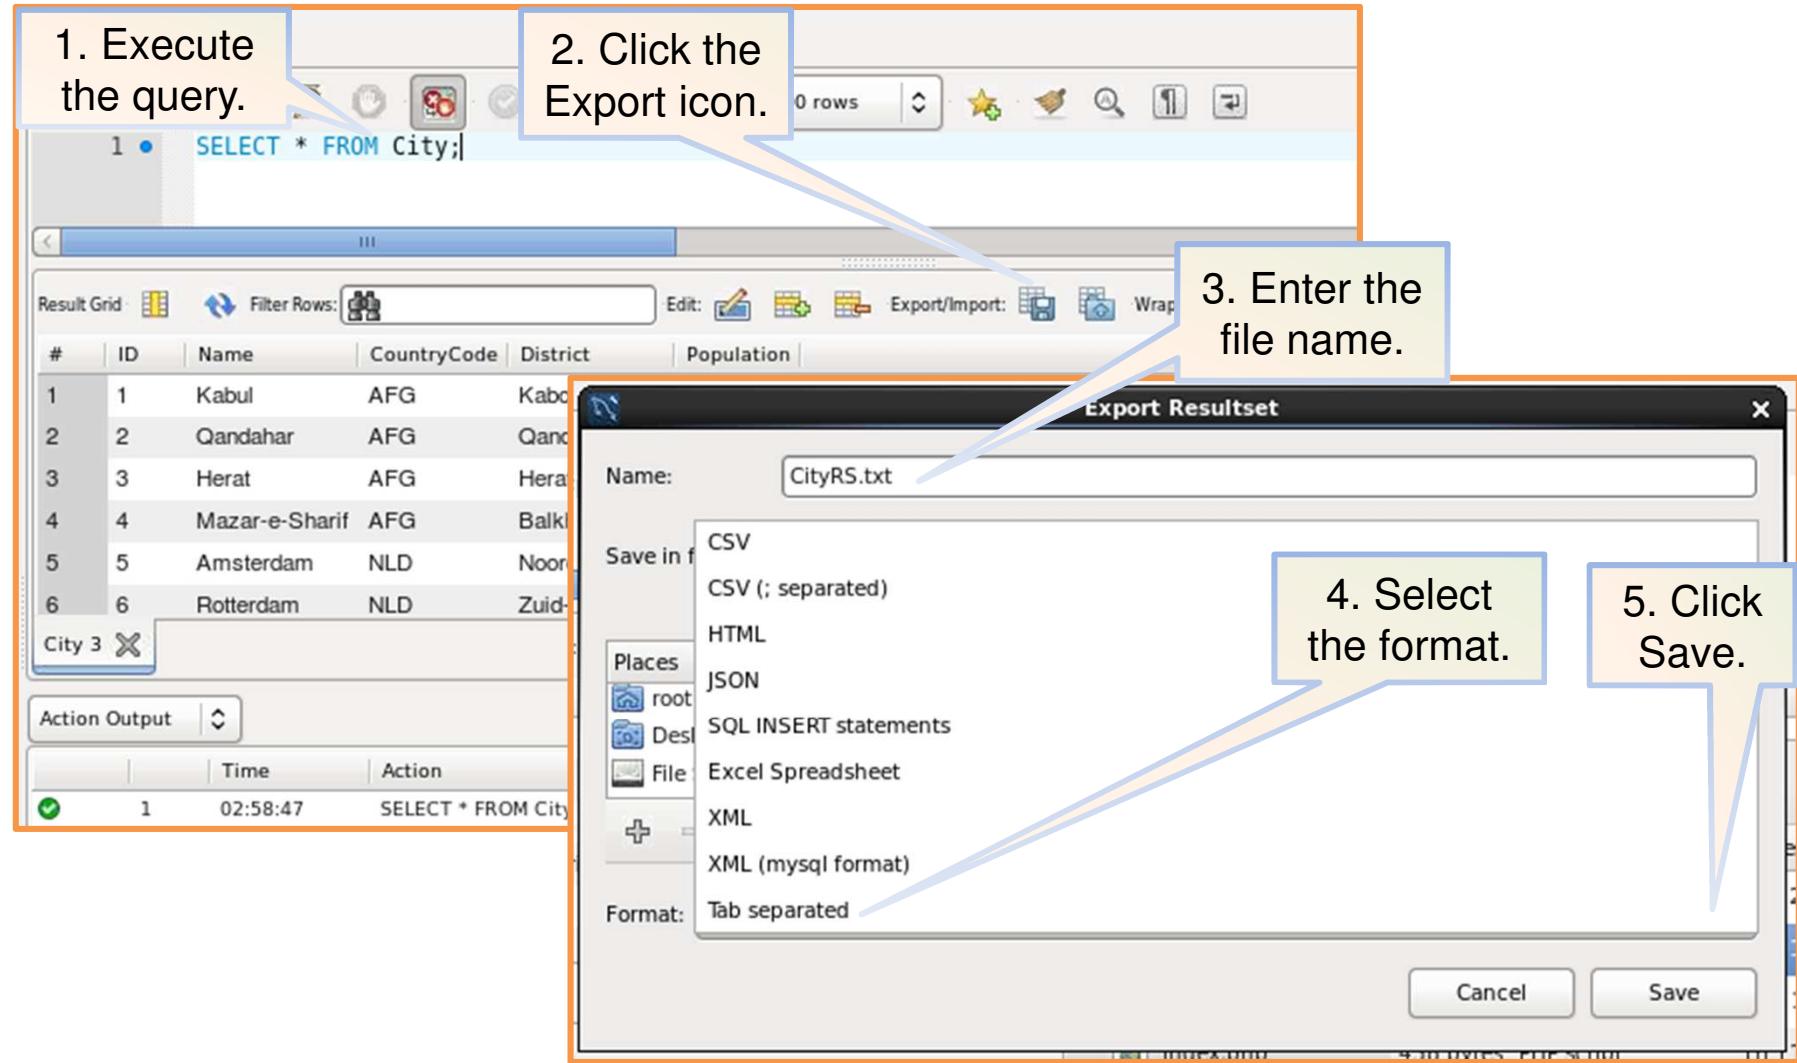
```
1 •  SELECT * INTO OUTFILE '/var/lib/mysql-files/City.csv'  
2   FIELDS TERMINATED BY ',' ENCLOSED BY ''  
3   LINES TERMINATED BY '\n'  
4   FROM City;  
5
```

Below the query editor is an 'Action Output' table with the following data:

	Time	Action	Message
1	01:39:10	SELECT * INTO OUTFILE '/var/lib/mysql-files/City.csv' FIELDS TERMINATED BY ',' ...	4079 row(s) affected

- Contents of the first few lines of the output file:
  - "1", "Kabul", "AFG", "Kabul", "1780000"
  - "2", "Qandahar", "AFG", "Qandahar", "237500"
  - "3", "Herat", "AFG", "Herat", "186800"
  - "4", "Mazar-e-Sharif", "AFG", "Balkh", "127800"

# Exporting Data from MySQL Workbench Result Set in Tab-Separated Format

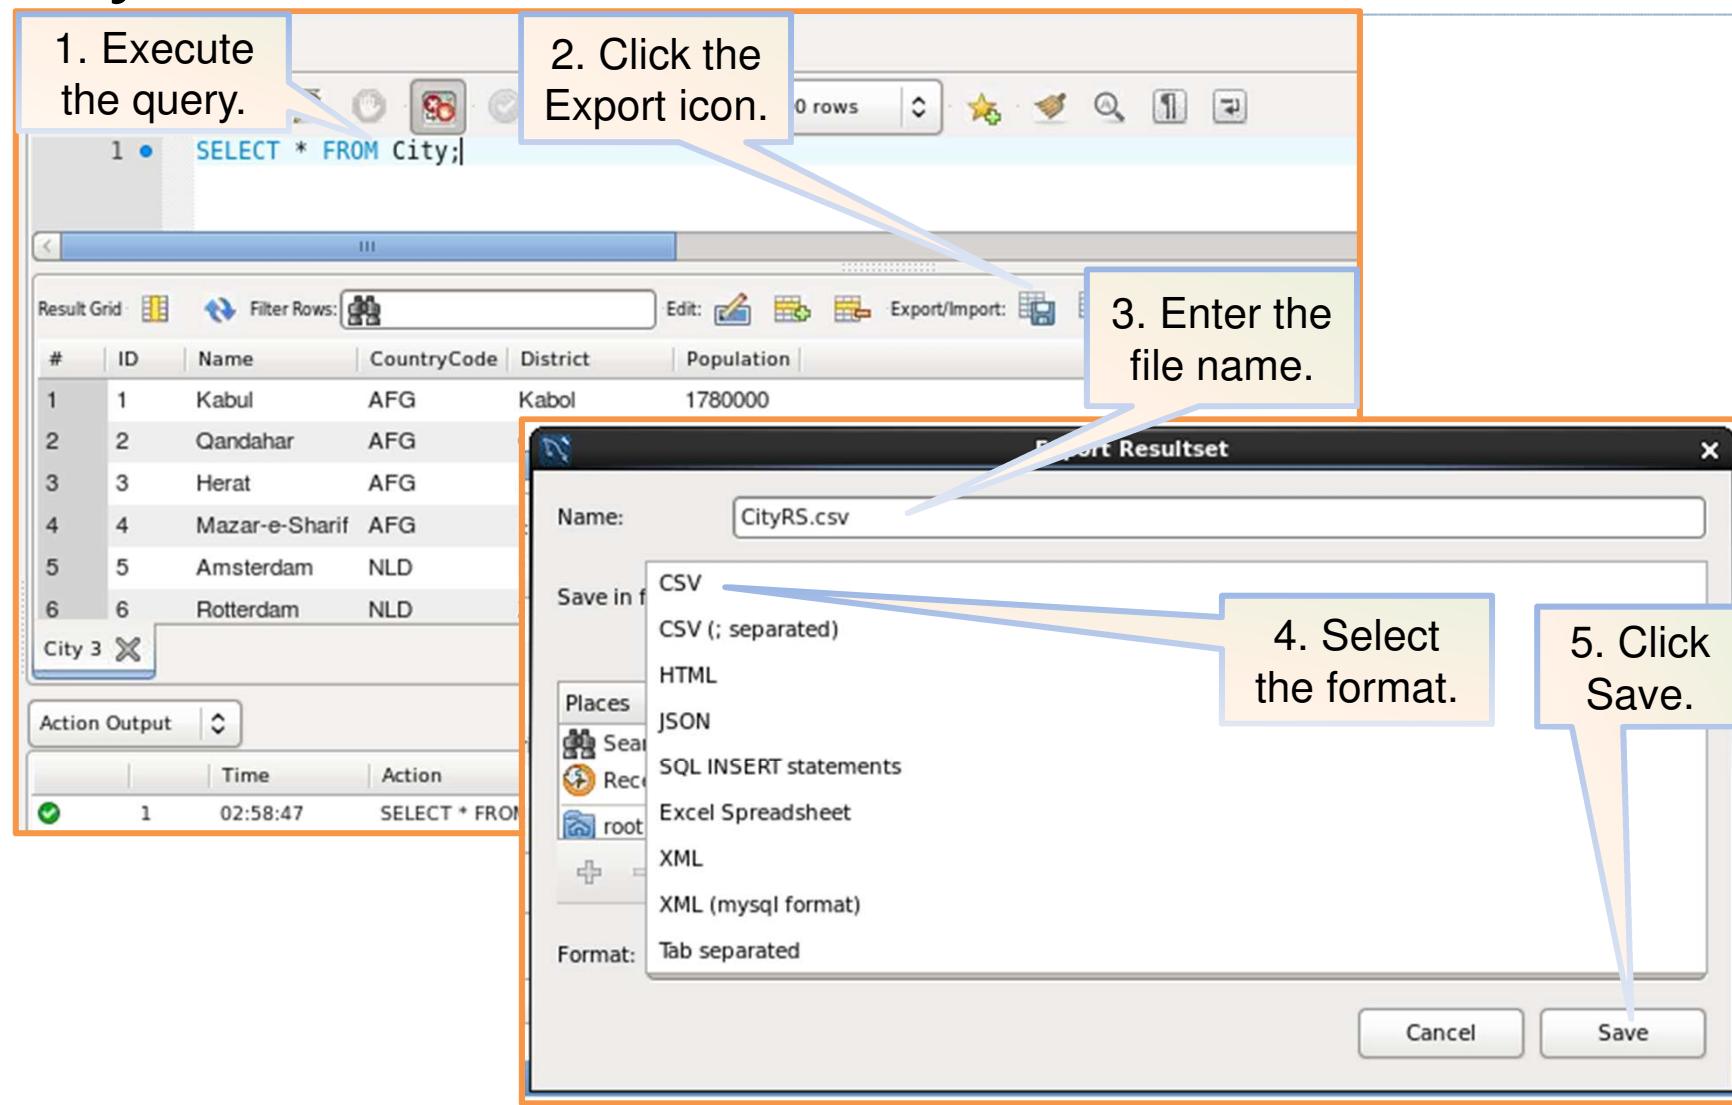


# Contents of the Tab-Separated File After Exporting the Result Set of SELECT \* FROM City;

- Contents of the output file:

ID	Name	CountryCode	District	Population
1	Kabul	AFG	Kabol	1780000
2	Qandahar	AFG	Qandahar	237500
3	Herat	AFG	Herat	186800
4	Mazar-e-Sharif	AFG	Balkh	127800
...				
•995	Probolinggo	IDN	"East Java"	120770
•996	Cilegon	IDN	"West Java"	117000
•997	Cianjur	IDN	"West Java"	114300
•998	Ciparay	IDN	"West Java"	111500
•999	Lhokseumawe	IDN	Aceh	109600
•1000	Taman	IDN	"East Java"	107000

# Exporting Data in CSV Format from the MySQL Workbench Result Set

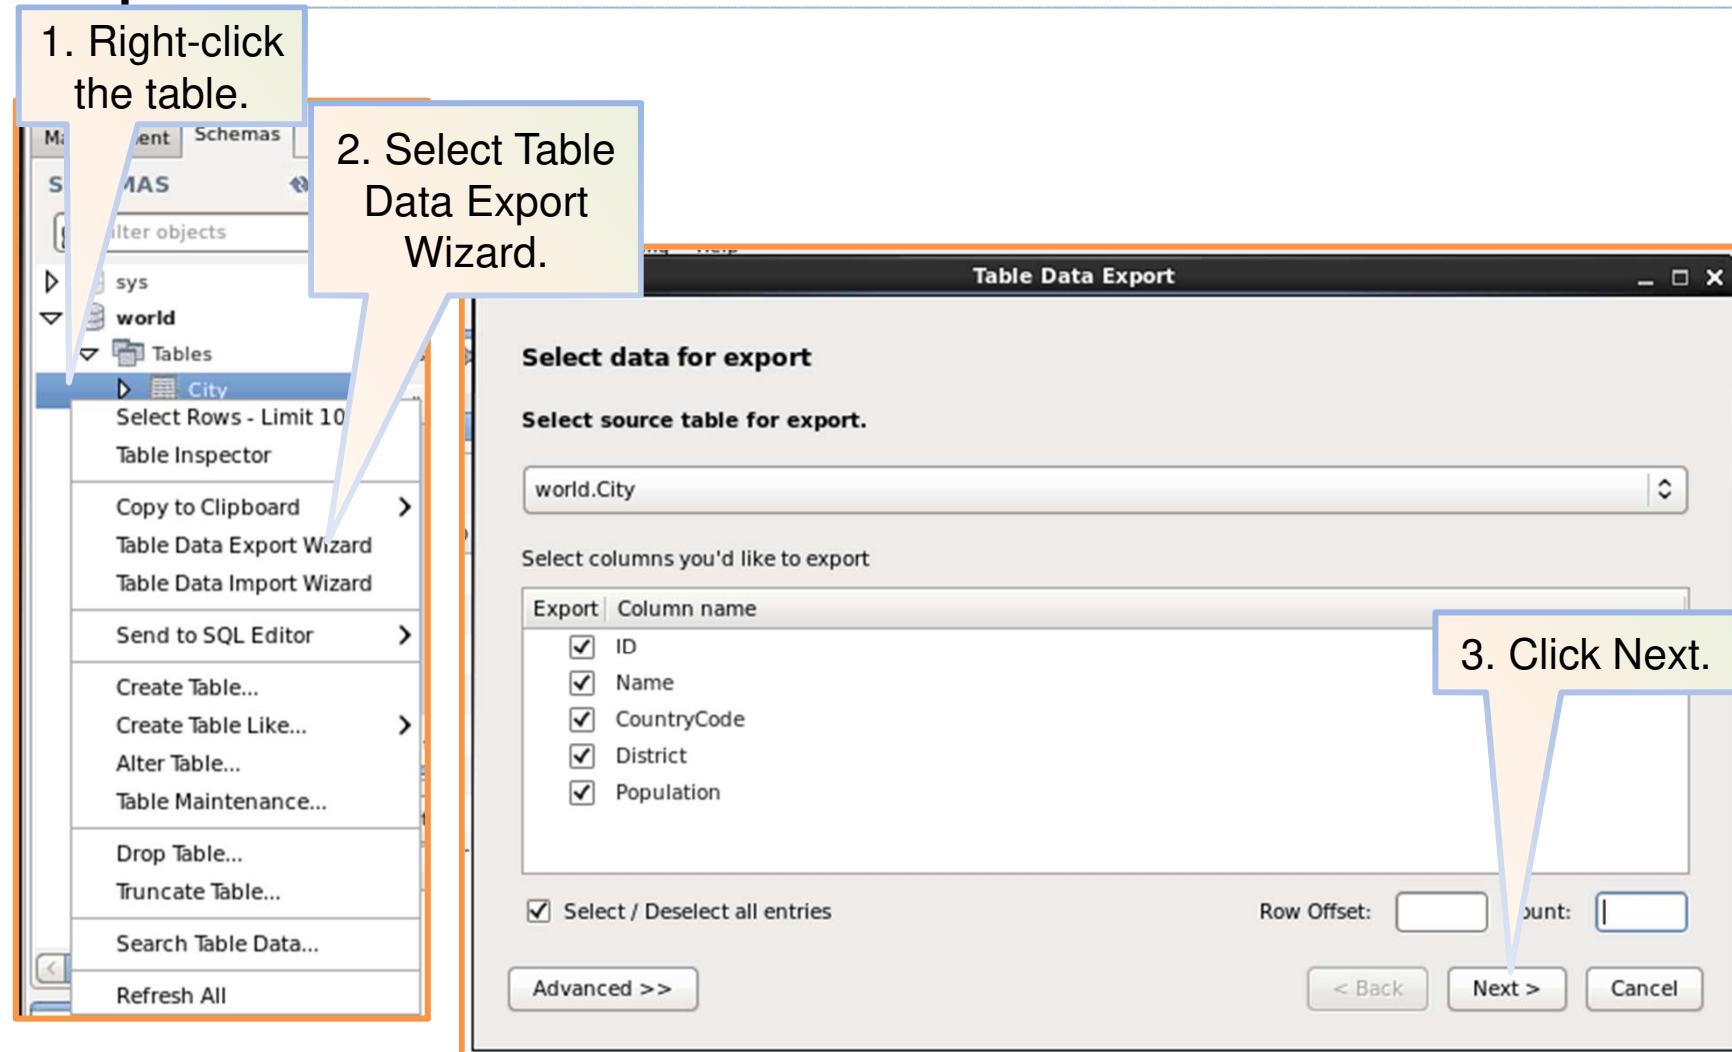


# Contents of the CSV File after Exporting the Result Set of SELECT \* FROM City;

- Contents of the output file:

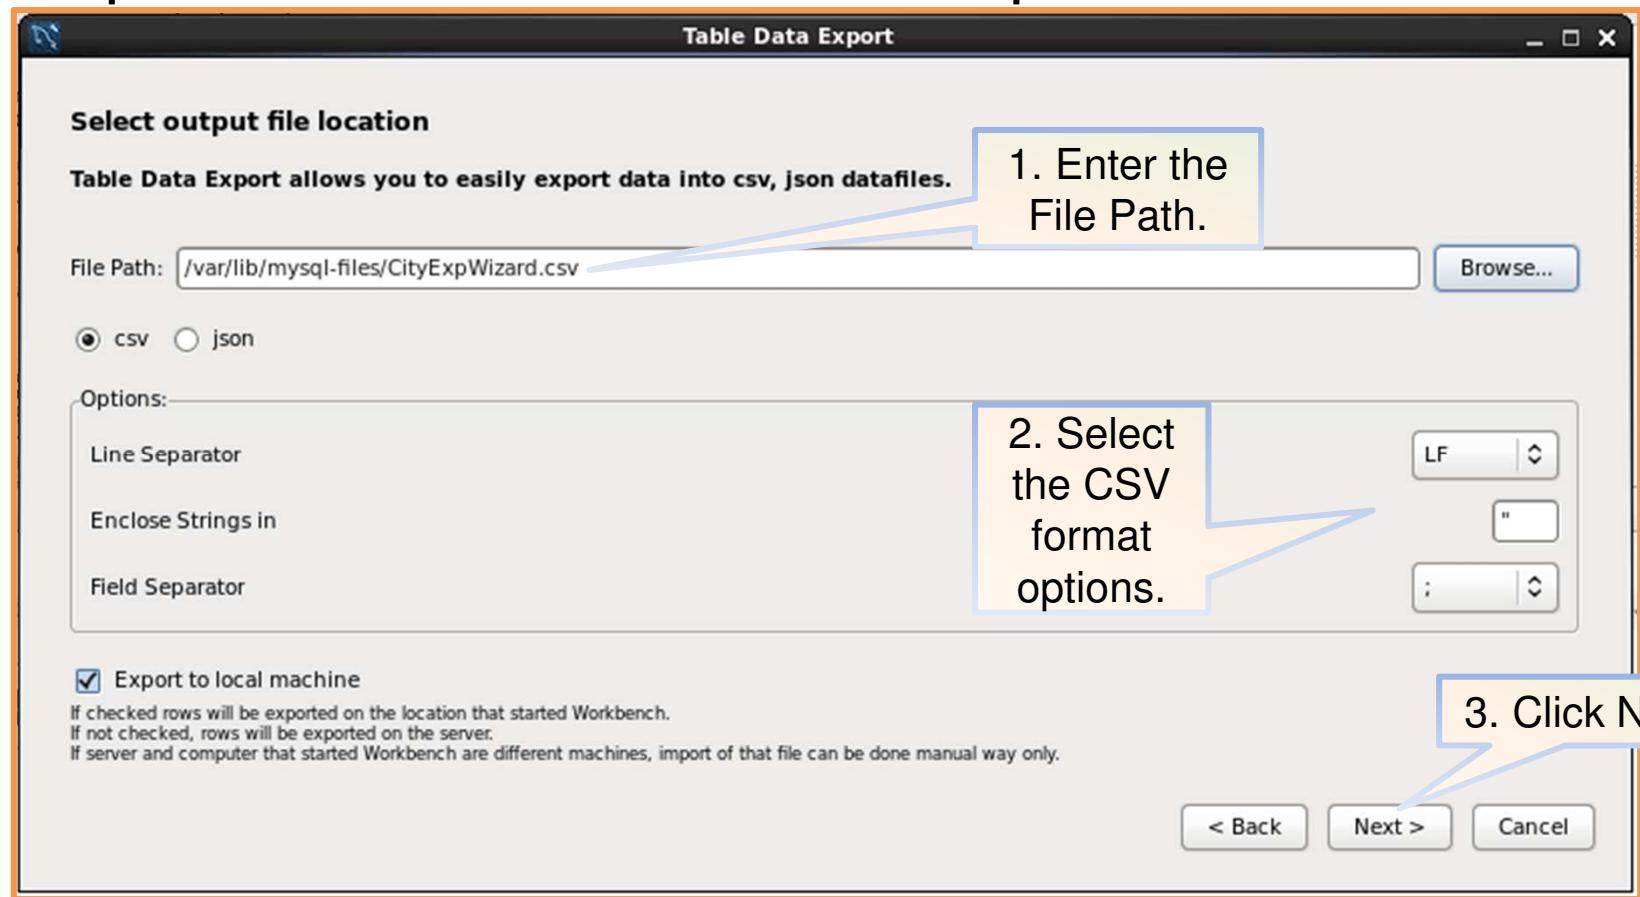
- ID, Name, CountryCode, District, Population
- 1, Kabul, AFG, Kabul, 1780000
- 2, Qandahar, AFG, Qandahar, 237500
- 3, Herat, AFG, Herat, 186800
- 4, Mazar-e-Sharif, AFG, Balkh, 127800
- ...
- 995, Probolinggo, IDN, "East Java", 120770
- 996, Cilegon, IDN, "West Java", 117000
- 997, Cianjur, IDN, "West Java", 114300
- 998, Ciparay, IDN, "West Java", 111500
- 999, Lhokseumawe, IDN, Aceh, 109600
- 1000, Taman, IDN, "East Java", 107000

# Starting the MySQL Workbench Table Data Export Wizard



# Selecting Options for the MySQL Workbench Table Data Export Wizard

- The Wizard exports data in CSV or JSON. Select options for the CSV format export.



# Contents of the CSV File from the MySQL Workbench Table Data Export Wizard

- Contents of the output file:

```
•"ID", "Name", "CountryCode", "District", "Population"  
•1, "Kabul", "AFG", "Kabol", 1780000  
•2, "Qandahar", "AFG", "Qandahar", 237500  
•3, "Herat", "AFG", "Herat", 186800  
•4, "Mazar-e-Sharif", "AFG", "Balkh", 127800  
•...  
•4074, "Gaza", "PSE", "Gaza", 353632  
•4075, "Khan Yunis", "PSE", "Khan Yunis", 123175  
•4076, "Hebron", "PSE", "Hebron", 119401  
•4077, "Jabaliya", "PSE", "North Gaza", 113901  
•4078, "Nablus", "PSE", "Nablus", 100231  
•4079, "Rafah", "PSE", "Rafah", 92020
```

# Importing Data from a Text File

---

- You can import a data file generated by MySQL or some other application into a MySQL table. You should know the following characteristics of the data file:
  - Column value separator
  - Order of the columns
  - Row separator
  - File system where the file resides
  - What characters enclose the values (example: double quotation marks)
  - Whether the file has a header row with the column names
  - If you need certain privileges to access the file
  - How you want to process any rows that have a unique key column that matches an existing value in the table



# Scenarios for Importing Data from a Text File

---

- A few of the possible scenarios for importing data from a text file include:
  - Populating a newly created, empty table
  - Adding data to an existing table
  - Updating data in an existing table
- When importing data into an existing table that has a primary key or unique key, rows that have the same value as the primary key or unique key cannot be added to the table. You can specify:
  - The data being imported replaces the row in the table with the same primary key or unique key.
  - The data being imported is ignored and the row in the table with the matching key is left unchanged.



# Importing with the LOAD DATA INFILE Statement

- **LOAD DATA INFILE** is the opposite of **SELECT...INTO OUTFILE**.
  - However, it uses similar clauses and format specifiers.
  - Syntax:

```
LOAD DATA INFILE '/path/filename'  
[REPLACE | IGNORE]  
INTO TABLE table  
[FIELDS TERMINATED BY 'delimiter'  
[OPTIONALLY] ENCLOSED BY 'character'  
LINES TERMINATED BY 'terminator' ]  
[IGNORE number LINES]  
[column_list];
```

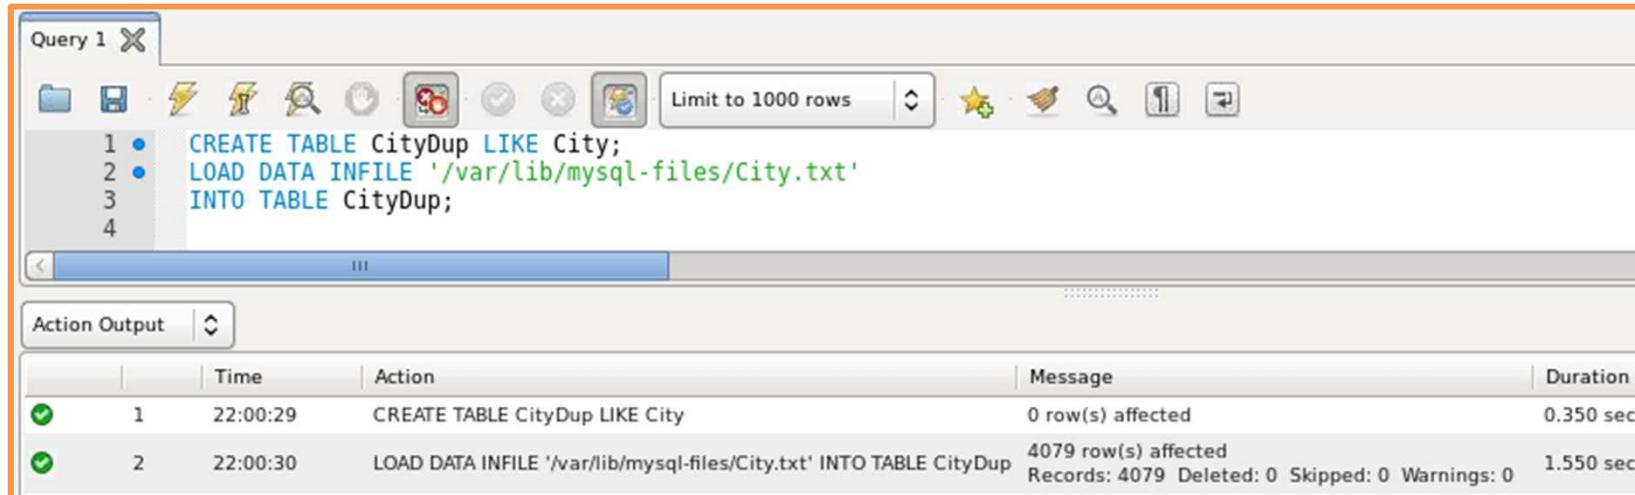
# Importing with LOAD DATA INFILE: File Control

---

- MySQL file assumptions and defaults:
  - File resides on the server host.
  - Fields are terminated by '\t'.
  - Each input line contains a value for each column in the table.
- Optional clauses give you more control:
  - Control of duplicate records/rows
  - Data file format specification (field and line terminators)
  - Skipping data file rows
  - Loading specific table columns



# Example of LOAD DATA INFILE into an Empty Table



The screenshot shows the MySQL Workbench interface with a query editor titled "Query 1". The query window contains the following SQL code:

```
1 • CREATE TABLE CityDup LIKE City;
2 • LOAD DATA INFILE '/var/lib/mysql-files/City.txt'
3   INTO TABLE CityDup;
4
```

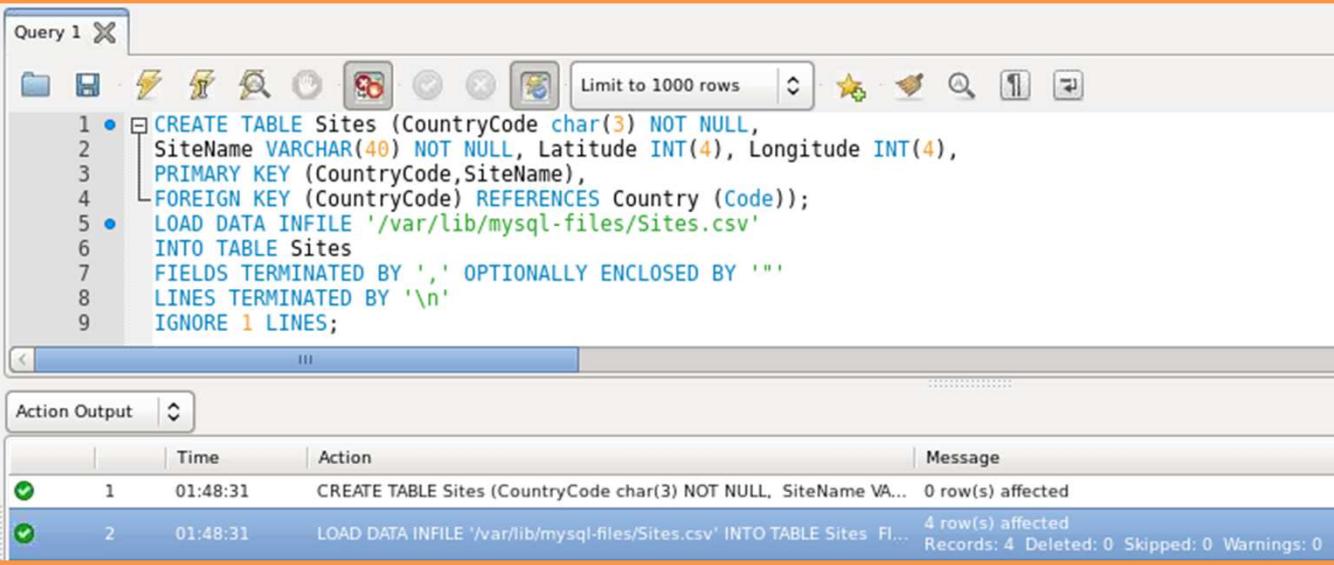
The "Action Output" tab displays the results of the execution:

	Time	Action	Message	Duration /
1	22:00:29	CREATE TABLE CityDup LIKE City	0 row(s) affected	0.350 sec
2	22:00:30	LOAD DATA INFILE '/var/lib/mysql-files/City.txt' INTO TABLE CityDup	4079 row(s) affected Records: 4079 Deleted: 0 Skipped: 0 Warnings: 0	1.550 sec

- Contents of the first few lines of the input file:

1	Kabul	AFG	Kabul	1780000
2	Qandahar	AFG	Qandahar	237500
3	Herat	AFG	Herat	186800
4	Mazar-e-Sharif	AFG	Balkh	
127800				
5	Amsterdam	NLD	Noord-Holland	
	731200			

# Example of LOAD DATA INFILE into an Empty Table from a CSV Formatted File



The screenshot shows the MySQL Workbench interface. The top part is a query editor titled "Query 1" containing the following SQL code:

```
1 • CREATE TABLE Sites (CountryCode char(3) NOT NULL,
2     SiteName VARCHAR(40) NOT NULL, Latitude INT(4), Longitude INT(4),
3     PRIMARY KEY (CountryCode,SiteName),
4     FOREIGN KEY (CountryCode) REFERENCES Country (Code));
5 • LOAD DATA INFILE '/var/lib/mysql-files/Sites.csv'
6     INTO TABLE Sites
7     FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY '"'
8     LINES TERMINATED BY '\n'
9     IGNORE 1 LINES;
```

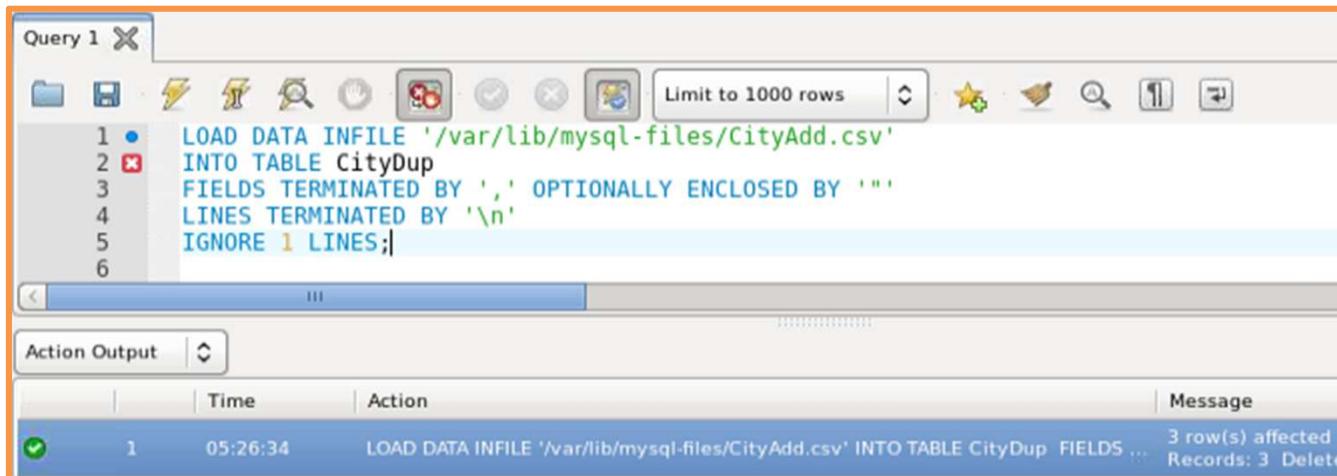
The bottom part is an "Action Output" table with two rows:

	Time	Action	Message
1	01:48:31	CREATE TABLE Sites (CountryCode char(3) NOT NULL, SiteName VA... 0 row(s) affected	
2	01:48:31	LOAD DATA INFILE '/var/lib/mysql-files/Sites.csv' INTO TABLE Sites Fl... 4 row(s) affected Records: 4 Deleted: 0 Skipped: 0 Warnings: 0	

- Contents of the input file:

```
"CountryCode", "SiteName", "Latitude", "Longitude"  
"FRA", "Eiffel Tower", 49, 2  
"USA", "Grand Canyon", 36, -114  
"IND", "Taj Mahal", 27, 78  
"AUS", "Sydney Opera House", -29, 167
```

# Example of LOAD DATA INFILE to Add Data to an Existing Table



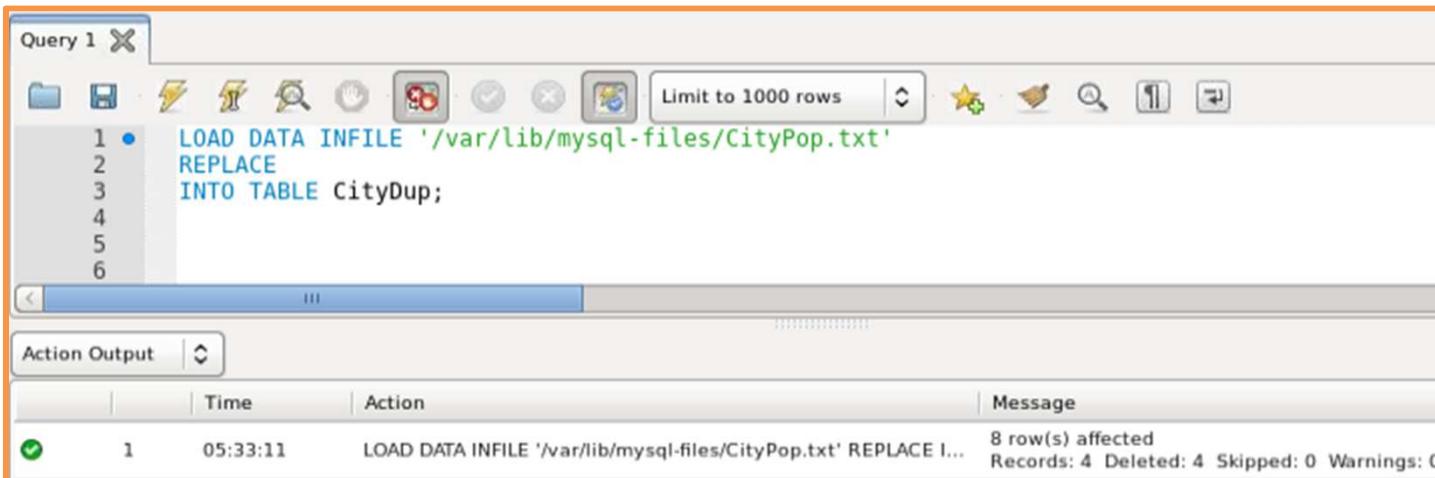
```
Query 1 X
LOAD DATA INFILE '/var/lib/mysql-files/CityAdd.csv'
INTO TABLE CityDup
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY '\"'
LINES TERMINATED BY '\n'
IGNORE 1 LINES;
```

Action Output	Time	Action	Message
1	05:26:34	LOAD DATA INFILE '/var/lib/mysql-files/CityAdd.csv' INTO TABLE CityDup FIELDS ...	3 row(s) affected Records: 3 Delete

- Contents of the input file:

- "ID", "Name", "CountryCode", "District", "Population"
- 4080, "Juneau", "USA", "Alaska", 32660
- 4081, "Dover", "USA", "Delaware", 37366
- 4082, "Helena", "USA", "Montana", 29596

# Example of LOAD DATA INFILE to Update Data in an Existing Table



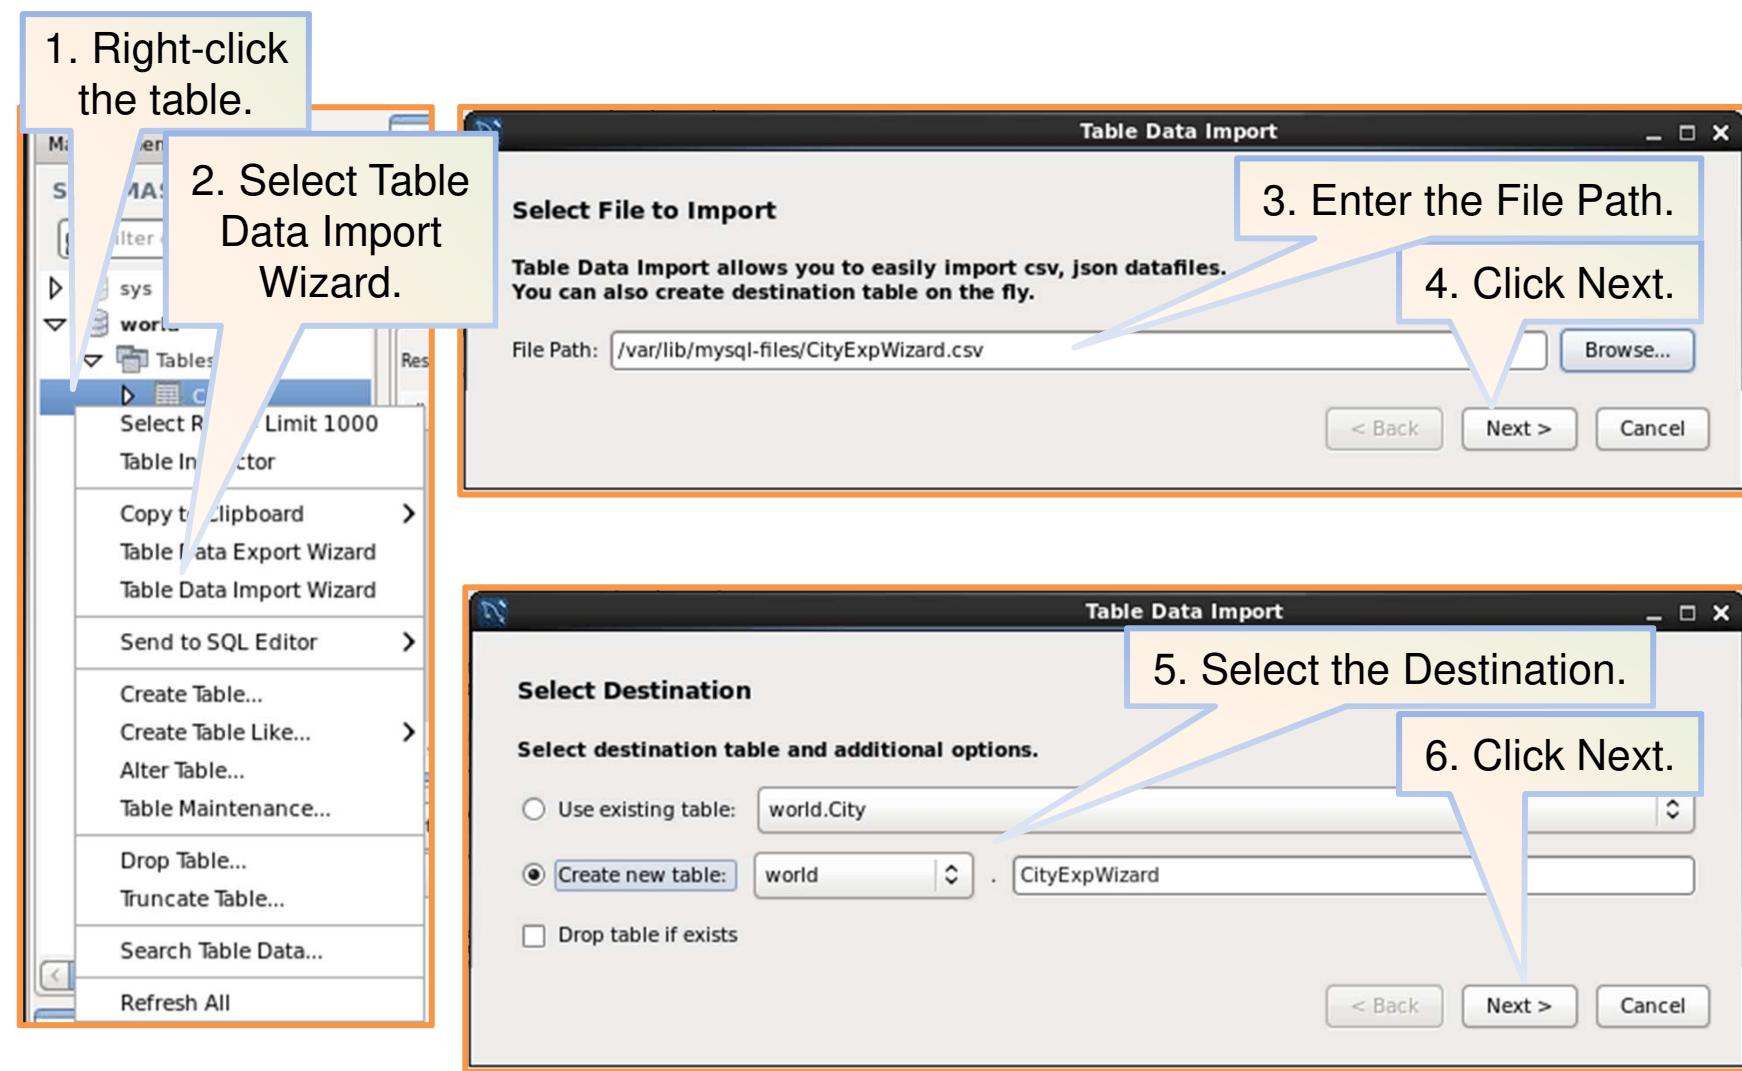
```
Query 1 X
LOAD DATA INFILE '/var/lib/mysql-files/CityPop.txt'
REPLACE
INTO TABLE CityDup;

Action Output
Time Action Message
1 05:33:11 LOAD DATA INFILE '/var/lib/mysql-files/CityPop.txt' REPLACE I... 8 row(s) affected
Records: 4 Deleted: 4 Skipped: 0 Warnings: 0
```

- Contents of the input file:

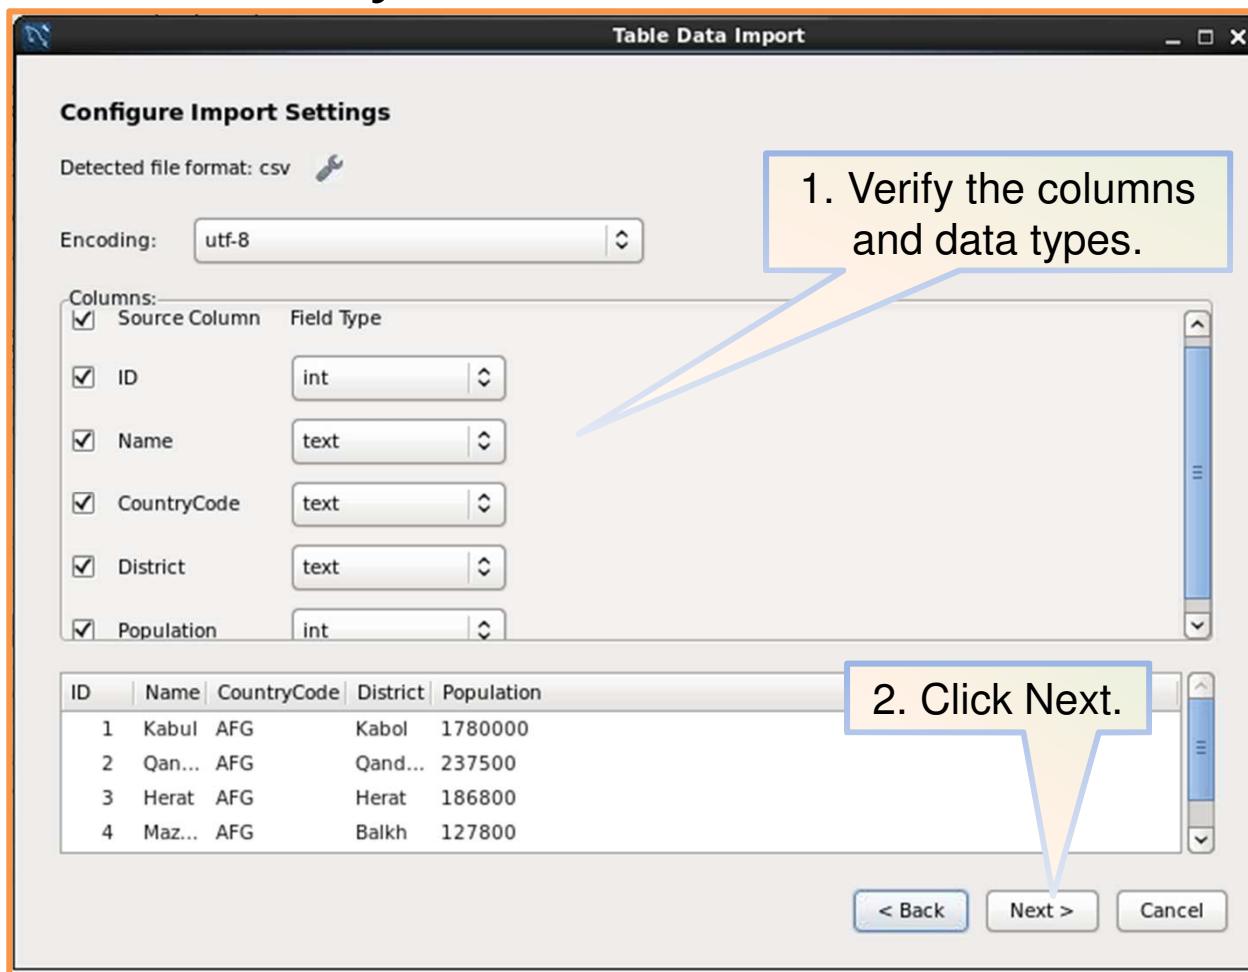
•1	Kabul	AFG	Kabol	3678000
•2	Qandahar		Qandahar	491500
•3	Herat	AFG	Herat	436300
•4	Mazar-e-Sharif		AFG	Balkh 693000

# Starting the MySQL Workbench Table Data Import Wizard



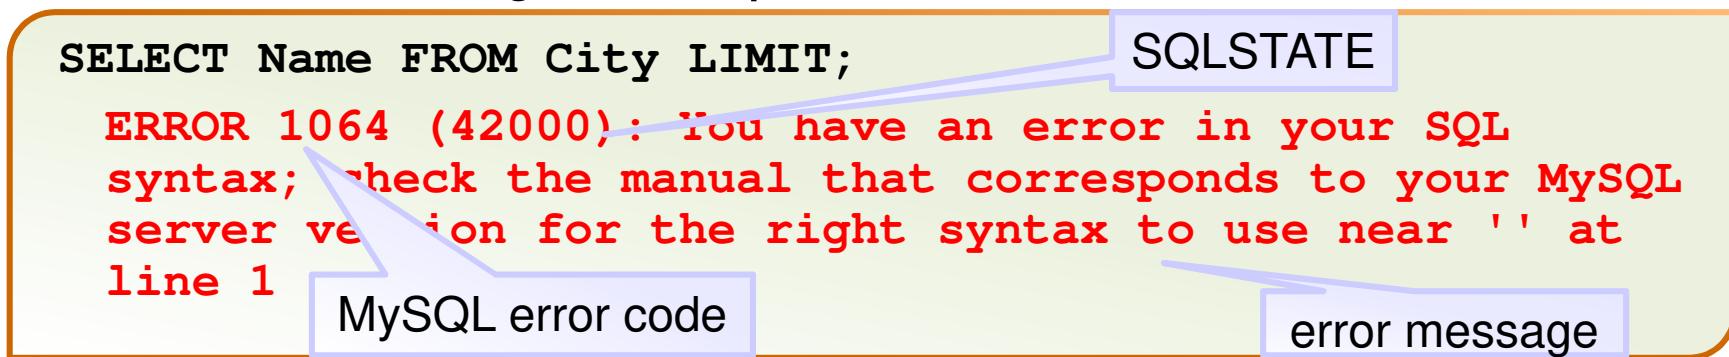
# Selecting Options for the MySQL Workbench Table Data Import Wizard

- The Wizard analyzes the file.



# Troubleshooting Messages

- MySQL provides error and warning messages.
- Error message example:



- This error is due to the missing `row_count` value for the `LIMIT` clause.

# Troubleshooting Message Options

---

- The following statements control display of errors and warnings:
  - **SHOW WARNINGS:** Displays any error, warning, and informational messages generated by the last statement
  - **SHOW ERRORS:** Displays error messages only
- Refer to the error code definitions in the *MySQL Reference Manual*.
- After you understand the cause of the error or warning, make any needed changes and run the SQL statement again.



# Index

# Index Types

---

- **Nonunique:** Values can appear multiple times in the index.
  - This is the default index type.
- **Unique:** Values must be unique (or NULL).
  - Primary keys are unique indexes and may not contain NULL.
- **Fulltext:** Values are character string data, and the index supports full text searching.
- **Spatial:** Values are spatial data types such as GEOMETRY, POINT, or POLYGON.



# MySQL Enterprise Monitor Query Analyzer

---

- Monitors SQL statements executed on the MySQL server:
  - Nature of query
  - Number of executions
  - Execution times
- Extracts this information from the Performance Schema using the MySQL Enterprise Monitor Agent by default
  - Can configure client applications to route database requests via MySQL Proxy and MySQL Aggregator
  - Can install a Connector plugin to send this information directly to the MySQL Enterprise Monitor Service Manager
- Normalizes queries for easier reporting
  - Combines similar queries with different literal values
- Enables you to drill into each query for detailed information



# Query Response Time Index

- MySQL Enterprise Monitor Query Analyzer generates a Query Response Time Index (QRTi) for each query.
  - Enables an “at-a-glance” indicator of query performance

Status	Default time value	Assigned Value	Meaning
Optimum	100ms	1.0 (100%)	All instances of the query were within the acceptable response time.
Acceptable	4 * Optimum: 100ms to 400ms	0.5 (50%)	Query execution time was less than 400 ms (the default acceptable response time).
Unacceptable	> Acceptable	0 (0%)	All instances of the query exceeded the acceptable response time.



# Query Analyzer User Interface

ORACLE MySQL Enterprise Monitor

Dashboards Events Query Analyzer Reports & Graphs Configuration Refresh: Off

Browse Queries

Query Analysis Reporting - Query Response Time Index (Aggregate)  
Zoom: 1h 2h 4h 6h 12h 1d 2d

Query Response Time Index

QR TI

11:15

qrqi

Query Response Time Index (QRTI)

Query	Database	Counts			Latency (hh:mm:ss)		
		Exec	Err	Warn	Total	Max	Avg
UPDATE `mem__inventory`...t` = ? WHERE `hid` = ? (1)	mem	1	0	0	1.00	0.000	0.000
SELECT COUNT (*) AS `...tate` IN (...) LIMIT ? (1)	mysql	2	0	0	1.00	0.017	0.012
INSERT INTO `mem__instr...fied` , `modified` ) (1)	mem	8	0	0	1.00	0.505	0.142
DELETE FROM `mem__instr...timestamp` < ? LIMIT ? (1)	mem	2	0	0	1.00	0.000	0.001
DELETE FROM `mem__instr...timestamp` < ? LIMIT ? (1)	mem	2	0	0	1.00	0.000	0.000
DELETE `es` FROM `mem__...D `e` . `isClosed` = ? (1)	mem	2	0	0	1.00	0.001	0.001
UPDATE `mem__inventory`...d` = ? WHERE `hid` = ? (1)	mem	2	0	0	1.00	0.001	0.004
UPDATE `mem__inventory`...n` = ? WHERE `hid` = ? (1)	mem	2	0	0	1.00	0.010	0.020

Copyright © 2005, 2013, Oracle and/or its affiliates. All rights reserved.

3.0.0.2880 - scissors (192.168.1.215) - Aug 22, 2013 11:34:45 am (Up Since: 2 days, 17 hours ago) - About

# Topics

---

- Identifying Slow Queries
- EXPLAIN Statement
- Working with Indexes
- Index Statistics



# Creating Indexes to Improve Query Performance

---

- Enables efficient access to specific rows
  - A query that searches for a value in an indexed field can immediately find rows containing that value.
    - A *seek*, which is efficient
  - A query that searches for a value in an unindexed field must read all rows to find rows containing that value.
    - A *scan*, which is inefficient
- Supports the following operations:
  - Direct value matching
    - Example: Find words such as `xenolith`.
  - Existence checking
    - Example: Determine that the word `xzzq` does not exist.
  - Range scans
    - Example: Find all words that begin with `xy`.



# Creating and Dropping Indexes on Existing Tables

---

- To change a table's primary key:

```
ALTER TABLE table DROP PRIMARY KEY,  
ADD PRIMARY KEY(col1, col2);
```

- To add a unique key:

```
ALTER TABLE table ADD UNIQUE(col3);  
CREATE UNIQUE INDEX index2 ON table(col4);
```

- To add an ordered (nonunique) index with a name:

```
CREATE INDEX index3 ON table(col5);
```

- To drop a nonprimary key by name:

```
ALTER TABLE table DROP INDEX indexname;  
DROP INDEX indexname ON table;
```



# Displaying Indexes with SHOW INDEXES FROM

```
mysql> SHOW INDEXES FROM departments\G
***** 1. row *****
      Table: departments
  Non_unique: 0
        Key_name: PRIMARY
  Seq_in_index: 1
    Column_name: dept_no
        Collation: A
   Cardinality: 9
      Sub_part: NULL
        Packed: NULL
        Null:
    Index_type: BTREE
        Comment:
Index_comment:
***** 2. row *****
      Table: departments
  Non_unique: 0
        Key_name: dept_name
  Seq_in_index: 1
    Column_name: dept_name
        Collation: A
   Cardinality: 9
      Sub_part: NULL
        Packed: NULL
        Null:
    Index_type: BTREE
        Comment:
Index_comment:
2 rows in set (0.00 sec)
```

